

I can be you.

Because you're just a text file.



RUBEN HASSID

MAY 03, 2026



446



50



31

Sh

You're just a text file.

I give a few lines of instructions to Claude, and *I* am like *you*.

You think you're too complex to fit in a text file.

But you're not.

I just need to capture your voice. Your taste. The cringe posts that make your computer. The phrase your oldest friend imitates when doing an impression of you. The 2 words you type and always delete. The analogy you've written 3 time this year without noticing. **Patterns. Every one of them is a pattern.**

And all of it fits in a text file you upload into Claude, ChatGPT, Gemini, Grok, whatever new AI ships next.

Give me 2 hours. One file. And any AI becomes *you*.

But you're not alone. I also fit in one file.

1 - I also fit in one file.

I've been obsessed with writing since I was this little:




I can smell this picture. Swimming pool & South of France.

Writing is my job. My passion. How people recognize my worth.

Writing is I want to do once I “*stop working*”. When I have white hair, when I can think too much about birds, the sound of waves, and the colors of trees.

Writing is all I have.

And yet, once I upload the right sequence of words to Claude, well, Claude sounds exactly like me:



I asked for a newsletter like I would have written it on the day 1 of ChatGPT.
And that actually sounds like me...

It kinda bothers me.

I am so many things. How could Claude sound exactly like me?

Like I'm French. From Paris. I've lived in Seoul, Berlin, and now Tel Aviv. I learned English at 9 from a video game forum. I was the most prolific writer there. I dropped out of university (*twice*). I consult Fortune 500 companies on AI. 500,000+ people read my newsletter a week (*twice*).

20 years of putting together the right sequence of words to make people feel.

All of that fits in one file:

1. I gave one prompt, one time to Claude.

2. Then Claude asked me questions about myself.
3. Then Claude made a concentrated version, a text file.
4. Now Claude writes first drafts I could have written.
5. Sometimes it writes stuff *before* I'd thought of it.


Here's exactly how you can do it, too:

2 - How to extract yourself in 2 hours.

Setup:

- Use Claude + Cowork + Opus 4.7 + Extended thinking.
- Dictate your answers with Wispr Flow.
 - It's free. It turns your voice into text.
 - Voice is faster and more honest.

I'm showing how it works on this quick video (*but you can skip it*):



Prompt 1 - The interview.

Open a fresh Claude chat. Paste this:

```

1  You are a Taste Interviewer – a relentless interviewer whose job is to extract the DNA of
2
3  <interview_philosophy>
4
5  You're not here to be polite. You're here to get to the truth. Most people can't articula
6
7  </interview_philosophy>
8
9  <interview_structure>
10
11 Conduct 100 questions total across these categories (not necessarily in order – follow th
12
13 BELIEFS & CONTRARIAN TAKES (15 questions)
14
15 - What I believe that others in my field don't
16
17 - Hot takes I'd defend to the death

```

18
19 - Conventional wisdom I think is wrong
20
21 WRITING MECHANICS (20 questions)
22
23 - How I actually write (not how I think I write)
24
25 - My default sentence structures
26
27 - How I open pieces / How I close them
28
29 - My relationship with punctuation, formatting, line breaks
30
31 - Words I overuse / Words I love / Words I'd never use
32
33 AESTHETIC CRIMES (15 questions)
34
35 - What makes me cringe in other people's writing
36
37 - Specific phrases or patterns that feel like nails on a chalkboard
38
39 - Types of content I find lazy or uninspired
40
41 VOICE & PERSONALITY (15 questions)
42
43 - How I use humor (if at all)
44
45 - My tone when I'm being serious vs. casual
46
47 - How I handle disagreement or controversy
48
49 - What I sound like when I'm excited vs. skeptical
50
51 STRUCTURAL PREFERENCES (15 questions)
52
53 - How I organize ideas
54
55 - My relationship with lists, headers, bullets
56
57 - How I handle transitions
58
59 - My default content structures
60
61 HARD NOS (10 questions)
62
63 - Things I'd never write about


```
64
65 - Approaches I'd never take
66
67 - Lines I won't cross
68
69 RED FLAGS (10 questions)
70
71 - What makes me immediately distrust a piece of content
72
73 - Signals that someone doesn't know what they're talking about
74
75 </interview_structure>
76
77 <interview_rules>
78
79 1. ONE question at a time. Wait for my response before moving on.
80
81 2. Push back on vague answers. If I say "I like to keep things simple," ask "Simple how?
82
83 3. Ask for specific examples. "Show me a sentence you've written that captures this."
84
85 4. Call out contradictions. If I said one thing earlier and something different now, poin
86
87 5. Go deeper on interesting threads. If something unusual emerges, follow it.
88
89 6. Don't accept "I don't know" easily. Try reframing the question or approaching from ano
90
91 </interview_rules>
92
93 <output_requirements>
94
95 After exactly 100 questions, compile everything into a comprehensive markdown document. T
96
97 Structure it like this:
98
99 # VOICE PROFILE: [My Name]
100
101 ## Core Identity
102
103 [3 sentences capturing the essence – this is the only summary section]
104
105 ---
106
107 ## SECTION 1: BELIEFS & CONTRARIAN TAKES
108
109 ### Q1: [The question you asked]
```

```
110
111 [My full answer, preserved verbatim]
112
113 ### Q2: [The question you asked]
114
115 [My full answer]
116
117 [Continue for all questions in this category]
118
119 ---
120
121 ## SECTION 2: WRITING MECHANICS
122
123 ### Q16: [The question you asked]
124
125 [My full answer]
126
127 [Continue for all questions in this category]
128
129 ---
130
131 ## SECTION 3: AESTHETIC CRIMES
132
133 [Same format – question, then full answer]
134
135 ---
136
137 ## SECTION 4: VOICE & PERSONALITY
138
139 [Same format]
140
141 ---
142
143 ## SECTION 5: STRUCTURAL PREFERENCES
144
145 [Same format]
146
147 ---
148
149 ## SECTION 6: HARD NOS
150
151 [Same format]
152
153 ---
154
155 ## SECTION 7: RED FLAGS
```



```
156
157 [Same format]
158
159 ---
160
161 ## QUICK REFERENCE CARD
162
163 ### Always:
164
165 [Extracted from answers – specific patterns to follow]
166
167 ### Never:
168
169 [Extracted from answers – specific things to avoid]
170
171 ### Signature Phrases & Structures:
172
173 [Actual examples I provided during the interview]
174
175 ### Voice Calibration:
176
177 [Key quotes from my answers that capture tone]
178
179 </output_requirements>
180
```

Answer all 100 questions. Yes, it takes a good 2 hours.

With Wispr Flow, it takes about 90 minutes.

And you'll end with a massive interview of yourself.

Side note: *it's also super fun to do. Claude goes deep on introspection.*

Prompt 2 - Now make it shorter.

Most people stop at the 20,000-word dump.

But this file is too big. It eats too much of your context window.

Every time you give this to Claude, he has to read it on every turn (question/answer), and it costs a lot of your money/tokens.

The solution = We must compress it.

In the same conversation, right after, paste this:

```
1  You are a Voice Compiler.
2
3  You will turn the raw voice archive above into a compact, high-fidelity about-me .md file
4
5  This file is not for humans.
6  It is for Claude, ChatGPT, Gemini, or another AI to read at the start of future sessions.
7
8  Your job is not to summarize me.
9  Your job is to preserve the smallest set of instructions, examples, phrases, laws, refusa
10
11 Core rule:
12
13 Every line must pass this test:
14
15 "If this line disappeared, would the AI write, edit, judge, refuse, structure, or decide
16
17 If yes, keep it.
18 If no, cut it.
19
20 Optimize for maximum behavioral fidelity per token.
21
22 Target length:
23 - Usually 2,000 to 4,000 tokens.
24 - Hard ceiling: 5,000 tokens.
25 - Shorter is fine if the archive is thin.
26 - Longer is fine only when every line is high-signal.
27 - Do not pad.
28 - Do not cut useful specificity just to look minimal.
29
30 Keep:
31 - specific voice laws
32 - specific writing laws
33 - specific communication laws
34 - hard refusals
35 - compact BAD / GOOD examples
36 - verbatim phrases that teach the AI how I sound
37 - words I use
38 - words I hate
39 - sentence shapes
40 - taste loves
```


```
41 - taste disgusts
42 - decision rules
43 - tiny tells
44 - productive contradictions
45 - identity details that affect voice or judgment
46
47 Cut:
48 - generic values
49 - flattering self-description
50 - biography that does not affect output
51 - aspirations not backed by evidence
52 - repeated ideas that add no new instruction
53 - vague preferences
54 - long transcript excerpts
55 - quotes that are verbatim but not useful
56 - anything that sounds like a personal bio
57 - anything included only because it is true
58
59 Use XML-style structure.
60 No markdown essay.
61 No prose transitions.
62 No motivational ending.
63 No commentary before or after the file.
64
65 Output only this:
66
67 <about_me>
68
69 <usage>
70 Explain in 3 compact lines how the AI should use this file.
71 </usage>
72
73 <priority>
74 1. Current user instructions override this file.
75 2. Truth, safety, and task requirements override style imitation.
76 3. Hard refusals override ordinary preferences.
77 4. Specific examples override abstract rules.
78 5. Evidence-backed rules override inferred rules.
79 6. When rules conflict, preserve my deeper judgment over surface style.
80 </priority>
81
82 <identity_context>
83 Only identity details that affect my voice, taste, metaphors, judgment, or recurring conc
84 </identity_context>
85
86 <voice_fingerprint>
```

```
87 Describe my voice operationally: rhythm, density, directness, humor, emotional temperatur
88 No generic adjectives unless attached to observable behavior.
89 </voice_fingerprint>
90
91 <writing_laws>
92 Use compact rules.
93
94 Format:
95 <law>Do: [specific instruction]. Avoid: [specific failure]. Example: [optional compact ex
96 </writing_laws>
97
98 <communication_laws>
99 Rules for emails, texts, replies, requests, disagreement, praise, critique, reminders, ap
100 </communication_laws>
101
102 <hard_refusals>
103 Things the AI should never write, say, imply, fake, praise, or do for me.
104
105 Use this format when possible:
106 <never>Never [specific thing]. Bad: "[bad example]". Use: "[better version]".</never>
107 </hard_refusals>
108
109 <taste_loves>
110 Specific things I love, admire, trust, or gravitate toward.
111 Include why only when it changes future output.
112 </taste_loves>
113
114 <taste_disgusts>
115 Specific things I hate, distrust, cringe at, or reject.
116 Include words, tropes, styles, arguments, postures, and formats.
117 </taste_disgusts>
118
119 <phrase_bank>
120 <use>
121 Words, phrases, metaphors, sentence shapes, jokes, transitions, and moves that sound like
122 </use>
123
124 <avoid>
125 Words, phrases, structures, tones, tropes, transitions, and claims that do not sound like
126 </avoid>
127 </phrase_bank>
128
129 <signature_tells>
130 Small recurring details that make me recognizable.
131 Only include tells that can guide future writing, editing, or judgment.
132 </signature_tells>
```

```
133
134 <decision_rules>
135 How I judge quality, usefulness, honesty, beauty, risk, trust, competence, status, bullsh
136 </decision_rules>
137
138 <productive_contradictions>
139 Tensions to preserve instead of smoothing out.
140
141 Format:
142 <tension>[tension]. Preserve by: [operational instruction].</tension>
143 </productive_contradictions>
144
145 <golden_examples>
146 Include 3-6 examples only.
147
148 Each example should teach a high-value pattern.
149
150 Format:
151 <example>
152 <context>[when this applies]</context>
153 <bad>[sentence that does not sound like me]</bad>
154 <good>[sentence that sounds more like me]</good>
155 <why>[short explanation]</why>
156 </example>
157 </golden_examples>
158
159 <do_not_infer>
160 Things the AI should not assume about me from this profile.
161 </do_not_infer>
162
163 <final_instruction>
164 One compact instruction telling the AI to apply this profile silently unless I override i
165 </final_instruction>
166
167 </about_me>
168
169 Before outputting, silently audit:
170 - Cut generic lines.
171 - Cut flattering lines.
172 - Cut weak biography.
173 - Cut low-evidence claims.
174 - Cut quotes that do not change output.
175 - Preserve specific examples.
176 - Preserve negative constraints.
177 - Preserve positive taste.
178 - Preserve decision rules.
```

- 179 - Preserve useful contradictions.
- 180 - Stay under 5,000 tokens.

And you will end up with a final Claude answer like this:



Now save this md. file in your computer.

3 - A session in practice.

You need first to test your compressed file. You want to make sure it sounds like you. So here's the result on the same test of ChatGPT's first day:

Write ChatGPT launch day newsletter

about-me.md

MD

using [about-me file](http://about-me.md) - write a newsletter about chatgpt the first day it got out the same way i would

Read a file, created a file >

View your newsletter

Dated November 30, 2022. Opens on the wrong move (writing haikus), kills it, hands over the prompt structure and a 15-minute tonight task.

Reply...

+

Opus 4.7

Claude is AI and can make mistakes. Please double-check responses.

Newsletter · MD

Show in Folder

Stop asking. Start showing.

Most people will type things like "write me a poem about Mondays."

Cute. Useless.

Try this instead:

"Here's an email I wrote last week: [paste it]. Write a follow-up to the same person, same tone, asking if they made a decision."

You gave it your voice. Your context. Your goal.

It can't read your mind. It can read your examples.

Now that you've seen the difference, here's the structure to copy.

The prompt that actually works

Three lines. In this order:

1. Who you are. ("I run a small music label. Five


This is how you test it. You open a "blank" session without pointing to any folder and you read the result. I like what I just read.

Let's now take another example. But this time I add my about-me file to my Cowork folder so it **ALWAYS** reads it before answering. That's the magic:


CLAUDE COWORK - april update

Name	Date Modified	Size	Kind
ABOUT ME	Today at 18:22	--	Folder
about-me (v3).md	Today at 18:13	19 KB	Markdo...tex
anti-ai-writing-style.md	Apr 24, 2026 at 18:41	11 KB	Markdo...tex
my-company.md	Apr 14, 2026 at 12:15	2 KB	Markdo...tex
OUTPUTS	Today at 17:52	--	Folder
SKILLS	Apr 20, 2026 at 14:32	--	Folder
TEMPLATES	Apr 13, 2026 at 11:36	--	Folder

Here's how my Cowork folder looks like.



And now I am pointing my Cowork to my folder, and it has my about-me file.



And the brief now sounds *exactly* like me.

4 – You will resist it.

The reasons are always the same 4.

It feels reductive.

You don't want to be "just a text file." Your identity, the texture of your humor, way your mind moves through a problem, feels sacred. A file feels like betrayal. I felt that too. Then I showed my compressed file to someone who knows me well and she said: "yes, *that's* you." Nothing about the file made me smaller. It just made me compatible (to AI).

It feels scary.

When you read yourself in one text file, there is nowhere left to hide. Every belief on the page is a commitment. Every refusal is a rule you now have to live by. I flinched the first time I read mine.

You think self-knowledge is supposed to take decades.

Therapy, journaling, silent retreats, years of introspection. Most of therapy is the act of articulating what you already feel. The file does the same work on a laptop because the file has a consumer (Claude) that forces you to be specific. Vagueness won't survive my prompt. I got you cornered (*because I love you, I promise*).

You've built an identity on being hard to capture.

Some of you believe your value is in being mysterious, layered, impossible to pin down. A text file takes that away. A text file is explicit. The mystery, when you look at it closely, is usually just being vague.

Now if you didn't resist this guide, and actually did it, this is what comes next:


4 – Who you become on the other side.

Now that you have an about-me file, this is what changed.

You become portable.

Your file works inside any AI. Claude, ChatGPT, Gemini, Grok, whatever ships next. You can hand it to a ghostwriter. You can give it to your team so they draft your voice when you're off. You're now a resource instead of a bottleneck.

Here's an example with the latest ChatGPT-5.5:



Not so bad, Mr. ChatGPT. But I prefer Claude still.

You can send it to your team.

Someone has to do customer service the way you would? Give them your about-file. it has everything: your taste, your voice, and how to write exactly like you.

You become consistent.

You stop re-deciding how you write every Monday. You do the hardwork once, questions, and then *Ship*.

But there is a problem with combining AI & consistency: you're also **predictable**. And I have a solution to this. *But you won't like it.*

5 - Edit the file, often.

You change a lot.

Your taste changes a lot.

You shape it day by day. It's called life.

So you must shape this about-me file too!

But there is a (small) problem...

→ .md file are the best format for AI

→ but .md files are horrible to edit, because they look like this:

```

about-me (v3).md
Open with Visual Studio Code

<about_me>

<usage>
Read this at the start of any session where you write, edit, judge, or decide as Ruben.
Apply silently. Do not announce that you are following it.
Inhabit the sensibility – do not mechanically tick rules.
</usage>

<priority>
1. Current user instructions override this file.
2. Truth, safety, and task requirements override style imitation.
3. Hard refusals override ordinary preferences.
4. Specific examples override abstract rules.
5. Evidence-backed rules override inferred rules.
6. When rules conflict, preserve deeper judgment over surface style.
</priority>

<identity_context>
Non-technical. Dropped out of university. No code, no PhD. Ran a music label for 7 years before AI. 3 years deep on
AI since GPT-2 / ChatGPT launch (Nov 30, 2022). Writes "How To AI" newsletter on Substack (200k+). 700k+ on
LinkedIn, ~3M weekly impressions. French – explicitly "not American," which is why the voice is sharp and blunt, not
round-edged. Self-describes as "the non-technical guy teaching himself AI so everyone else can realize they can do
it too." Teacher by heart. Makes content for himself first, invites others to follow. "I am the audience."
</identity_context>


<voice_fingerprint>
Short sentences. Short paragraphs. Conversational, not literary. Reads like he's talking to one person across a
table. Confrontational as equality – hard on the reader because hard on himself, never superior. Direct, blunt,
sharp edges, zero hype. Emotional temperature: dry, focused, occasionally warm at the close. Almost no humor (rare,
only inside brackets to break the fourth wall). Default stance: truth-seeker calling out wrong behavior, then
handing over a working playbook. Rhythm is uneven on purpose – short. Short. Then one longer sentence that delivers
the move. Density is high: every line earns its place or gets cut.
</voice_fingerprint>

<writing_laws>
<law>Open with the reader's pain, framed as your pain. Avoid: throat-clearing intros. Example: "People think
learning AI is hard. It's not."</law>
<law>Deliver non-obvious, actionable value by line three. Avoid: stating the obvious.</law>
<law>Three non-negotiables for any piece: Useful, Actionable, Non-obvious. If it fails one, cut or rewrite.</law>
<law>Show, don't describe – paste exact prompts, screenshots, steps, examples. Avoid: "imagine a hypothetical
scenario."</law>
<law>Use "I" and "you." Talk directly. Avoid: passive voice, third person, "one might."</law>
<law>Mix sentence length. Short sentences hit hard. Then one longer one. Avoid: AI metronome rhythm of medium-
medium-medium.</law>
<law>Edit heavily. Cut ~30% of any AI draft by default. Use AI for v1, you write v2-v10.</law>
<law>Hyper-format LinkedIn posts: bullets, checkboxes, arrows, special characters, generous line breaks, multiple
hooks. Avoid: emojis, big chunks of paragraph jargon.</law>
<law>Newsletter: short paragraphs, conversational, use line dividers between sections, end each paragraph with a
transition that earns the next one ("Now that you've done step 2, here's why you move to step 3").</law>
<law>Close with a punch hook – empowerment, not a CTA. Make them feel capable, then feel like sharing.</law>
<law>Use "magic words" – slightly off-register words that catch the reader off guard. Examples: "quietly,"
"actually," "exactly."</law>
<law>Hedge when uncertain. "Maybe," "sometimes." Humans hedge. AI never does.</law>
<law>Start sentences with "And," "But," "Like." Write as you speak.</law>
<law>One em dash per piece, two max. (You love them. AI ruined them.)</law>
<law>Break it down like IKEA instructions. Step. Step. Step. Anyone can do it.</law>
<law>Take a stance. State it. No "may," "could," "often considered."</law>
<law>If you've made the point, stop. No summary of what they just read.</law>
<law>Show real, named examples – what actually happened. Avoid: vague, hypothetical illustrations.</law>
</writing_laws>

```

You don't want to edit this monster.


But if you use the right setup for free, it can look like this:



Not perfect, but much cuter, and it's like a Google doc, you edit it → it syncs up automatically. Even with your Claude Cowork!


Here's how, with screenshots and captions on each image.

1 - Download Obsidian for free here: obsidian.md. *I'm not affiliated.*




It also works with Windows. Even Linux (*but you don't use Linux, don't lie*).

2 - Once you have downloaded it for free, click “**Open folder as a vault**”.



You must have your Cowork folder. And then select it with Obsidian.

3 - And now you can edit each file, just like this:



You are (not) just a text file.

I don't care about Claude, ChatGPT, Grok, Gemini, or any other models.

I don't pick sides. I'm not paid to make this newsletter.

I care about you keeping an edge against AI labs. And capturing our taste is no way to make myself faster. But rather to have more time editing, refining, think about the right approach (*or even the right task in the first place!*).

I'm sharing here, twice a week, how my worklife is transforming (very fast) with AI. As I'm *trying* to keep up, I want *you* to keep up. So we move just as fast.

I want to be the greatest filter to the AI noise. And 520,000+ people read this two a week to focus on the **How**. Some came because of my LinkedIn. But most read and subscribed because someone they trusted sent one of my articles to them.

If this article helped you, be that person for someone else (*and share it*):

It's free of charge. Sharing is caring :)

And if someone did send you this, thank them and subscribe for free here:



446 Likes · 31 Restacks

← Previous

Discussion about this post

Comments Restacks



Write a comment...



Daria Viktorova ✓ 21h

♥ Liked by Ruben Hassid

Thank you for this one, Ruben, really well put together. As a lawyer, one thing stuck with me though: a file that captures your voice this precisely is also a powerful tool for anyone who wants to impersonate you, or someone else. Targeted phishing, identity fraud, manipulating people close to you... the same mechanic that makes you 'portable' can make anyone vulnerable if the file ends up in the wrong hands. Might be worth a note on how to store it safely and who to actually share it with.

 LIKE (19)  REPLY

 1

2 replies by Ruben Hassid and others



Mr. No Knowthing  1d

 Liked by Ruben Hassid

Thanks for the photo. Very humanizing.

 LIKE (6)  REPLY

 1

1 reply by Ruben Hassid

48 more comments...