

# AI-ML for Mechanical Engineers

## Yogesh Haribhau Kulkarni

### Session 4: Pandas

#### Introduction to Pandas

##### Pandas

**pandas** introduces two new data structures to Python - **Series** and **DataFrame**, both of which are built on top of NumPy. Throughout this session, lines starting with `>>>` (or `...` for a continued line) are what you type at the Python prompt; everything else is the output pandas prints back.

```
>>> import pandas as pd
>>> import numpy as np
>>> import matplotlib.pyplot as plt
>>> pd.set_option('display.max_columns', 50)
```

##### Series

Series is a one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.). The axis labels are collectively referred to as the index. The basic method to create a Series is to call:

Here, data can be many different things:

- a Python dict
- an ndarray
- a scalar value (like 5)

```
>>> s = pd.Series(data, index=index)
```

##### Creating a Series from a List

- A Series is a one-dimensional object similar to an array, list, or column in a table.
- It will assign a labeled index to each item in the Series.
- By default, each item will receive an index label from 0 to N, where N is the length of the Series minus one.

```
>>> # create a Series with an arbitrary list
>>> s = pd.Series([7, 'Heisenberg', 3.14,
                  -1789710578, 'Happy Eating!'])
>>> s
```

##### Series

Output from Previous Slide

```
0          7
1    Heisenberg
2         3.14
3   -1789710578
4    Happy Eating!
dtype: object
```

##### Series with a Custom Index

Alternatively, you can specify an index to use when creating the Series.

```
>>> s = pd.Series([7, 'Heisenberg', 3.14,
                  -1789710578,
                  'Happy Eating!'],
                  index=['A', 'Z', 'C', 'Y', 'E'])
>>> s

A          7
Z    Heisenberg
C         3.14
Y   -1789710578
E    Happy Eating!
dtype: object
```

##### Series

The Series constructor can convert a dictionary as well, using the keys of the dictionary as its index.

```
>>> d = {'Chicago': 1000, 'New York': 1300,
        'Portland': 900, 'San Francisco':
        1100, 'Austin': 450, 'Boston': None}
>>> cities = pd.Series(d)
>>> cities

Austin          450
Boston          NaN
Chicago        1000
New York       1300
Portland         900
San Francisco   1100
dtype: float64
```

##### Series

You can use the index to select specific items from the Series

```
>>> cities['Chicago']
1000.0
```

##### Series

```
>>> cities[['Chicago', 'Portland', 'San
          Francisco']]
Chicago          1000
Portland          900
San Francisco    1100
dtype: float64
```

##### Series

You can use *boolean indexing* for selection.

```
>>> cities[cities < 1000]
Austin          450
Portland         900
dtype: float64
```

##### Boolean Indexing, Explained

That last one might be a little strange, so let's make it more clear - `cities < 1000` returns a Series of `True/False` values, which we then pass to our Series `cities`, returning the corresponding `True` items.

##### Boolean Indexing: Selecting Values

```
>>> less_than_1000 = cities < 1000
>>> print(less_than_1000)
>>> print('\n')
>>> print(cities[less_than_1000])
Austin          True
Boston          False
Chicago          False
New York         False
Portland          True
San Francisco    False
dtype: bool

Austin          450
Portland         900
dtype: float64
```

## Modifying Series Values by Index

You can also change the values in a Series on the fly.

```
>>> # changing based on the index
>>> print('Old value:', cities['Chicago'])
>>> cities['Chicago'] = 1400
>>> print('New value:', cities['Chicago'])
Old value: 1000.0
New value: 1400.0
```

## Modifying Values via Boolean Logic

Changing values using boolean logic

```
>>> print(cities[cities < 1000])
>>> print('\n')
>>> cities[cities < 1000] = 750
>>> print(cities[cities < 1000])
Austin      450
Portland    900
dtype: float64

Austin      750
Portland    750
dtype: float64
```

## Working with Series

What if you aren't sure whether an item is in the Series? You can check using idiomatic Python.

```
>>> print('Seattle' in cities)
>>> print('San Francisco' in cities)
False
True
```

## Series Arithmetic with Scalars

Mathematical operations can be done using scalars and functions.

```
>>> # divide city values by 3
>>> cities / 3
Austin      250.000000
Boston      NaN
Chicago     466.666667
New York    433.333333
Portland    250.000000
San Francisco 366.666667
dtype: float64
```

## Series Arithmetic with NumPy Functions

```
>>> # square city values
>>> np.square(cities)
Austin      562500
Boston      NaN
Chicago     1960000
New York    1690000
Portland    562500
San Francisco 1210000
dtype: float64
```

## Adding Two Series Together

You can add two Series together, which returns a union of the two Series with the addition occurring on the shared index values. Values on either Series that did not have a shared index will produce a NULL/NaN (not a number).

```
>>> print(cities[['Chicago', 'New York',
                  'Portland']])
>>> print('\n')
>>> print(cities[['Austin', 'New York']])
>>> print('\n')
>>> print(cities[['Chicago', 'New York',
                  'Portland']] + cities[['Austin', 'New York']])
```

## Adding Two Series: Output

```
Chicago     1400
New York    1300
Portland     750
dtype: float64

Austin      750
New York    1300
dtype: float64

Austin      NaN
Chicago     NaN
New York    2600
Portland    NaN
dtype: float64
```

## Working with Series

### NULL Checking

- Notice that because Austin, Chicago, and Portland were not found in both Series, they were returned with NULL/NaN values.
- NULL checking can be performed with `isnull()` and `notnull()`.

## Intuition: Index Alignment

### Intuition

Adding two Series is like matching up name tags at two different parties and adding values only where the *same* name appears at both. Anyone who only showed up to one party (an index value present in just one Series) gets a NaN back: there's nothing on the other side to add them to.

## Checking for Non-Null Values

Return a boolean series indicating which values aren't NULL

```
>>> cities.notnull()
Austin      True
Boston      False
Chicago     True
New York    True
Portland    True
San Francisco True
dtype: bool
```

## Checking for Null Values

Using boolean logic to grab the NULL cities

```
>>> print(cities.isnull())
>>> print('\n')
>>> print(cities[cities.isnull()])
Austin      False
Boston      True
Chicago     False
New York    False
Portland    False
San Francisco False
dtype: bool

Boston      NaN
dtype: float64
```

## Quick Check: Series

### Think About It

`cities['Chicago']` returns a single value, but `cities[['Chicago', 'Portland']]` returns a Series. What's the difference in the syntax, and why does it matter?

### Answer

A single string key does scalar/label lookup and returns *one* value. A *list* of keys is fancy indexing and always returns a Series, even with just one element: the double brackets are doing two things at once: the outer `[]` is the indexing operation, the inner `[]` is a literal Python list being passed as the key.

# Dataframes in Pandas

## Creating Dataframes

Creating a DataFrame by passing a numpy array, with a date-time index and labeled columns:

```
>>> dates = pd.date_range('20130101', periods=6)
>>> dates
DatetimeIndex(['2013-01-01', '2013-01-02',
               '2013-01-03', '2013-01-04',
               '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')

>>> df_demo = pd.DataFrame(np.random.randn(6, 4),
                           index=dates, columns=list('ABCD'))
>>> df_demo
```

|            | A         | B         | C         | D         |
|------------|-----------|-----------|-----------|-----------|
| 2013-01-01 | 0.469112  | -0.282863 | -1.509059 | -1.135632 |
| 2013-01-02 | 1.212112  | -0.173215 | 0.119209  | -1.044236 |
| 2013-01-03 | -0.861849 | -2.104569 | -0.494929 | 1.071804  |
| 2013-01-04 | 0.721555  | -0.706771 | -1.039575 | 0.271860  |
| 2013-01-05 | -0.424972 | 0.567020  | 0.276232  | -1.087401 |
| 2013-01-06 | -0.673690 | 0.113648  | -1.478427 | 0.524988  |

## Creating Dataframes

Creating a DataFrame by passing a dict of objects that can be converted to series-like.

```
>>> df2 = pd.DataFrame({'A': 1.,
...                     'B':
                        pd.Timestamp('20130102'),
...                     'C': pd.Series(1,
                        index=list(range(4)), dtype='float32'),
...                     'D': np.array([3] * 4,
                        dtype='int32'),
...                     'E':
                        pd.Categorical(["test", "train", "test",
                        "train"]),
...                     'F': 'foo'})
>>> df2
```

|   | A   | B          | C   | D | E     | F   |
|---|-----|------------|-----|---|-------|-----|
| 0 | 1.0 | 2013-01-02 | 1.0 | 3 | test  | foo |
| 1 | 1.0 | 2013-01-02 | 1.0 | 3 | train | foo |
| 2 | 1.0 | 2013-01-02 | 1.0 | 3 | test  | foo |
| 3 | 1.0 | 2013-01-02 | 1.0 | 3 | train | foo |

## Reading Data from Files

Every DataFrame so far has been typed in by hand. In practice, almost all data starts life in a file: `pd.read_csv()` and `pd.read_excel()` load it straight into a DataFrame, inferring column names and types automatically. `machine_logs.csv`

is provided upfront (no need to type it): copy it into your working folder before class.

```
>>> df = pd.read_csv('machine_logs.csv')
>>> df.head(3)
```

|   | machine_id | shift   | temperature | vibration | operator_notes |
|---|------------|---------|-------------|-----------|----------------|
| 0 | M1         | Morning | 72.5        | 0.021     | ok             |
| 1 | M2         | Morning | 75.1        | 0.034     | check bearing  |
| 2 | M1         | Evening | 78.3        | 0.041     | ok             |

## Quick Inspection

Before doing anything else with a newly loaded DataFrame, check its shape, column types, and non-null counts.

```
>>> df.shape
(6, 5)

>>> df.dtypes
machine_id      object
shift           object
temperature     float64
vibration       float64
operator_notes  object
dtype: object
```

## Quick Inspection: Full Structure

`df.info()` combines shape, dtypes, and non-null counts into one summary.

```
>>> df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   machine_id            6 non-null     object
1   shift                 6 non-null     object
2   temperature           6 non-null     float64
3   vibration             6 non-null     float64
4   operator_notes        6 non-null     object
dtypes: float64(2), object(3)
```

## Viewing Data

See the top & bottom rows of the frame

```
>>> df_demo.head()
```

|            | A         | B         | C         | D         |
|------------|-----------|-----------|-----------|-----------|
| 2013-01-01 | 0.469112  | -0.282863 | -1.509059 | -1.135632 |
| 2013-01-02 | 1.212112  | -0.173215 | 0.119209  | -1.044236 |
| 2013-01-03 | -0.861849 | -2.104569 | -0.494929 | 1.071804  |
| 2013-01-04 | 0.721555  | -0.706771 | -1.039575 | 0.271860  |
| 2013-01-05 | -0.424972 | 0.567020  | 0.276232  | -1.087401 |

```
>>> df_demo.tail(3)
```

|            | A         | B         | C         | D         |
|------------|-----------|-----------|-----------|-----------|
| 2013-01-04 | 0.721555  | -0.706771 | -1.039575 | 0.271860  |
| 2013-01-05 | -0.424972 | 0.567020  | 0.276232  | -1.087401 |
| 2013-01-06 | -0.673690 | 0.113648  | -1.478427 | 0.524988  |

## Viewing Data

Display the index, columns, and the underlying numpy data

```
>>> df_demo.index
DatetimeIndex(['2013-01-01', '2013-01-02',
               '2013-01-03', '2013-01-04',
               '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')

>>> df_demo.columns
Index(['A', 'B', 'C', 'D'], dtype='object')

>>> df_demo.values
array([[ 0.4691, -0.2829, -1.5091, -1.1356],
       [ 1.2121, -0.1732,  0.1192, -1.0442],
       [-0.8618, -2.1046, -0.4949,  1.0718],
       [ 0.7216, -0.7068, -1.0396,  0.2719],
       [-0.425 ,  0.567 ,  0.2762, -1.0874],
       [-0.6737,  0.1136, -1.4784,  0.525 ]])
```

## Viewing Data

Describe shows a quick statistic summary of your data

```
>>> df_demo.describe()
```

|       | A         | B         | C         | D         |
|-------|-----------|-----------|-----------|-----------|
| count | 6.000000  | 6.000000  | 6.000000  | 6.000000  |
| mean  | 0.073711  | -0.431125 | -0.687758 | -0.233103 |
| std   | 0.843157  | 0.922818  | 0.779887  | 0.973118  |
| min   | -0.861849 | -2.104569 | -1.509059 | -1.135632 |
| 25%   | -0.611510 | -0.600794 | -1.368714 | -1.076610 |
| 50%   | 0.022070  | -0.228039 | -0.767252 | -0.386188 |
| 75%   | 0.658444  | 0.041933  | -0.034326 | 0.461706  |
| max   | 1.212112  | 0.567020  | 0.276232  | 1.071804  |

## Viewing Data

Sorting by an axis

```
>>> df_demo.sort_index(axis=1, ascending=False)
```

|            | D         | C         | B         | A         |
|------------|-----------|-----------|-----------|-----------|
| 2013-01-01 | -1.135632 | -1.509059 | -0.282863 | 0.469112  |
| 2013-01-02 | -1.044236 | 0.119209  | -0.173215 | 1.212112  |
| 2013-01-03 | 1.071804  | -0.494929 | -2.104569 | -0.861849 |
| 2013-01-04 | 0.271860  | -1.039575 | -0.706771 | 0.721555  |
| 2013-01-05 | -1.087401 | 0.276232  | 0.567020  | -0.424972 |

```
2013-01-06  0.524988 -1.478427  0.113648 -0.673690
```

## Viewing Data

Sorting by values

```
>>> df_demo.sort_values(by='B')
              A          B          C          D
2013-01-03 -0.861849 -2.104569 -0.494929  1.071804
2013-01-04  0.721555 -0.706771 -1.039575  0.271860
2013-01-01  0.469112 -0.282863 -1.509059 -1.135632
2013-01-02  1.212112 -0.173215  0.119209 -1.044236
2013-01-06 -0.673690  0.113648 -1.478427  0.524988
2013-01-05 -0.424972  0.567020  0.276232 -1.087401
```

## Value Counts

`value_counts()` is the fastest way to see how a categorical column is distributed: here, how many readings each machine contributed.

```
>>> df['machine_id'].value_counts()
M1    2
M2    2
M3    2
Name: machine_id, dtype: int64
```

## Selection of Data

Selecting a single column, which yields a Series, equivalent to `df_demo.A`

```
>>> df_demo['A']
2013-01-01    0.469112
2013-01-02    1.212112
2013-01-03   -0.861849
2013-01-04    0.721555
2013-01-05   -0.424972
2013-01-06   -0.673690
Freq: D, Name: A, dtype: float64
```

## Selection of Data

Selecting via `[]`, which slices the rows.

```
>>> df_demo[0:3]
              A          B          C          D
2013-01-01  0.469112 -0.282863 -1.509059 -1.135632
2013-01-02  1.212112 -0.173215  0.119209 -1.044236
2013-01-03 -0.861849 -2.104569 -0.494929  1.071804

>>> df_demo['20130102':'20130104']
              A          B          C          D
2013-01-02  1.212112 -0.173215  0.119209 -1.044236
```

```
2013-01-03 -0.861849 -2.104569 -0.494929  1.071804
2013-01-04  0.721555 -0.706771 -1.039575  0.271860
```

## Selection of Data

Selecting on a multi-axis by label

```
>>> df_demo.loc[:, ['A', 'B']]
              A          B
2013-01-01  0.469112 -0.282863
2013-01-02  1.212112 -0.173215
2013-01-03 -0.861849 -2.104569
2013-01-04  0.721555 -0.706771
2013-01-05 -0.424972  0.567020
2013-01-06 -0.673690  0.113648
```

## Selection of Data

By integer slices, acting similar to numpy/python

```
>>> df_demo.iloc[3:5, 0:2]
              A          B
2013-01-04  0.721555 -0.706771
2013-01-05 -0.424972  0.567020
```

## Selection of Data

By lists of integer position locations, similar to the numpy/python style

```
>>> df_demo.iloc[[1, 2, 4], [0, 2]]
              A          C
2013-01-02  1.212112  0.119209
2013-01-03 -0.861849 -0.494929
2013-01-05 -0.424972  0.276232
```

## Intuition: Label vs. Position

### Intuition

Label-based selection is like finding a book by its title on the shelf: you ask for “Wednesday’s row” or “column A,” and pandas looks up that name. Position-based selection is like grabbing the third book from the left, regardless of what it’s called. Both answer valid but different questions, so pandas gives you separate ways to ask each one instead of guessing your intent.

## Boolean Indexing

Using a single column’s values to select data.

```
>>> df_demo[df_demo.A > 0]
              A          B          C          D
2013-01-01  0.469112 -0.282863 -1.509059 -1.135632
2013-01-02  1.212112 -0.173215  0.119209 -1.044236
2013-01-04  0.721555 -0.706771 -1.039575  0.271860
```

## Setting Values

Two changes happen here that carry into every slide from this point on: column D is overwritten entirely, and a new column F is added by assigning a Series that only covers *some* of the dates: alignment (the same mechanism from the Series section) fills in the rest with NaN.

```
>>> df_demo.loc[:, 'D'] = np.array([5] *
                                     len(df_demo))

>>> s1 = pd.Series([1, 2, 3, 4, 5, 6],
                   index=pd.date_range('20130102', periods=6))
>>> df_demo['F'] = s1
>>> df_demo
              A          B          C  D  F
2013-01-01  0.469112 -0.282863 -1.509059  5  NaN
2013-01-02  1.212112 -0.173215  0.119209  5  1.0
2013-01-03 -0.861849 -2.104569 -0.494929  5  2.0
2013-01-04  0.721555 -0.706771 -1.039575  5  3.0
2013-01-05 -0.424972  0.567020  0.276232  5  4.0
2013-01-06 -0.673690  0.113648 -1.478427  5  5.0
```

## Missing Data

Missing values show up as NaN. Blanking out an earlier value and reindexing to a smaller set of dates creates some gaps to demonstrate on:

```
>>> df_demo.loc[dates[0], ['A', 'B']] = 0
>>> df1 = df_demo.reindex(index=dates[0:4], columns=list(df_demo.columns) +
                          ['E'])
>>> df1.loc[dates[0]:dates[1], 'E'] = 1
>>> df1
              A          B          C  D  F  E
2013-01-01  0.000000  0.000000 -1.509059  5  NaN  1.0
2013-01-02  1.212112 -0.173215  0.119209  5  1.0  1.0
2013-01-03 -0.861849 -2.104569 -0.494929  5  2.0  NaN
2013-01-04  0.721555 -0.706771 -1.039575  5  3.0  NaN
```

## Missing Data: Drop

To drop any rows that have missing data:

```
>>> df1.dropna(how='any')
              A          B          C  D  F  E
2013-01-02  1.212112 -0.173215  0.119209  5  1.0  1.0
```

## Missing Data: Fill

Filling missing data:

```
>>> df1.fillna(value=5)
      A      B      C  D  F  E
2013-01-01  0.000000  0.000000 -1.509059  5  5.0  1.0
2013-01-02  1.212112 -0.173215  0.119209  5  1.0  1.0
2013-01-03 -0.861849 -2.104569 -0.494929  5  2.0  5.0
2013-01-04  0.721555 -0.706771 -1.039575  5  3.0  5.0
```

## Stats

Operations in general exclude missing data. Performing a descriptive statistic:

```
>>> df_demo.mean()
A    -0.004474
B    -0.383981
C    -0.687758
D      5.000000
F      3.000000
dtype: float64
```

## Stats: The Other Axis

Same operation on the other axis:

```
>>> df_demo.mean(1)
2013-01-01    0.872735
2013-01-02    1.431621
2013-01-03    0.707731
2013-01-04    1.395042
2013-01-05    1.883656
2013-01-06    1.592306
Freq: D, dtype: float64
```

## Apply

Applying functions to the data

```
>>> df_demo.apply(np.cumsum)
      A      B      C  D  F
2013-01-01  0.000000  0.000000 -1.509059  5  NaN
2013-01-02  1.212112 -0.173215 -1.389850 10  1.0
2013-01-03  0.350263 -2.277784 -1.884779 15  3.0
2013-01-04  1.071818 -2.984555 -2.924354 20  6.0
2013-01-05  0.646846 -2.417535 -2.648122 25 10.0
2013-01-06 -0.026844 -2.303886 -4.126549 30 15.0

>>> df_demo.apply(lambda x: x.max() - x.min())
A    2.073961
B    2.671590
C    1.785291
D    0.000000
F    4.000000
dtype: float64
```

## Renaming & Dropping Columns

Cleaning up column names and removing ones you don't need is usually the first step of any real analysis.

```
>>> df = df.rename(columns={'vibration':
                             'vibration_mm_s'})
>>> df = df.drop(columns=['operator_notes'])
>>> df.columns
Index(['machine_id', 'shift', 'temperature',
       'vibration_mm_s'], dtype='object')
```

## Creating New Columns

New columns are computed from existing ones with plain vectorized arithmetic and comparisons, no loops needed.

```
>>> df['vibration_in'] = df['vibration_mm_s'] /
25.4
>>> df['high_vibration'] = df['vibration_mm_s'] >
0.025
>>> df[['machine_id', 'vibration_mm_s',
       'high_vibration']]

   machine_id  vibration_mm_s  high_vibration
0          M1           0.021             False
1          M2           0.034              True
2          M1           0.041              True
3          M2           0.028              True
4          M3           0.019             False
5          M3           0.052              True
```

## Quick Check: DataFrames

### Think About It

`df['A']` returns a Series (a column), but `df[0:2]` returns a DataFrame (the first two rows). Same `[]` operator, same object: why does pandas treat a string key and a slice so differently?

### Answer

The plain `[]` operator on a DataFrame is intentionally overloaded: a single label looks up a *column* (matching the `df.A` shortcut), while a slice looks up *rows*, mirroring plain Python/NumPy slicing. Pandas infers your intent from the *type* of the key rather than a fixed row/column rule, which is exactly why `.loc/.iloc` exist: they remove the ambiguity by always meaning “by label” or “by position,” regardless of what kind of key you pass.

## GroupBy: Split-Apply-Combine

`groupby()` splits rows into groups by a key column. Nothing is computed until you chain on an aggregate.

### Intuition

Think of `groupby()` as sorting mail into pigeonholes by recipient. Nothing gets read or totalled yet: that's just the sorting step. The real work happens when a clerk (an aggregate like `.mean()` or `.sum()`) processes each pigeonhole separately and the results are stacked back into a single tray.

```
>>> df.groupby('machine_id')
<pandas.core.groupby.generic.DataFrameGroupBy
object at 0x...>
```

## GroupBy with Aggregation

Chain `.mean()`, `.sum()`, or `.agg()` to actually run the per-group computation.

```
>>> df.groupby('machine_id')['temperature'].mean()
M1    75.40
M2    74.55
M3    75.50
Name: temperature, dtype: float64

>>> df.groupby('machine_id').agg({'temperature':
                                'mean', 'vibration_mm_s': 'max'})
      temperature  vibration_mm_s
machine_id
M1             75.40           0.041
M2             74.55           0.034
M3             75.50           0.052
```

## Quick Check: GroupBy

### Think About It

`df.groupby('machine_id')` on its own prints only an object description, not a table of numbers. Why doesn't pandas compute anything at that point?

### Answer

`groupby()` only performs the *split*: it doesn't yet know what you want to *apply* to each group. Waiting until you chain on `.mean()`, `.sum()`, or `.agg()` means the same grouped object can be reused for several different aggregations without repeating the split step, and pandas never computes a result you never asked for.

## Concatenating DataFrames

`pd.concat()` stacks DataFrames with the same columns on top of each other, useful when the same kind of data arrives in separate batches (e.g., one file per shift).

```
>>> morning = df[df['shift'] ==
...     'Morning'][['machine_id', 'temperature']]
>>> evening = df[df['shift'] ==
...     'Evening'][['machine_id', 'temperature']]
>>> pd.concat([morning, evening])
  machine_id  temperature
0          M1           72.5
1          M2           75.1
4          M3           69.8
2          M1           78.3
3          M2           74.0
5          M3           81.2
```

## Merging DataFrames

`pd.merge()` joins two DataFrames on a shared key column: like a VLOOKUP, or a SQL JOIN.

```
>>> machine_info = pd.DataFrame({
...     'machine_id': ['M1', 'M2', 'M3'],
...     'machine_type': ['Lathe', 'Mill', 'Drill']
... })
>>> pd.merge(df, machine_info, on='machine_id')[
...     ['machine_id', 'machine_type',
...     'temperature']].head(3)
  machine_id machine_type  temperature
```

|   |    |       |      |
|---|----|-------|------|
| 0 | M1 | Lathe | 72.5 |
| 1 | M2 | Mill  | 75.1 |
| 2 | M1 | Lathe | 78.3 |

## Plotting with Pandas

`.plot()` is a thin wrapper around the `matplotlib` imported all the way back on the very first Pandas slide, no separate charting library needed for a first look at the data.

```
>>> avg_temp =
...     df.groupby('machine_id')['temperature'].mean()
>>> avg_temp.plot(kind='bar')
>>> plt.ylabel('Avg Temperature (deg C)')
>>> plt.show()
```

## Putting It Together: Machine Health Check

The same five-step pattern (read, inspect, clean, transform, summarize) shows up in almost every dataset you'll touch in this course.

```
>>> df = pd.read_csv('machine_logs.csv')
>>> df = df.rename(columns={'vibration':
...     'vibration_mm_s'})
>>> df = df.drop(columns=['operator_notes'])
>>> df['high_vibration'] = df['vibration_mm_s'] >
...     0.025
>>> summary = df.groupby('machine_id').agg(
```

```
...     avg_temp=('temperature', 'mean'),
...     flagged_readings=('high_vibration', 'sum')
... )
>>> summary
      avg_temp  flagged_readings
machine_id
M1           75.40                1
M2           74.55                2
M3           75.50                1
```

## Putting It Together: Visualizing the Summary

One line turns the summary table into a chart flagging which machine needs a closer look.

```
>>> summary['avg_temp'].plot(kind='bar',
...     color='steelblue')
>>> plt.title('Average Temperature by Machine')
>>> plt.ylabel('deg C')
>>> plt.show()
```

## Try It Yourself

Using the same `machine_logs.csv` you already have, no new data or setup needed:

- Which machine has the highest average vibration?
- Filter the DataFrame to only the Evening shift rows.
- Add a column flagging rows where `temperature > 76`.