Zero to Hero in Retrieval Augmented Generation

Yogesh Haribhau Kulkarni

Zero to Hero in RAG - 15 Day Training

■ Foundation   ■ NLP & ML   ■ RAG Implementation
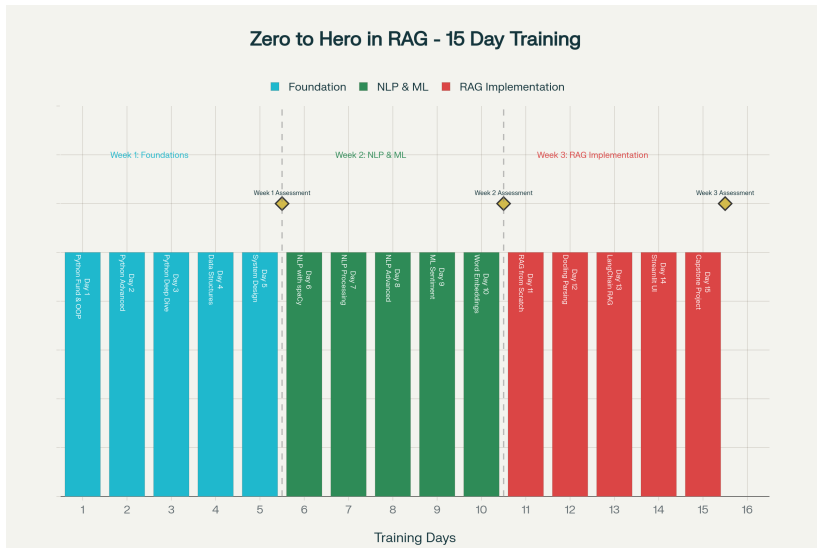
Learning Path: Python $\rightarrow$ DSA $\rightarrow$ System Design $\rightarrow$ NLP $\rightarrow$ RAG $\rightarrow$ Production

Program Overview

- ▶ Duration: 15 working days (3 weeks)
- ▶ Daily: 8 hours (4 learning + 4 coding)
- ▶ Target: Software engineers with basic programming knowledge

What You'll Build

- ▶ 15 progressive projects
- ▶ 1 production capstone
- ▶ Complete end-to-end RAG

Portfolio Highlights:

- ▶ Week 1: 5 Python projects + System design
- ▶ Week 2: NLP pipeline + Embeddings + Vector stores
- ▶ Week 3: Production RAG system with UI

Weekly Structure:

- ▶ Saturday (Pre-week): Online overview + QnA (2 hours)
- ▶ Monday-Friday: Daily learning + coding (8 hours)
- ▶ Final Saturday: Capstone review + QnA (4 hours)

YHK

Program Architecture

Phase 1: Foundations (Days 1-5)
- Python Fundamentals
- File Handling
- Async Programming
- DSA Essentials
- System Design Basics

Phase 2: NLP & ML (Days 6-10)
- Text Processing
- NLP with spaCy
- Vectorization
- Sentiment Analysis
- Word Embeddings

Phase 3: RAG (Days 11-15)
- RAG from Scratch
- Docling Parsing
- LangChain RAG
- Streamlit UI
- Capstone Project

Daily Structure:
- 09:00-13:00: Learning (tutorials + documentation)
- 14:00-18:00: Coding (2 project options - pick 1)

Weekly Milestones:
- Pre-Week Saturday: Online overview session + QnA (2 hrs)
- End-Week Assessment: Project review + readiness check
- Final Saturday (Week 3): All capstone reviews + QnA (4 hrs)

YHK

Day 1: Python Fundamentals & OOP

Morning Session (4 hrs)

- ▶ Variables & data types
- ▶ Control flow (if, loops)
- ▶ Functions & scope
- ▶ Object-oriented programming
- ▶ Classes & objects
- ▶ Inheritance & polymorphism

Afternoon: Pick 1 Project

- ▶ **Library Management:** Build a system with books, members, lending with OOP classes and inheritance
- ▶ **Shopping Cart System:** Create product catalog, cart operations, checkout with discount calculations using OOP

Learning Resources:

- ▶ Tutorial: freeCodeCamp Python OOP (https://www.freecodecamp.org/news/object-oriented-programming-python/)
- ▶ Documentation: Python Official Docs (https://docs.python.org/3/tutorial/classes.html)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working code + README

YHK

Day 2: Python File Handling & Decorators

Morning Session (4 hrs)
- ▸ File I/O operations
- ▸ Reading/writing text files
- ▸ JSON handling
- ▸ CSV processing
- ▸ Exception handling
- ▸ Decorators basics
- ▸ Context managers

Afternoon: Pick 1 Project
- ▸ **CSV Data Processor:** Read CSV files, filter/transform data, export results with error handling and logging decorator
- ▸ **Config File Manager:** Build JSON/YAML config reader/writer with validation, backup system using context managers

Learning Resources:
- ▸ Tutorial: freeCodeCamp File Handling (https://www.freecodecamp.org/news/file-handling-in-python/)
- ▸ Documentation: Python I/O Docs (https://docs.python.org/3/tutorial/inputoutput.html)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working code + README

YHK

Day 3: Python Advanced Deep Dive

Morning Session (4 hrs)

- ▶ List/dict/set comprehensions
- ▶ Generator expressions
- ▶ Async/await basics
- ▶ Asyncio fundamentals
- ▶ Concurrent operations
- ▶ Async file operations

Afternoon: Pick 1 Project

- ▶ **Async Web Scraper:** Fetch multiple URLs concurrently, parse HTML, save results using asyncio and aiohttp
- ▶ **Concurrent File Processor:** Process multiple large files in parallel with async I/O and progress tracking

Learning Resources:

- ▶ Tutorial: freeCodeCamp Async Python (https://www.freecodecamp.org/news/python-async-await-tutorial/)
- ▶ Documentation: Asyncio Docs (https://docs.python.org/3/library/asyncio.html)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working code + README

YHK

Day 4: Data Structures & Algorithms

Morning Session (4 hrs)
- Arrays & lists
- Hash tables/dictionaries
- Stacks & queues
- Binary search
- Sorting algorithms
- Big O complexity

Afternoon: Pick 1 Project
- **LeetCode Problem Set:** Solve 5 medium problems on arrays, hashmaps, and sorting with optimized solutions
- **Custom Data Structures:** Implement stack, queue, and hashmap from scratch with test cases and complexity analysis

Learning Resources:
- Tutorial: freeCodeCamp DSA (https://www.freecodecamp.org/news/data-structures-and-algorithms-in-python/)
- Practice: LeetCode (https://leetcode.com/explore/learn/)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working code + README

YHK

Day 5: System Design Fundamentals

Morning Session (4 hrs)

- ▶ Requirements analysis
- ▶ Scalability patterns
- ▶ Caching strategies
- ▶ Load balancing
- ▶ Database design
- ▶ API design

Afternoon: Pick 1 Project

- ▶ **URL Shortener:** Design and implement TinyURL-like service with hashing, database, and caching layer
- ▶ **API Rate Limiter:** Build token bucket/sliding window rate limiter with Redis-like cache and API endpoints

Learning Resources:

- ▶ Tutorial: freeCodeCamp System Design (https://www.freecodecamp.org/news/systems-design-for-interviews/)
- ▶ Guide: System Design Primer (https://github.com/donnemartin/system-design-primer)

Difficulty: Medium | **Time:** 4 hours | **Deliverable:** Design doc + Working code

YHK

Week 1 Milestone Assessment

Learning Outcomes

- ▶ Python fundamentals mastery
- ▶ File & async operations
- ▶ DSA problem-solving
- ▶ System design thinking

Assessment Activities:

- ▶ Code review of all 5 projects
- ▶ Technical Q&A on concepts
- ▶ System design explanation
- ▶ Readiness verification for Week 2

Assessment Rubric (100 pts):

- ▶ Functionality (50 pts) - All features working correctly
- ▶ Code Quality (20 pts) - Clean, documented, readable
- ▶ Design Principles (20 pts) - Architecture and patterns
- ▶ Documentation (10 pts) - README and comments

Success Criteria

- ▶ 5 projects completed
- ▶ Code quality ¿80%
- ▶ Core concepts understood
- ▶ Ready for NLP phase

YHK

Day 6: Text Processing Fundamentals

Morning Session (4 hrs)

- Text preprocessing
- Tokenization
- Lowercasing & cleaning
- Stop word removal
- Stemming & lemmatization
- Regular expressions

Afternoon: Pick 1 Project

- **Text Preprocessing Pipeline:** Build modular pipeline for cleaning, tokenizing, and normalizing text data with configurable steps
- **Document Cleaner:** Create tool to extract clean text from messy documents with noise removal and format standardization

Learning Resources:

- Tutorial: freeCodeCamp NLP Basics (https://www.freecodecamp.org/news/natural-language-processing-tutorial/)
- Documentation: NLTK Docs (https://www.nltk.org/)

Difficulty: Medium | **Time:** 4 hours | **Deliverable:** Working code + README

YHK

Day 7: NLP with spaCy

Morning Session (4 hrs)

- ▶ spaCy pipeline basics
- ▶ Part-of-speech tagging
- ▶ Named Entity Recognition
- ▶ Dependency parsing
- ▶ Text similarity
- ▶ Pipeline customization

Afternoon: Pick 1 Project

- ▶ **Named Entity Extractor:** Build NER system to extract people, organizations, locations from documents with visualization
- ▶ **Text Similarity Engine:** Create document comparison tool using spaCy embeddings with ranking and scoring

Learning Resources:

- ▶ Tutorial: freeCodeCamp spaCy (https://www.freecodecamp.org/news/spacy-tutorial-nlp/)
- ▶ Documentation: spaCy Official Docs (https://spacy.io/usage)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working code + README

YHK

Day 8: Text Vectorization

Morning Session (4 hrs)

- Bag of Words
- TF-IDF vectorization
- Feature extraction
- Cosine similarity
- Vector operations
- Scikit-learn vectorizers

Afternoon: Pick 1 Project

- **TF-IDF Search Engine:** Build keyword-based search using TF-IDF with document ranking and relevance scoring
- **Document Clusterer:** Create text clustering system using vectorization with visualization of document groups

Learning Resources:

- Tutorial: freeCodeCamp Text Vectorization (https://www.freecodecamp.org/news/text-classification-and-embeddings/)
- Documentation: Scikit-learn Text (https://scikit-learn.org/stable/modules/feature_extraction.html)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working code + README

Day 9: Sentiment Analysis

Morning Session (4 hrs)

- ▶ Sentiment analysis basics
- ▶ Feature engineering
- ▶ ML classifiers (Naive Bayes, SVM)
- ▶ Model training & evaluation
- ▶ Cross-validation
- ▶ Model metrics

Afternoon: Pick 1 Project

- ▶ **Review Sentiment Analyzer:**
  Build classifier for product reviews
  with training pipeline and accuracy
  metrics
- ▶ **Tweet Sentiment Dashboard:**
  Create real-time sentiment analysis
  tool with classification and trend
  visualization

Learning Resources:

- ▶ Tutorial: freeCodeCamp Sentiment Analysis
  (https://www.freecodecamp.org/news/sentiment-analysis-with-python/)
- ▶ Documentation: Scikit-learn ML
  (https://scikit-learn.org/stable/supervised_learning.html)

Difficulty: Medium | **Time:** 4 hours | **Deliverable:** Working code + README

YHK

Day 10: Word Embeddings

Morning Session (4 hrs)

- Word embeddings concepts
- Word2Vec (CBOW, Skip-Gram)
- Pre-trained embeddings
- Embedding similarity
- Sentence embeddings
- Vector databases basics

Afternoon: Pick 1 Project

- **Semantic Search Engine:** Build document search using embeddings with similarity ranking and query expansion
- **Word Analogy Solver:** Create tool for word analogies and relationships using pre-trained Word2Vec models

Learning Resources:

- Tutorial: freeCodeCamp Word Embeddings
  (https://www.freecodecamp.org/news/word-embeddings-for-nlp/)
- Documentation: Gensim Word2Vec
  (https://radimrehurek.com/gensim/models/word2vec.html)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working code + README

YHK

Week 2 Milestone Assessment

Learning Outcomes
- Text processing mastery
- NLP pipeline creation
- ML classification skills
- Embeddings & similarity

Assessment Activities:
- NLP pipeline demonstration
- Sentiment model evaluation
- Embedding system walkthrough
- Component integration check

Assessment Rubric (100 pts):
- Functionality (50 pts) - All NLP components working
- Model Quality (20 pts) - Accuracy and performance
- Code Quality (20 pts) - Clean and documented
- Integration (10 pts) - Components work together

Success Criteria
- 5 NLP projects completed
- ML model ¿75% accuracy
- Vector operations working
- Ready for RAG phase

YHK

Day 11: RAG from Scratch

Morning Session (4 hrs)
- ▶ RAG architecture overview
- ▶ Document chunking strategies
- ▶ Vector stores (FAISS/ChromaDB)
- ▶ Retrieval mechanisms
- ▶ Re-ranking techniques
- ▶ Basic RAG pipeline

Afternoon: Pick 1 Project
- ▶ **Basic RAG System:** Build end-to-end RAG with document loading, chunking, embedding, retrieval, and generation
- ▶ **Multi-Document RAG:** Create RAG system handling multiple documents with metadata filtering and source attribution

Learning Resources:
- ▶ Tutorial: freeCodeCamp RAG (https://www.freecodecamp.org/news/retrieval-augmented-generation-rag-tutorial/)
- ▶ Documentation: LangChain RAG (https://python.langchain.com/docs/tutorials/rag/)

Key Pipeline: Load â†' Chunk â†' Embed â†' Store â†' Retrieve â†' Generate
Difficulty: Medium | **Time:** 4 hours | **Deliverable:** Working RAG + README

YHK

Day 12: Document Parsing with Docling

Morning Session (4 hrs)
- ▶ Docling architecture
- ▶ PDF/DOCX/PPTX parsing
- ▶ Layout analysis
- ▶ Table extraction
- ▶ Image extraction
- ▶ Metadata handling

Afternoon: Pick 1 Project
- ▶ **Multi-Format Parser:** Build document parser supporting PDF/DOCX/PPTX with text, tables, and image extraction
- ▶ **Structured Data Extractor:** Create tool to extract hierarchical content with headings, sections, and metadata preservation

Learning Resources:
- ▶ Tutorial: freeCodeCamp Document Processing (https://www.freecodecamp.org/news/document-parsing-python/)
- ▶ Documentation: Docling Official Docs (https://ds4sd.github.io/docling/)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working parser + README

YHK

Day 13: LangChain RAG Implementation

Morning Session (4 hrs)

- LangChain Essentials course
- Document loaders
- Text splitters
- Embeddings integration
- Vector stores
- Retrieval chains
- LCEL expressions

Afternoon: Pick 1 Project

- **PDF Q&A System:** Build LangChain-based RAG for PDF documents with chat history and source citations
- **Conversational RAG:** Create multi-turn RAG chatbot with memory, context awareness, and follow-up questions

Learning Resources:

- Course: LangChain Academy Essentials (https://academy.langchain.com/)
- Documentation: LangChain Official Docs (https://python.langchain.com/docs/introduction/)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working RAG + README

YHK

Morning Session (4 hrs)

- ▶ Streamlit basics
- ▶ Layout components
- ▶ Interactive widgets
- ▶ File uploads
- ▶ Session state management
- ▶ Chat interfaces

Afternoon: Pick 1 Project

- ▶ **RAG Chatbot Interface:** Create interactive chat UI with file upload, message history, and streaming responses
- ▶ **Document Q&A Dashboard:** Build multi-page app with document upload, processing status, and query interface

Learning Resources:

- ▶ Tutorial: freeCodeCamp Streamlit (https://www.freecodecamp.org/news/build-a-web-app-with-streamlit/)
- ▶ Documentation: Streamlit Official Docs (https://docs.streamlit.io/)

Difficulty: Medium |**Time:** 4 hours |**Deliverable:** Working UI + README

YHK

Morning Session (4 hrs)

- ▶ Architecture planning
- ▶ Component integration
- ▶ Testing strategies
- ▶ Performance optimization
- ▶ Error handling
- ▶ Deployment preparation

Pick 1 Full-Stack Project

- ▶ **Document Intelligence System:** Multi-format parser + RAG + chat UI with source attribution and metadata
- ▶ **Enterprise Knowledge Base:** Multi-document RAG with advanced search, filters, and analytics dashboard

Integration Requirements:

- ▶ Docling: Multi-format document parsing (PDF/DOCX/PPTX)
- ▶ LangChain: Complete RAG orchestration with chains
- ▶ Streamlit: Professional interactive web interface
- ▶ All components seamlessly integrated

Capstone Evaluation (100 pts): Functionality (40) |Code Quality (30) |Architecture (20) |Documentation (10)
Minimum Passing Score: 70/100

YHK

Week 3 Milestone & Final Assessment

Learning Outcomes

- ► Complete RAG systems
- ► Document processing mastery
- ► LLM orchestration
- ► Production-ready apps

Success Criteria

- ► Capstone fully functional
- ► All components integrated
- ► Code well-documented
- ► Production-ready quality

Production Readiness Checklist:

- ► Error handling & validation implemented
- ► Performance optimized (chunking, caching)
- ► Comprehensive documentation
- ► User-friendly interface
- ► Code quality ¿80%

Final Assessment Activities:

- ► Individual capstone project demonstrations (20 mins each)
- ► Architecture & design explanation
- ► Live code walkthrough
- ► Q&A on technical decisions
- ► Future enhancements discussion

YHK

Final Saturday: Capstone Reviews & QnA

Session Structure (4 hours):
- 10:00-12:30: Individual capstone presentations (20 mins each)
- 12:30-13:00: Break
- 13:00-14:00: Open QnA & technical discussions

Presentation Format (20 mins):
- Project demo (8 mins) - Live demonstration
- Architecture overview (5 mins) - System design & components
- Code walkthrough (5 mins) - Key implementation details
- Q&A (2 mins) - Questions from instructors

Evaluation Criteria:
- **Functionality (40 pts):** All features working, error-free
- **Code Quality (30 pts):** Clean, modular, documented
- **Architecture (20 pts):** Sound design, scalable structure
- **Documentation (10 pts):** Clear README, setup instructions

QnA Topics:
- Technical challenges faced & solutions
- RAG optimization techniques

YHK

Congratulations! From RAG to Riches

Technical Skills Acquired

- ▶ Advanced Python programming
- ▶ NLP & text processing
- ▶ RAG architecture & implementation
- ▶ Full-stack AI development
- ▶ Production deployment

Career Pathways

- ▶ AI/ML Engineer
- ▶ RAG Solutions Architect
- ▶ LLM Application Developer
- ▶ AI Product Engineer
- ▶ NLP Engineer

Recommended Certifications:

- ▶ LangChain Academy Certification (academy.langchain.com)
- ▶ Azure AI Fundamentals (AI-900)
- ▶ AWS Machine Learning Specialty

Continuous Learning Resources:

- ▶ LangChain Academy (academy.langchain.com)
- ▶ DeepLearning.AI Short Courses (deeplearning.ai)
- ▶ Hugging Face NLP Course (huggingface.co/learn/nlp-course)
- ▶ LangChain Documentation (python.langchain.com)
- ▶ Explore open-source RAG projects on GitHub

YHK

Weekly Saturday Sessions Schedule

Pre-Week Saturday Sessions (2 hours each):

Saturday Week 0 (Before Day 1):
- ▶ Program overview & expectations
- ▶ Week 1 preview: Python fundamentals to system design
- ▶ Development environment setup
- ▶ QnA session

Saturday Week 1 (Before Day 6):
- ▶ Week 1 recap & learnings
- ▶ Week 2 preview: NLP & ML foundations
- ▶ Introduction to text processing and embeddings
- ▶ QnA session

Saturday Week 2 (Before Day 11):
- ▶ Week 2 recap & portfolio review
- ▶ Week 3 preview: RAG systems & production
- ▶ RAG architecture deep dive
- ▶ QnA session

Final Saturday (After Day 15):

YHK