

ML FOR MECH ENG (CoEP)

Yogesh Haribhau Kulkarni

Outline

1 SESSION 4: PANDAS

About Me

Yogesh Haribhau Kulkarni

Bio:

- ▶ 20+ years in CAD/Engineering software development
- ▶ Got Bachelors, Masters and Doctoral degrees in Mechanical Engineering (specialization: Geometric Modeling Algorithms).
- ▶ Currently doing Coaching in fields such as Data Science, Artificial Intelligence Machine-Deep Learning (ML/DL) and Natural Language Processing (NLP).
- ▶ Feel free to follow me at:
 - ▶ Github (github.com/yogeshhk)
 - ▶ LinkedIn (www.linkedin.com/in/yogeshkulkarni/)
 - ▶ Medium (yogeshharibhaukulkarni.medium.com)
 - ▶ Send email to yogeshkulkarni at yahoo dot com



Office Hours:
Saturdays, 2 to 5pm
(IST); Free-Open to all;
email for appointment.

Introduction to Pandas

Pandas

pandas introduces two new data structures to Python - **Series** and **DataFrame**, both of which are built on top of NumPy. Throughout this session, lines starting with `>>>` (or `...` for a continued line) are what you type at the Python prompt; everything else is the output `pandas` prints back.

```
1 >>> import pandas as pd
  >>> import numpy as np
3 >>> import matplotlib.pyplot as plt
  >>> pd.set_option('display.max_columns', 50)
```

Series

Series is a one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.). The axis labels are collectively referred to as the index. The basic method to create a Series is to call:

Here, data can be many different things:

- ▶ a Python dict
- ▶ an ndarray
- ▶ a scalar value (like 5)

```
>>> s = pd.Series(data, index=index)
```

Creating a Series from a List

- ▶ A Series is a one-dimensional object similar to an array, list, or column in a table.
- ▶ It will assign a labeled index to each item in the Series.
- ▶ By default, each item will receive an index label from 0 to N, where N is the length of the Series minus one.

```
1 >>> # create a Series with an arbitrary list
>>> s = pd.Series([7, 'Heisenberg', 3.14, -1789710578, 'Happy Eating!'])
3 >>> s
```

Series

Output from Previous Slide

```
1 0          7
2 1    Heisenberg
3 2          3.14
4 3    -1789710578
5 4    Happy Eating!
dtype: object
```

Series with a Custom Index

Alternatively, you can specify an index to use when creating the Series.

```
>>> s = pd.Series([7, 'Heisenberg', 3.14, -1789710578,  
2 ... 'Happy Eating!'],  
... index=['A', 'Z', 'C', 'Y', 'E'])  
4 >>> s  
  
6 A              7  
Z      Heisenberg  
8 C              3.14  
Y      -1789710578  
10 E    Happy Eating!  
dtype: object
```

Series

The Series constructor can convert a dictionary as well, using the keys of the dictionary as its index.

```
1 >>> d = {'Chicago': 1000, 'New York': 1300, 'Portland': 900, 'San Francisco':  
          1100, 'Austin': 450, 'Boston': None}  
>>> cities = pd.Series(d)  
3 >>> cities  
Austin          450  
5 Boston         NaN  
Chicago         1000  
7 New York       1300  
Portland         900  
9 San Francisco  1100  
dtype: float64
```

Series

You can use the index to select specific items from the Series

```
>>> cities['Chicago']  
2 1000.0
```

Series

```
>>> cities[['Chicago', 'Portland', 'San Francisco']]
2 Chicago          1000
Portland          900
4 San Francisco    1100
dtype: float64
```

Series

You can use ***boolean indexing*** for selection.

```
1 >>> cities[cities < 1000]
Austin      450
3 Portland   900
dtype: float64
```

Boolean Indexing, Explained

That last one might be a little strange, so let's make it more clear - `cities < 1000` returns a Series of True/False values, which we then pass to our Series `cities`, returning the corresponding True items.

Boolean Indexing: Selecting Values

```
>>> less_than_1000 = cities < 1000
2 >>> print(less_than_1000)
>>> print('\n')
4 >>> print(cities[less_than_1000])
Austin           True
6 Boston         False
Chicago          False
8 New York       False
Portland         True
10 San Francisco False
dtype: bool
12
Austin          450
14 Portland      900
dtype: float64
```

Modifying Series Values by Index

You can also change the values in a Series on the fly.

```
1 >>> # changing based on the index
>>> print('Old value:', cities['Chicago'])
3 >>> cities['Chicago'] = 1400
>>> print('New value:', cities['Chicago'])
5 Old value: 1000.0
New value: 1400.0
```

Modifying Values via Boolean Logic

Changing values using boolean logic

```
>>> print(cities[cities < 1000])
2 >>> print('\n')
>>> cities[cities < 1000] = 750
4 >>> print(cities[cities < 1000])
Austin      450
6 Portland   900
dtype: float64
8
Austin      750
10 Portland  750
dtype: float64
```

Working with Series

What if you aren't sure whether an item is in the Series? You can check using idiomatic Python.

```
1 >>> print('Seattle' in cities)
   >>> print('San Francisco' in cities)
3 False
   True
```

Series Arithmetic with Scalars

Mathematical operations can be done using scalars and functions.

```
>>> # divide city values by 3
2 >>> cities / 3
4 Austin          250.000000
6 Boston          NaN
8 Chicago         466.666667
New York         433.333333
Portland         250.000000
San Francisco    366.666667
dtype: float64
```

Series Arithmetic with NumPy Functions

```
1 >>> # square city values
2 >>> np.square(cities)
3 Austin          562500
4 Boston          NaN
5 Chicago         1960000
6 New York        1690000
7 Portland         562500
8 San Francisco   1210000
9 dtype: float64
```

Adding Two Series Together

You can add two Series together, which returns a union of the two Series with the addition occurring on the shared index values. Values on either Series that did not have a shared index will produce a NULL/NaN (not a number).

```
1 >>> print(cities[['Chicago', 'New York', 'Portland']])
>>> print('\n')
3 >>> print(cities[['Austin', 'New York']])
>>> print('\n')
5 >>> print(cities[['Chicago', 'New York', 'Portland']] + cities[['Austin', 'New
    York']])
```

Adding Two Series: Output

```
1 Chicago      1400
  New York     1300
3 Portland      750
  dtype: float64
5
  Austin        750
7 New York     1300
  dtype: float64
9
  Austin      NaN
11 Chicago    NaN
   New York   2600
13 Portland   NaN
   dtype: float64
```

Working with Series

NULL Checking

- ▶ Notice that because Austin, Chicago, and Portland were not found in both Series, they were returned with NULL/NaN values.
- ▶ NULL checking can be performed with `isnull()` and `notnull()`.

Intuition: Index Alignment

INTUITION

Adding two Series is like matching up name tags at two different parties and adding values only where the *same* name appears at both. Anyone who only showed up to one party (an index value present in just one Series) gets a NaN back: there's nothing on the other side to add them to.

Checking for Non-Null Values

Return a boolean series indicating which values aren't NULL

```
>>> cities.notnull()  
2  
Austin                True  
4 Boston              False  
Chicago              True  
6 New York            True  
Portland              True  
8 San Francisco       True  
dtype: bool
```

Checking for Null Values

Using boolean logic to grab the NULL cities

```
1 >>> print(cities.isnull())
>>> print('\n')
3 >>> print(cities[cities.isnull()])
Austin           False
5 Boston          True
Chicago          False
7 New York        False
Portland          False
9 San Francisco   False
dtype: bool
11
Boston    NaN
13 dtype: float64
```

Quick Check: Series

THINK ABOUT IT

`cities['Chicago']` returns a single value, but `cities[['Chicago', 'Portland']]` returns a Series. What's the difference in the syntax, and why does it matter?

Quick Check: Series

ANSWER

A single string key does scalar/label lookup and returns *one* value. A *list* of keys is fancy indexing and always returns a Series, even with just one element: the double brackets are doing two things at once: the outer `[]` is the indexing operation, the inner `[]` is a literal Python list being passed as the key.

Dataframes in Pandas

Creating Dataframes

Creating a DataFrame by passing a numpy array, with a datetime index and labeled columns:

```
>>> dates = pd.date_range('20130101', periods=6)
2 >>> dates
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
4             '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')
6
>>> df_demo = pd.DataFrame(np.random.randn(6, 4), index=dates,
                           columns=list('ABCD'))
8 >>> df_demo
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401
2013-01-06	-0.673690	0.113648	-1.478427	0.524988

Creating Dataframes

Creating a DataFrame by passing a dict of objects that can be converted to series-like.

```
1 >>> df2 = pd.DataFrame({'A': 1.,
...                       'B': pd.Timestamp('20130102'),
3 ...                       'C': pd.Series(1, index=list(range(4)),
...                                     dtype='float32'),
...                       'D': np.array([3] * 4, dtype='int32'),
5 ...                       'E': pd.Categorical(["test", "train", "test",
...                                     "train"]),
...                       'F': 'foo'})
7 >>> df2
```

	A	B	C	D	E	F
9 0	1.0	2013-01-02	1.0	3	test	foo
1 1	1.0	2013-01-02	1.0	3	train	foo
11 2	1.0	2013-01-02	1.0	3	test	foo
3	1.0	2013-01-02	1.0	3	train	foo

Reading Data from Files

Every DataFrame so far has been typed in by hand. In practice, almost all data starts life in a file: `pd.read_csv()` and `pd.read_excel()` load it straight into a DataFrame, inferring column names and types automatically. `machine_logs.csv` is provided upfront (no need to type it): copy it into your working folder before class.

```
>>> df = pd.read_csv('machine_logs.csv')
2 >>> df.head(3)
```

	machine_id	shift	temperature	vibration	operator_notes
0	M1	Morning	72.5	0.021	ok
6 1	M2	Morning	75.1	0.034	check bearing
2	M1	Evening	78.3	0.041	ok

Quick Inspection

Before doing anything else with a newly loaded DataFrame, check its shape, column types, and non-null counts.

```
1 >>> df.shape
  (6, 5)
3
5 >>> df.dtypes
machine_id      object
shift           object
7 temperature   float64
vibration       float64
9 operator_notes object
dtype: object
```

Quick Inspection: Full Structure

`df.info()` combines shape, dtypes, and non-null counts into one summary.

```
>>> df.info()
2 <class 'pandas.core.frame.DataFrame'>
   RangeIndex: 6 entries, 0 to 5
4 Data columns (total 5 columns):
   #   Column          Non-Null Count  Dtype
6 ---  -
   0   machine_id       6 non-null     object
   8   shift            6 non-null     object
   2   temperature      6 non-null     float64
  10   vibration        6 non-null     float64
   4   operator_notes   6 non-null     object
  12 dtypes: float64(2), object(3)
```

Viewing Data

See the top & bottom rows of the frame

```
>>> df_demo.head()
2
2013-01-01    0.469112 -0.282863 -1.509059 -1.135632
4
2013-01-02    1.212112 -0.173215  0.119209 -1.044236
2013-01-03   -0.861849 -2.104569 -0.494929  1.071804
6
2013-01-04    0.721555 -0.706771 -1.039575  0.271860
2013-01-05   -0.424972  0.567020  0.276232 -1.087401
8

>>> df_demo.tail(3)
10
2013-01-04    0.721555 -0.706771 -1.039575  0.271860
12
2013-01-05   -0.424972  0.567020  0.276232 -1.087401
2013-01-06   -0.673690  0.113648 -1.478427  0.524988
```

Viewing Data

Display the index, columns, and the underlying numpy data

```
1 >>> df_demo.index
   DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
3               '2013-01-05', '2013-01-06'],
   dtype='datetime64[ns]', freq='D')
5
7 >>> df_demo.columns
   Index(['A', 'B', 'C', 'D'], dtype='object')
9
11 >>> df_demo.values
   array([[ 0.4691, -0.2829, -1.5091, -1.1356],
13         [ 1.2121, -0.1732,  0.1192, -1.0442],
14         [-0.8618, -2.1046, -0.4949,  1.0718],
15         [ 0.7216, -0.7068, -1.0396,  0.2719],
16         [-0.425 ,  0.567 ,  0.2762, -1.0874],
17         [-0.6737,  0.1136, -1.4784,  0.525 ]])
```

Viewing Data

Describe shows a quick statistic summary of your data

```
1 >>> df_demo.describe()
3      count      A      B      C      D
5  mean    0.073711 -0.431125 -0.687758 -0.233103
7  std     0.843157  0.922818  0.779887  0.973118
9  min    -0.861849 -2.104569 -1.509059 -1.135632
25%   -0.611510 -0.600794 -1.368714 -1.076610
50%    0.022070 -0.228039 -0.767252 -0.386188
75%    0.658444  0.041933 -0.034326  0.461706
max     1.212112  0.567020  0.276232  1.071804
```

Viewing Data

Sorting by an axis

```
>>> df_demo.sort_index(axis=1, ascending=False)
```

	D	C	B	A
2 2013-01-01	-1.135632	-1.509059	-0.282863	0.469112
4 2013-01-02	-1.044236	0.119209	-0.173215	1.212112
2013-01-03	1.071804	-0.494929	-2.104569	-0.861849
6 2013-01-04	0.271860	-1.039575	-0.706771	0.721555
2013-01-05	-1.087401	0.276232	0.567020	-0.424972
8 2013-01-06	0.524988	-1.478427	0.113648	-0.673690

Viewing Data

Sorting by values

```
>>> df_demo.sort_values(by='B')
```

	A	B	C	D
2 2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
4 2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
6 2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-06	-0.673690	0.113648	-1.478427	0.524988
8 2013-01-05	-0.424972	0.567020	0.276232	-1.087401

Value Counts

`value_counts()` is the fastest way to see how a categorical column is distributed: here, how many readings each machine contributed.

```
>>> df['machine_id'].value_counts()
2  M1    2
   M2    2
4  M3    2
Name: machine_id, dtype: int64
```

Selection of Data

Selecting a single column, which yields a Series, equivalent to `df_demo.A`

```
1 >>> df_demo['A']  
2013-01-01    0.469112  
3 2013-01-02    1.212112  
2013-01-03   -0.861849  
5 2013-01-04    0.721555  
2013-01-05   -0.424972  
7 2013-01-06   -0.673690  
Freq: D, Name: A, dtype: float64
```

Selection of Data

Selecting via `[]`, which slices the rows.

```
>>> df_demo[0:3]
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804


```
>>> df_demo['20130102':'20130104']
```

	A	B	C	D
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860

Selection of Data

Selecting on a multi-axis by label

```
1 >>> df_demo.loc[:, ['A', 'B']]
      A      B
3 2013-01-01  0.469112 -0.282863
  2013-01-02  1.212112 -0.173215
5 2013-01-03 -0.861849 -2.104569
  2013-01-04  0.721555 -0.706771
7 2013-01-05 -0.424972  0.567020
  2013-01-06 -0.673690  0.113648
```

Selection of Data

By integer slices, acting similar to numpy/python

```
>>> df_demo.iloc[3:5, 0:2]
```

	A	B
2	2013-01-04	0.721555 -0.706771
4	2013-01-05	-0.424972 0.567020

Selection of Data

By lists of integer position locations, similar to the numpy/python style

```
>>> df_demo.iloc[[1, 2, 4], [0, 2]]
```

	A	C
2013-01-02	1.212112	0.119209
2013-01-03	-0.861849	-0.494929
2013-01-05	-0.424972	0.276232

Intuition: Label vs. Position

INTUITION

Label-based selection is like finding a book by its title on the shelf: you ask for “Wednesday’s row” or “column A,” and pandas looks up that name.

Position-based selection is like grabbing the third book from the left, regardless of what it’s called. Both answer valid but different questions, so pandas gives you separate ways to ask each one instead of guessing your intent.

Boolean Indexing

Using a single column's values to select data.

```
1 >>> df_demo[df_demo.A > 0]
3      A      B      C      D
3 2013-01-01  0.469112 -0.282863 -1.509059 -1.135632
3 2013-01-02  1.212112 -0.173215  0.119209 -1.044236
5 2013-01-04  0.721555 -0.706771 -1.039575  0.271860
```

Setting Values

Two changes happen here that carry into every slide from this point on: column D is overwritten entirely, and a new column F is added by assigning a Series that only covers *some* of the dates: alignment (the same mechanism from the Series section) fills in the rest with NaN.

```

1 >>> df_demo.loc[:, 'D'] = np.array([5] * len(df_demo))
3 >>> s1 = pd.Series([1, 2, 3, 4, 5, 6], index=pd.date_range('20130102',
    periods=6))
>>> df_demo['F'] = s1
5 >>> df_demo

```

	A	B	C	D	F
2013-01-01	0.469112	-0.282863	-1.509059	5	NaN
2013-01-02	1.212112	-0.173215	0.119209	5	1.0
2013-01-03	-0.861849	-2.104569	-0.494929	5	2.0
2013-01-04	0.721555	-0.706771	-1.039575	5	3.0
2013-01-05	-0.424972	0.567020	0.276232	5	4.0
2013-01-06	-0.673690	0.113648	-1.478427	5	5.0

Missing Data

Missing values show up as NaN. Blanking out an earlier value and reindexing to a smaller set of dates creates some gaps to demonstrate on:

```
>>> df_demo.loc[dates[0], ['A', 'B']] = 0
2 >>> df1 = df_demo.reindex(index=dates[0:4], columns=list(df_demo.columns) + ['E'])
>>> df1.loc[dates[0]:dates[1], 'E'] = 1
4 >>> df1
```

	A	B	C	D	F	E
6 2013-01-01	0.000000	0.000000	-1.509059	5	NaN	1.0
2013-01-02	1.212112	-0.173215	0.119209	5	1.0	1.0
8 2013-01-03	-0.861849	-2.104569	-0.494929	5	2.0	NaN
2013-01-04	0.721555	-0.706771	-1.039575	5	3.0	NaN

Missing Data: Drop

To drop any rows that have missing data:

```
1 >>> df1.dropna(how='any')
```

```
3      A      B      C  D  F  E  
2013-01-02  1.212112 -0.173215  0.119209  5  1.0  1.0
```

Missing Data: Fill

Filling missing data:

```
1 >>> df1.fillna(value=5)
3      A      B      C  D  F  E
3 2013-01-01  0.000000  0.000000 -1.509059  5  5.0  1.0
3 2013-01-02  1.212112 -0.173215  0.119209  5  1.0  1.0
5 2013-01-03 -0.861849 -2.104569 -0.494929  5  2.0  5.0
2013-01-04  0.721555 -0.706771 -1.039575  5  3.0  5.0
```

Stats

Operations in general exclude missing data. Performing a descriptive statistic:

```
>>> df_demo.mean()
2 A    -0.004474
  B    -0.383981
4 C    -0.687758
  D     5.000000
6 F     3.000000
dtype: float64
```

Stats: The Other Axis

Same operation on the other axis:

```
1 >>> df_demo.mean(1)
2013-01-01    0.872735
3 2013-01-02    1.431621
4 2013-01-03    0.707731
5 2013-01-04    1.395042
6 2013-01-05    1.883656
7 2013-01-06    1.592306
Freq: D, dtype: float64
```

Apply

Applying functions to the data

```
>>> df_demo.apply(np.cumsum)
2          A          B          C  D    F
2013-01-01  0.000000  0.000000 -1.509059  5  NaN
4 2013-01-02  1.212112 -0.173215 -1.389850 10  1.0
2013-01-03  0.350263 -2.277784 -1.884779 15  3.0
6 2013-01-04  1.071818 -2.984555 -2.924354 20  6.0
2013-01-05  0.646846 -2.417535 -2.648122 25 10.0
8 2013-01-06 -0.026844 -2.303886 -4.126549 30 15.0

10 >>> df_demo.apply(lambda x: x.max() - x.min())
12 A    2.073961
13 B    2.671590
14 C    1.785291
15 D    0.000000
16 F    4.000000
dtype: float64
```

Renaming & Dropping Columns

Cleaning up column names and removing ones you don't need is usually the first step of any real analysis.

```
>>> df = df.rename(columns={'vibration': 'vibration_mm_s'})
2 >>> df = df.drop(columns=['operator_notes'])
>>> df.columns
4 Index(['machine_id', 'shift', 'temperature', 'vibration_mm_s'], dtype='object')
```

Creating New Columns

New columns are computed from existing ones with plain vectorized arithmetic and comparisons, no loops needed.

```
>>> df['vibration_in'] = df['vibration_mm_s'] / 25.4
2 >>> df['high_vibration'] = df['vibration_mm_s'] > 0.025
>>> df[['machine_id', 'vibration_mm_s', 'high_vibration']]
4
6      machine_id  vibration_mm_s  high_vibration
8 0           M1          0.021          False
9 1           M2          0.034           True
10 2           M1          0.041           True
11 3           M2          0.028           True
12 4           M3          0.019          False
13 5           M3          0.052           True
```

Quick Check: DataFrames

THINK ABOUT IT

`df['A']` returns a Series (a column), but `df[0:2]` returns a DataFrame (the first two rows). Same `[]` operator, same object: why does pandas treat a string key and a slice so differently?

Quick Check: DataFrames

ANSWER

The plain `[]` operator on a DataFrame is intentionally overloaded: a single label looks up a *column* (matching the `df.A` shortcut), while a slice looks up *rows*, mirroring plain Python/NumPy slicing. Pandas infers your intent from the *type* of the key rather than a fixed row/column rule, which is exactly why `.loc/.iloc` exist: they remove the ambiguity by always meaning “by label” or “by position,” regardless of what kind of key you pass.

GroupBy: Split-Apply-Combine

`groupby()` splits rows into groups by a key column. Nothing is computed until you chain on an aggregate.

INTUITION

Think of `groupby()` as sorting mail into pigeonholes by recipient. Nothing gets read or totalled yet: that's just the sorting step. The real work happens when a clerk (an aggregate like `.mean()` or `.sum()`) processes each pigeonhole separately and the results are stacked back into a single tray.

```
>>> df.groupby('machine_id')  
2 <pandas.core.groupby.generic.DataFrameGroupBy object at 0x...>
```

GroupBy with Aggregation

Chain `.mean()`, `.sum()`, or `.agg()` to actually run the per-group computation.

```
>>> df.groupby('machine_id')['temperature'].mean()
2 M1    75.40
  M2    74.55
4 M3    75.50
  Name: temperature, dtype: float64
6
>>> df.groupby('machine_id').agg({'temperature': 'mean', 'vibration_mm_s':
  'max'})
8
   temperature  vibration_mm_s
machine_id
10 M1          75.40          0.041
   M2          74.55          0.034
12 M3          75.50          0.052
```

Quick Check: GroupBy

THINK ABOUT IT

`df.groupby('machine_id')` on its own prints only an object description, not a table of numbers. Why doesn't pandas compute anything at that point?

Quick Check: GroupBy

ANSWER

`groupby()` only performs the *split*: it doesn't yet know what you want to *apply* to each group. Waiting until you chain on `.mean()`, `.sum()`, or `.agg()` means the same grouped object can be reused for several different aggregations without repeating the split step, and pandas never computes a result you never asked for.

Concatenating DataFrames

`pd.concat()` stacks DataFrames with the same columns on top of each other, useful when the same kind of data arrives in separate batches (e.g., one file per shift).

```
>>> morning = df[df['shift'] == 'Morning'][['machine_id', 'temperature']]
2 >>> evening = df[df['shift'] == 'Evening'][['machine_id', 'temperature']]
>>> pd.concat([morning, evening])
4   machine_id  temperature
0          M1          72.5
6   1          M2          75.1
4   4          M3          69.8
8   2          M1          78.3
   3          M2          74.0
10  5          M3          81.2
```

Merging DataFrames

`pd.merge()` joins two DataFrames on a shared key column: like a VLOOKUP, or a SQL JOIN.

```
>>> machine_info = pd.DataFrame({
...     'machine_id': ['M1', 'M2', 'M3'],
...     'machine_type': ['Lathe', 'Mill', 'Drill']
... })
>>> pd.merge(df, machine_info, on='machine_id')[
...     ['machine_id', 'machine_type', 'temperature']].head(3)
```

	machine_id	machine_type	temperature
0	M1	Lathe	72.5
1	M2	Mill	75.1
2	M1	Lathe	78.3

Plotting with Pandas

`.plot()` is a thin wrapper around the `matplotlib` imported all the way back on the very first Pandas slide, no separate charting library needed for a first look at the data.

```
>>> avg_temp = df.groupby('machine_id')['temperature'].mean()  
2 >>> avg_temp.plot(kind='bar')  
>>> plt.ylabel('Avg Temperature (deg C)')  
4 >>> plt.show()
```

Putting It Together: Machine Health Check

The same five-step pattern (read, inspect, clean, transform, summarize) shows up in almost every dataset you'll touch in this course.

```
>>> df = pd.read_csv('machine_logs.csv')
2 >>> df = df.rename(columns={'vibration': 'vibration_mm_s'})
>>> df = df.drop(columns=['operator_notes'])
4 >>> df['high_vibration'] = df['vibration_mm_s'] > 0.025

6 >>> summary = df.groupby('machine_id').agg(
...     avg_temp=('temperature', 'mean'),
8 ...     flagged_readings=('high_vibration', 'sum')
... )
10 >>> summary
```

	avg_temp	flagged_readings
machine_id		
M1	75.40	1
M2	74.55	2
M3	75.50	1

Putting It Together: Visualizing the Summary

One line turns the summary table into a chart flagging which machine needs a closer look.

```
1 >>> summary['avg_temp'].plot(kind='bar', color='steelblue')  
>>> plt.title('Average Temperature by Machine')  
3 >>> plt.ylabel('deg C')  
>>> plt.show()
```

Try It Yourself

Using the same `machine_logs.csv` you already have, no new data or setup needed:

- ▶ Which machine has the highest average vibration?
- ▶ Filter the DataFrame to only the Evening shift rows.
- ▶ Add a column flagging rows where `temperature > 76`.

Thanks ...

- ▶ Office Hours: Saturdays, 3 to 5 pm (IST); Free-Open to all; email for appointment to yogeshkulkarni at yahoo dot com
- ▶ Call + 9 1 9 8 9 0 2 5 1 4 0 6



(<https://www.linkedin.com/in/yogeshkulkarni/>)



(<https://medium.com/@yogeshharibhaukulkarni>)



(<https://www.github.com/yogeshhk/>)