

Machine Learning for Mechanical Engineers at CoEP Session 8: Feature Selection

Session 8: FeatSel

Feature Extraction and Selection

From Raw Data to Useful Features

- Before training any model: which numbers do you feed it, and how many?
- **Feature Extraction**: raw data → a manageable set of numeric features.
- **Feature Selection**: given candidate features, which ones actually help?

Feature Extraction: Statistical Features

- Real signals often arrive as a stream: a sensor reading, a week of sales, a semester of test scores.
- A model needs one row per example, not 1000 raw readings – so summarize the stream instead.
- Simplest, most common summaries: mean, standard deviation, minimum, maximum, range.

Intuition

Describing a month of temperatures in one sentence? You wouldn't list 30 numbers – “averaged 22°C, swung between 18 and 27”. That sentence *is* mean, min, max: statistical features.

Feature Extraction: A Worked Example

A sensor logs one reading a minute, for 10 minutes:

```
import numpy as np

readings = np.array([21.1, 21.3, 21.0, 21.6, 22.0,
                    22.4, 22.1, 21.8, 21.5, 21.2])

features = {
    'mean': readings.mean(),
    'std': readings.std(),
    'min': readings.min(),
    'max': readings.max(),
    'range': readings.max() - readings.min(),
}

print(features)

# Output:
# {'mean': 21.6, 'std': 0.44, 'min': 21.0, 'max': 22.4, 'range': 1.4}
```

10 numbers became 5 features. Per sensor, per time window – several sensors, dozens of candidates.

Feature Selection: Which Features Actually Help?

Feature Selection: Ranking (Filter Methods)

- Score each feature *on its own* against the target, rank, keep the top few.
- “Score”: a statistical test, e.g. an F-test (ANOVA) for a numeric feature vs. a class label.
- Fast – one pass, no model training. Called a **filter**: features filtered before any model sees them.
- Blind spot: scores features in isolation, can miss a feature that only helps *in combination*.

Ranking: Loading the Data

Pima Indians Diabetes dataset: 768 patients, 8 measurements, one outcome.

```
import pandas as pd

url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indians-diabetes.data.csv"
names = ['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class']
df = pd.read_csv(url, names=names)
df.head(3)


# Output:
#   preg  plas  pres  skin  test  mass  pedi  age  class
# 0    6   148    72   35    0   33.6  0.627  50     1
# 1    1    85    66   29    0   26.6  0.351  31     0
# 2    8   183    64    0    0   23.3  0.672  32     1
```

class: 1 = diabetic, 0 = not. Which of the 8 columns actually predicts it?

Ranking: Correlation with the Outcome

Simplest ranking: correlation of each column with class.

```
corr = df.corr()['class'].drop('class').sort_values(
    ascending=False)
print(corr.round(2))
```



Ranking: F-score (a Sharper Test)

Correlation is one lens; an F-test (ANOVA) is another, built for exactly this:

```
from sklearn.feature_selection import SelectKBest, f_classif

X, y = df.values[:, 0:8], df.values[:, 8]
skb = SelectKBest(score_func=f_classif, k=4).fit(X, y)
ranked = sorted(zip(names[0:8], skb.scores_), key=lambda t: -t[1])

# F-score, highest to lowest:
# plas 213 mass 72 age 46 preg 40
# pedi 24 test 13 skin 4 pres 3
```

Same story as the correlation plot: plas dominates, skin/pres barely register.

Feature Selection: Wrapper Methods

- Filter scores each feature alone. **Wrapper**: train/cross-validate the model on different feature *subsets*, let that score decide.
- Slower – many models trained – but catches combinations a filter would miss.
- Which subsets to try? Four classic strategies:
 - Exhaustive Search
 - Greedy Forward Selection
 - Greedy Backward Elimination
 - Best-First Search

Wrapper: Exhaustive Search

- Try *every* non-empty subset, keep the best cross-validated score.
- n features $\rightarrow 2^n - 1$ subsets. Pima's 8 $\rightarrow 255$ – small enough to actually run.

```
import itertools
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score, KFold

model = LogisticRegression(max_iter=1000)
kfold = KFold(n_splits=5, shuffle=True, random_state=7)
best_score, best_subset = -1, None

for r in range(1, 9):
    for combo in itertools.combinations(range(8), r):
        s = cross_val_score(model, X[:, combo], y, cv=
            kfold).mean()
        if s > best_score:
            best_score, best_subset = s, combo

# Output (255 subsets evaluated):
# Best: all features except 'skin', accuracy=0.775
# All 8 features: accuracy=0.775
```

Proves **skin** adds nothing: dropping it costs zero accuracy.

Wrapper: Why Exhaustive Search Doesn't Scale

- 8 features $\rightarrow 255$ subsets: a few seconds.
- 3 sensors $\times$ 5 statistical features each = 15-20 candidates already.
- 20 features $\rightarrow 2^{20} - 1 \approx 1$ million subsets – roughly $4,000\times$ more work.
- Realistic only for a small, already-narrowed feature set.

Wrapper: Greedy Forward Selection

- Start empty. Each round: add whichever remaining feature helps most.
- Stop when nothing left improves the score.

Same model and cross validation, on Pima:

Step	Feature added	Accuracy so far
1	plas	0.750
2	mass	0.767
3	pedi	0.773
4	pres	0.772

3 features reach 0.773 – within 0.002 of exhaustive search's best, at a fraction of the cost.

Wrapper: Greedy Backward Elimination

- Start with *all* features. Each round: remove whichever hurts least (or helps).
- Stop once every remaining removal costs accuracy.

On Pima:

Step	Feature removed	Accuracy after
start	(all 8)	0.775
1	skin	0.775
2	age	0.772
3	preg	0.775

First move – drop **skin** – already matches what exhaustive search proved was right.

Wrapper: Best-First Search

- Greedy commits forever: once added or removed, never reconsidered.
- Best-first keeps a pool of promising subsets, not just one – can expand any of them next.
- Can back out of an early choice greedy would be stuck with.
- Costs more than greedy, far less than exhaustive.

Intuition

Greedy: always take the steepest visible path uphill. Best-first: keep a mental map of your 3 best viewpoints, and choose which to climb from next.

Filter vs. Wrapper: Summary

Method	Looks at	Cost
Ranking (Filter)	One feature at a time	Cheapest
Exhaustive	Every possible subset	Guaranteed best, but explodes
Greedy Forward	Builds up, one add at a time	Cheap, no backtracking
Greedy Backward	Trims down, one remove at a time	Cheap, no backtracking
Best-First	A pool of candidate subsets	Middle ground

Intuition

No method is “correct.” Ranking: fast first pass. Wrapper: costs more, sees combinations. Pick by budget.

Quick Check: Feature Selection

Think About It

The filter ranking on Pima placed **preg** (number of pregnancies) 4th out of 8 individually. Yet Greedy Backward Elimination's best 5-feature subset dropped **preg** entirely without hurting accuracy. Are these two results in conflict? Why or why not?

Quick Check: Feature Selection

Answer

No conflict – they answer different questions. The filter score measures how strongly **preg** relates to the outcome *by itself*, ignoring every other feature; on that measure it ranks 4th. Backward elimination asks a different question: given everything else already in the model, does **preg** add anything *on top of it*? If the information **preg** carries is already captured by other features (e.g. **age** correlates with number of pregnancies), removing it costs nothing even though it looked useful alone. This is exactly the wrapper methods' advantage over filters: they can see redundancy that isolated ranking cannot.

Recap: Feature Extraction and Selection

- Extraction:** raw data $\rightarrow$ manageable features (mean, std, min, max, range).
- Selection:** which candidate features earn a place in the model.
- Filter (Ranking):** fast, scores features independently.
- Wrapper** (Exhaustive, Greedy Forward/Backward, Best-First): slower, sees combinations filters miss.
- Narrows the search – you still decide what ships.

Hyperparameter Tuning

Hyperparameter Tuning

- Parameters:** learned from data during **fit()** – e.g. regression coefficients, a tree's split points.
- Hyperparameters:** set *before* training, control how learning happens – e.g. k in KNN, **max_depth** in a Decision Tree, C in an SVM.
- Picking by hand (“try $k = 3$, try $k = 5$, ...”) doesn't scale past one hyperparameter.
- Fix: try many combinations, score each with cross validation, keep the best.

Intuition

Parameters: what the model learns. Hyperparameters: the knobs you set before it starts learning – oven temperature and bake time, chosen before the cake goes in.

Hyperparameter Tuning: Grid Search

Try every combination on a fixed grid, keep the best cross-validated score.

```
import pandas as pd
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier

url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indians-diabetes.data.csv"
names = ['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class']
df = pd.read_csv(url, names=names)
X, y = df.values[:, 0:8], df.values[:, 8]

model = KNeighborsClassifier()
param_grid = {'n_neighbors': [3, 5, 7, 9, 11, 15, 21]}

grid = GridSearchCV(model, param_grid, cv=5, scoring='accuracy')
grid.fit(X, y)

# Output: Best k 11, accuracy 0.749
```

Hyperparameter Tuning: Grid Search, Cost

- 7 candidate values of $k \times 5$ folds = 35 model fits.
- All run automatically – no manual loop.
- One hyperparameter is cheap. Cost multiplies fast with more: 3 hyperparameters, 10 values each, = 1000 fits.

Hyperparameter Tuning: Randomized Search

- Grid Search checks every combination – fine for one hyperparameter, explodes with more.
- RandomizedSearchCV: sample a fixed number of random combinations instead.
- Trades a guarantee of the single best combination for a large cut in compute.

```
from sklearn.model_selection import RandomizedSearchCV

param_dist = {'n_neighbors': list(range(1, 30))}
random_search = RandomizedSearchCV(model, param_dist,
                                   n_iter=5,
                                   cv=5, scoring='accuracy', random_state=42)
random_search.fit(X, y)

# Output: tried 5 of 29 possible values -> Best k 13, accuracy 0.755
```

Hyperparameter Tuning: Randomized Search, Result

- Just 5 candidates out of 29 possible values of k .
- Matched Grid Search's result ($k = 11$ vs. $k = 13$, 0.749 vs. 0.755) at a fraction of the cost.
- In practice: often close to Grid Search's answer, rarely worse in any way that matters.

Quick Check: Hyperparameter Tuning

Think About It

Why is it important to use cross-validated scores, not a single train/test split, when comparing hyperparameter values during Grid or Randomized Search?

Quick Check: Hyperparameter Tuning

Answer

A single split can hand a particular hyperparameter value a “lucky” or “unlucky” test set by chance. Averaging over k folds gives a far more reliable estimate of how each candidate value actually generalizes, so the value selected as “best” is best on average, not best on one arbitrary split.

Hands-On: End-to-End Mini Pipeline

The Dataset

- Breast Cancer Wisconsin dataset, built into scikit-learn – no download needed.
- 569 patients, 30 measurements each, one target: malignant or benign.
- 30 features – too many to exhaustively search, a real test for ranking and wrapper methods.
- Goal: pandas to load and inspect, sklearn to preprocess/select/train/evaluate – a revision lap before the algorithm-by-algorithm sessions start.

Step 1: Load and Look

```
import pandas as pd
from sklearn.datasets import load_breast_cancer

data = load_breast_cancer()
df = pd.DataFrame(data.data, columns=data.feature_names)
df['target'] = data.target
print(df.shape)
df.head(3)

# Output: (569, 31)
```

```
#   mean radius  mean texture  ...  worst fractal
#   dimension  target
# 0      17.99         0      10.38 ...
# 1      20.57         0      17.77 ...
# 0      0.0890         0
```

Step 2: Missing Values and Class Balance

```
print(df.isna().sum().sum(), "missing values")
print(df['target'].value_counts())

# Output:
# 0 missing values
# target
# 1    357   (benign)
# 0    212   (malignant)
```

Step 3: Train/Test Split

Stratify on the target so both splits keep the same 63/37 class balance.

```
from sklearn.model_selection import train_test_split

X = df.drop(columns='target').values
y = df['target'].values
feature_names = df.drop(columns='target').columns.tolist()

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=42, stratify=y)

# Output: X_train (426, 30)   X_test (143, 30)
```

Step 4: Scale the Features

Fit the scaler on training data only – the same no-leakage rule from Data Preparation.

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler().fit(X_train)
X_train_s = scaler.transform(X_train)
X_test_s = scaler.transform(X_test)
```

Step 5: Baseline, All 30 Features

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

model_all = LogisticRegression(max_iter=5000)
model_all.fit(X_train_s, y_train)
pred_all = model_all.predict(X_test_s)
print("Accuracy:", accuracy_score(y_test, pred_all))

# Output: Accuracy: 0.986
```

Step 6: Rank Features (Filter)

Same SelectKBest approach as before, now on 30 candidates instead of 8.

```
from sklearn.feature_selection import SelectKBest, f_classif

skb = SelectKBest(score_func=f_classif, k=10).fit(
    X_train_s, y_train)
top10 = [feature_names[i] for i in skb.get_support(indices=True)]
print(top10)

# Output: ['mean radius', 'mean perimeter', 'mean area',
#         'mean concavity', 'mean concave points', 'worst radius',
#         'worst perimeter', 'worst area', 'worst concavity',
#         'worst concave points']
```

Step 7: Retrain on the Top 10

```
idx = skb.get_support(indices=True)
model_10 = LogisticRegression(max_iter=5000)
model_10.fit(X_train_s[:, idx], y_train)
pred_10 = model_10.predict(X_test_s[:, idx])
print("Accuracy:", accuracy_score(y_test, pred_10))

# Output: Accuracy: 0.951
```

Step 8: Filter Result – Not a Free Lunch

- 30 features: 0.986 test accuracy.
- Filter top 10: 0.951 test accuracy.
- Unlike Pima's skin column, these 20 dropped features weren't dead weight – accuracy fell.
- Real trade-off: 20 fewer features for ~3.5 points of accuracy. Worth it or not depends on the deployment.

Step 9: Try a Wrapper Instead

Greedy Forward Selection, cross-validated, on all 30 features:

```
from sklearn.model_selection import cross_val_score, KFold

Xs = StandardScaler().fit_transform(X)
kfold = KFold(n_splits=5, shuffle=True, random_state=7)
selected, remaining = [], list(range(30))

for step in range(6):
    scores = [(f, cross_val_score(LogisticRegression(
        max_iter=5000),
```

```
    Xs[:, selected + [f]], y, cv=kfold).mean())
    for f in remaining]
    best_f, best_s = max(scores, key=lambda t: t[1])
    selected.append(best_f); remaining.remove(best_f)
```

```
# Output: 6 rounds -> worst area, worst smoothness, worst texture,
#         worst compactness, mean compactness, mean fractal dimension
```

Step 10: Filter vs. Wrapper vs. Everything

Feature set	# Features	5-fold CV Accuracy
All features	30	0.975
Filter (top by F-score)	10	0.951
Wrapper (greedy forward)	6	0.979

Intuition

6 wrapper-chosen features beat all 30 – fewer, better-combined features can outperform throwing everything at the model. The filter's top 10, chosen one at a time, missed that combination entirely.

Step 11: Evaluate the Final Model

```
from sklearn.metrics import confusion_matrix, classification_report

wrapper_cols = ['worst area', 'worst smoothness', 'worst texture',
               'worst compactness', 'mean compactness', 'mean fractal dimension']
idx6 = [feature_names.index(f) for f in wrapper_cols]

model_6 = LogisticRegression(max_iter=5000)
model_6.fit(X_train_s[:, idx6], y_train)
pred_6 = model_6.predict(X_test_s[:, idx6])
print(confusion_matrix(y_test, pred_6))

# Output: [[50  3]
#         [ 4 86]]
```

Step 11: Classification Report

```
print(classification_report(y_test, pred_6,
                           target_names=['malignant', 'benign']))

# Output:
#               precision    recall  f1-score   support
#  malignant       0.93      0.94      0.93         53
#    benign       0.97      0.96      0.96         90
```

6 features, 95% test accuracy, on par with the 10-feature filter model – and better on cross validation.

Quick Check: The Pipeline

Think About It

Step 6 fits SelectKBest on `X_train_s` (the scaled training set), not the full scaled dataset. Step 9's greedy forward selection, however, scores each candidate using cross validation on *all* of `X` (via `Xs`, not `X_train_s`). Is Step 9 leaking test information in a way Step 6 avoids?

Quick Check: The Pipeline

Answer

No – both are safe, but for different reasons. Step 6 never touches `X_test_s` at all. Step 9 does use every row of `X`, including what would later become test rows, but cross validation itself holds out a fold at a time during scoring – so no fold's held-out data influences its own score. The real risk it does *not* avoid: the final held-out test set (`X_test_s`) from Step 3 was never part of Step 9's cross validation either, so the eventual Step 11 evaluation is still a fair, unseen check on the wrapper-selected features.

Try It Yourself

- Change `k=10` to `k=5` in Step 6 – does accuracy hold up with even fewer features?
- Run Greedy Backward Elimination instead of Forward on all 30 features – same 6 features at the end?
- Swap `LogisticRegression` for a different classifier – does the winning feature set change?
- Plot the correlation matrix of the 6 wrapper-chosen features – are any of them redundant with each other?

Recap: The Pipeline Pattern

- Load → Inspect → Split → Scale → Select → Train → Evaluate → Compare.
- Every step reused today: pandas for loading/inspecting, sklearn for the rest.
- This shape doesn't change once you swap in a real ML algorithm – only Step 5/7's model does.
- Next: the algorithms themselves, one at a time, starting with Linear Regression.