ML for Mech Eng (CoEP)

Yogesh Haribhau Kulkarni

YHK

Outline

About Me

YHK

Yogesh Haribhau Kulkarni

Bio:

- 20+ years in CAD/Engineering software development
- Got Bachelors, Masters and Doctoral degrees in Mechanical Engineering (specialization: Geometric Modeling Algorithms).
- Currently doing Coaching in fields such as Data Science, Artificial Intelligence Machine-Deep Learning (ML/DL) and Natural Language Processing (NLP).
- Feel free to follow me at:
  - Github (github.com/yogeshhk)
  - LinkedIn (www.linkedin.com/in/yogeshkulkarni/)
  - Medium (yogeshharibhaukulkarni.medium.com)
  - Send email to yogeshkulkarni at yahoo dot com

Office Hours: Saturdays, 2 to 5pm (IST); Free-Open to all; email for appointment.

YHK

Applications to Mechanical Engineering

YHK

The Announcement

"Revitalizing manufacturing through AI"

- AI is already transforming the IT industry.
- It is now time to for an AI-powered society.
- We will initially focus on the manufacturing industry.

(Reference: "Medium" - Andrew Ng, Dec 14,2017)

YHK

Why Manufacturing?

- Manufacturing touches every part of society.
- Need to help not just bits and pixels but physical world.

YHK

Challenges facing manufacturing

- Variable quality and yield,
- Inflexible production line design
- Inability to manage capacity
- and rising production costs.

ML to help address them.

YHK

Challenges facing manufacturing

More help in

- Shorten design cycles,
- Remove supply-chain bottlenecks,
- Reduce materials and energy waste,
- and improve production yields.

YHK

Why now?

"Industry today is experiencing a never seen increase in available data"

Sources:

- Environmental data
- Sensor data from production line
- Machine tool parameters
- Logistics, transaction data
- Customer demand data.

(Reference: What is smart manufacturing? Time Magazine - Chand & Davis, 2010)

YHK

How ML can help?

- Ability to handle high-dimensional problems
- SVM handling high dimensionality (> 1000) very well.
- Concerns: Over-fitting

(Ref: "Machine learning in manufacturing: advantages, challenges, and applications"- Thorsten Wuest, et al)

YHK

How ML can help?

- Ability to reduce complexity
- Present transparent and concrete advice for practitioners
- E.g. monitor XX and parameter YY at checkpoint ZZ
- ML finds important features, thresholds

(Ref: "Machine learning in manufacturing: advantages, challenges, and applications"- Thorsten Wuest, et al)

YHK

How ML can help?

- Ability to adapt to changing environment
- New data, new training, new model.

(Ref: "Machine learning in manufacturing: advantages, challenges, and applications"- Thorsten Wuest, et al)

YHK

New wave: IoT

"Internet of Things"

- IOT is connecting things to the Internet
- Not the same as connecting things to the cellular network!
- IOT personal (Health apps)
- IOT Machine (Machine 2 Machine M2M)
- IoT Everywhere (5G, wifi)

YHK

New wave: IoT

IoT Data

- Sensors everywhere.
- By 2020,50 billion connected devices.
- Wearables, cars, machines, cities.

YHK

New wave: IoT

IoT and Machine Learning

- Basic idea of machine learning is to build a mathematical model based on training data(learning stage) , predict results for new data(prediction stage) and tweak the model based on new conditions
- IoT creates a lot of contextual data
- Use the scale of IoT data to gain new insights

YHK

New wave: IoT

IoT ML Applications

- Consumer : bio sensors(real time tracking)
- Enterprise : Complex machinery (preventative maintenance), asset efficiency.
- Public infrastructure: Smart Cities, Dynamically adjust traffic lights

YHK

Sample Applications: Tool Wear Estimation

Hidden Markov Model

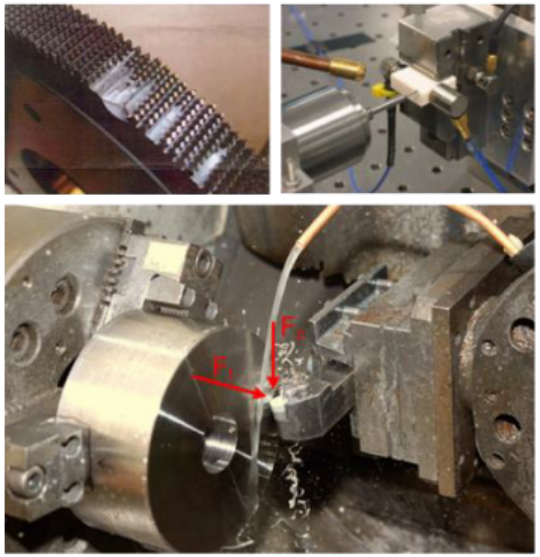

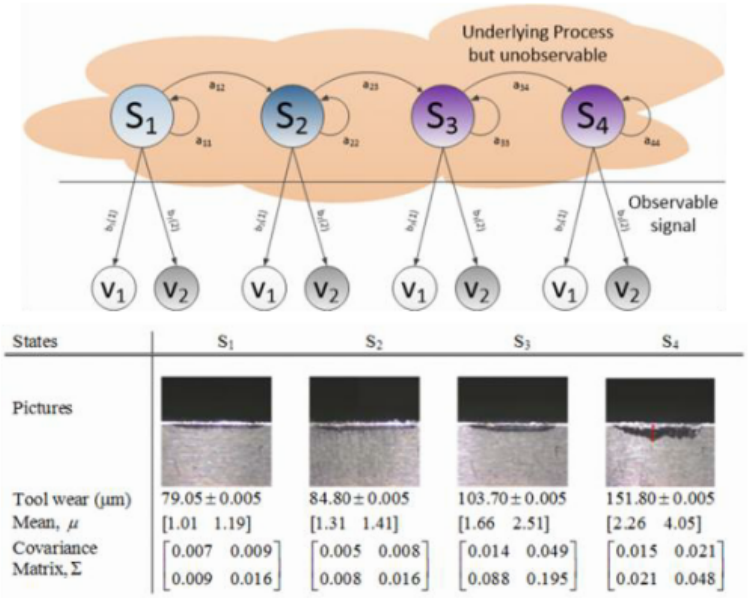

States	S_1	S_2	S_3	S_4
Pictures				
Tool wear (μm)	79.05 ± 0.005	84.80 ± 0.005	103.70 ± 0.005	151.80 ± 0.005
Mean, μ	[1.01 1.19]	[1.31 1.41]	[1.66 2.51]	[2.26 4.05]
Covariance Matrix, Σ	$\begin{bmatrix} 0.007 & 0.009 \\ 0.009 & 0.016 \end{bmatrix}$	$\begin{bmatrix} 0.005 & 0.008 \\ 0.008 & 0.016 \end{bmatrix}$	$\begin{bmatrix} 0.014 & 0.049 \\ 0.088 & 0.195 \end{bmatrix}$	$\begin{bmatrix} 0.015 & 0.021 \\ 0.021 & 0.048 \end{bmatrix}$

S. Lee, L. Li*, and J. Ni, 2010, "Online Degradation Assessment and Adaptive Fault Detection Using Modified Hidden Markov Model," ASME Journal of Manufacturing Science and Engineering, 132(2), pp. 021010-11. [PDF]

YHK

Ideas for new applications

Project ideas?

- Kaggle: https://www.kaggle.com/datasets
- UCI: http://archive.ics.uci.edu/ml/index.php
- Other: http://deeplearning.net/datasets/

YHK

Linear Regression — House Price Prediction

(Ref: scikit-learn User Guide — Generalized Linear Models)

YHK

Problem Info

- Dataset: California Housing (`sklearn.datasets.fetch_california_housing`, built-in)
- Features: median income, house age, avg rooms, avg bedrooms, population, latitude, longitude
- Target: Median house value (in 100,000s USD)
- Goal: Predict house prices using Linear Regression; compare Ridge and Lasso regularization

YHK

Import Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
```

YHK

Load and Explore Data

```
data = fetch_california_housing()
df = pd.DataFrame(data.data, columns=data.feature_names)
df['MedHouseVal'] = data.target

print(df.shape)       # (20640, 9)
print(df.describe())
print(df.head())
```

YHK

Train Linear Regression

```
X_train, X_test, y_train, y_test = train_test_split(
    data.data, data.target, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test  = scaler.transform(X_test)

lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)

print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
print("R2  :", r2_score(y_test, y_pred))
```

YHK

Compare Regularization

```
for name, model in [("Linear", LinearRegression()),
                    ("Ridge",  Ridge(alpha=1.0)),
                    ("Lasso",  Lasso(alpha=0.1))]:
    model.fit(X_train, y_train)
    pred = model.predict(X_test)
    rmse = np.sqrt(mean_squared_error(y_test, pred))
    r2   = r2_score(y_test, pred)
    print(f"{name:8s}  RMSE={rmse:.4f}  R2={r2:.4f}")
```

YHK

Predicted vs Actual

Points close to the diagonal indicate accurate predictions.

```
plt.figure(figsize=(6, 5))
plt.scatter(y_test, y_pred, alpha=0.3, s=10)
plt.plot([y_test.min(), y_test.max()],
         [y_test.min(), y_test.max()], 'r--', lw=2)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Linear Regression: Predicted vs Actual")
plt.tight_layout()
plt.show()
```

YHK

SVM — Handwritten Digit Recognition

(Ref: scikit-learn SVM example — Digit Classification)

YHK

Problem Info

- Dataset: Digits (`sklearn.datasets.load_digits`, built-in)
- 1,797 images of handwritten digits (0–9), each 8×8 pixels $= 64$ features
- Goal: Classify each image into one of 10 digit classes using SVM
- Demonstrates the power of SVM's RBF kernel for non-linear image classification

YHK

Import Libraries

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
```

YHK

Load and Visualize Data

```
digits = load_digits()
print("Data shape:", digits.data.shape)   # (1797, 64)
print("Classes   :", digits.target_names) # [0 1 2 ... 9]

fig, axes = plt.subplots(2, 5, figsize=(10, 4))
for ax, img, label in zip(axes.ravel(),
                          digits.images, digits.target):
    ax.imshow(img, cmap='gray_r')
    ax.set_title(str(label))
    ax.axis('off')
plt.suptitle("Sample Digits")
plt.show()
```

YHK

Train SVM Classifier

```
X_train, X_test, y_train, y_test = train_test_split(
    digits.data, digits.target,
    test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test  = scaler.transform(X_test)

svm = SVC(kernel='rbf', C=10, gamma=0.001)
svm.fit(X_train, y_train)
y_pred = svm.predict(X_test)

print(classification_report(y_test, y_pred))
```

YHK

Confusion Matrix

```
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=digits.target_names,
            yticklabels=digits.target_names)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("SVM Digit Recognition: Confusion Matrix")
plt.tight_layout()
plt.show()
```

YHK

Visualize Misclassifications

```
wrong = np.where(y_pred != y_test)[0]
fig, axes = plt.subplots(2, 5, figsize=(10, 4))
for ax, idx in zip(axes.ravel(), wrong[:10]):
    ax.imshow(X_test[idx].reshape(8, 8), cmap='gray_r')
    ax.set_title(f"True:{y_test[idx]} Pred:{y_pred[idx]}")
    ax.axis('off')
plt.suptitle("Misclassified Digits")
plt.tight_layout()
plt.show()
```

YHK

K-Means Clustering — Customer Segmentation

(Ref: scikit-learn Clustering User Guide)

YHK

Problem Info

- Scenario: A retail company wants to segment its customers for targeted marketing
- Features: Annual Income (k USD) and Spending Score (1–100)
- Data: Synthetic dataset mimicking a retail loyalty programme
- Goal: Discover natural customer groups using K-Means; choose optimal K with the Elbow method

YHK

Import Libraries and Generate Data

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

np.random.seed(42)
# Simulate 3 customer segments: budget, standard, premium
centres = [(30, 20), (55, 55), (80, 85)]
X = np.vstack([
    np.random.randn(120, 2) * [8, 10] + c
    for c in centres])
df = pd.DataFrame(X, columns=['AnnualIncome', 'SpendingScore'])
print(df.describe())
```

YHK

Elbow Method: Choosing K

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df)

inertia = []
K_range = range(1, 10)
for k in K_range:
    km = KMeans(n_clusters=k, random_state=42, n_init=10)
    km.fit(X_scaled)
    inertia.append(km.inertia_)

plt.plot(K_range, inertia, marker='o')
plt.xlabel("Number of Clusters K")
plt.ylabel("Inertia (WCSS)")
plt.title("Elbow Method")
plt.grid(True); plt.show()
```

YHK

Fit K-Means with Optimal K

```
km = KMeans(n_clusters=3, random_state=42, n_init=10)
df['Segment'] = km.fit_predict(X_scaled)

labels = {0: 'Budget', 1: 'Standard', 2: 'Premium'}
colours = ['#e07b54', '#5b9bd5', '#70ad47']

plt.figure(figsize=(7, 5))
for seg, (name, colour) in zip([0, 1, 2],
                               zip(labels.values(), colours)):
    mask = df['Segment'] == seg
    plt.scatter(df.loc[mask, 'AnnualIncome'],
                df.loc[mask, 'SpendingScore'],
                label=name, color=colour, alpha=0.7)
plt.xlabel("Annual Income (k USD)")
plt.ylabel("Spending Score")
plt.title("Customer Segments (K-Means, K=3)")
plt.legend(); plt.tight_layout(); plt.show()
```

YHK

Segment Profiles

- **Budget**: low income, low spending — price-sensitive customers
- **Standard**: mid income, mid spending — core customer base
- **Premium**: high income, high spending — loyalty programme targets

```
profile = df.groupby('Segment')[
    ['AnnualIncome','SpendingScore']].mean()
profile.index = [labels[i] for i in profile.index]
print(profile.round(1))
```

YHK

PCA + K-Means — Handwritten Digit Clustering

(Ref: scikit-learn Decomposition and Clustering User Guide)

YHK

Problem Info

- Dataset: Digits (`sklearn.datasets.load_digits`, built-in) — 1,797 images, 64 features each
- Part A: Use PCA to reduce 64 dimensions to 2; visualise digit clusters in 2D
- Part B: Apply K-Means (K=10) on PCA-reduced data; measure cluster quality
- Part C: Vary the number of PCA components (2, 10, 30) and observe effect on clustering

YHK

Import Libraries

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_digits
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score, adjusted_rand_score
```

YHK

Part A: PCA Visualisation

```
digits = load_digits()
X = StandardScaler().fit_transform(digits.data)

pca2 = PCA(n_components=2)
X2   = pca2.fit_transform(X)

plt.figure(figsize=(8, 6))
scatter = plt.scatter(X2[:, 0], X2[:, 1],
                      c=digits.target, cmap='tab10',
                      s=10, alpha=0.7)
plt.colorbar(scatter, label='Digit')
plt.title("Digits projected to 2 PCA components")
plt.xlabel("PC 1"); plt.ylabel("PC 2")
plt.tight_layout(); plt.show()

print("Variance explained:", pca2.explained_variance_ratio_.sum())
```

YHK

Part B: K-Means on PCA Data

- **Silhouette score**: measures cluster compactness (higher is better, max 1.0)
- **Adjusted Rand Index**: compares clusters to true digit labels (1.0 = perfect)

```
for n_comp in [2, 10, 30]:
    Xr = PCA(n_components=n_comp).fit_transform(X)
    km = KMeans(n_clusters=10, random_state=42, n_init=10)
    labels = km.fit_predict(Xr)
    sil = silhouette_score(Xr, labels)
    ari = adjusted_rand_score(digits.target, labels)
    print(f"PCA={n_comp:2d}  Silhouette={sil:.3f}  "
          f"ARI={ari:.3f}")
```

YHK

Assignment Tasks

- Report the variance explained by 2, 10, and 30 PCA components
- Plot the explained variance ratio (scree plot) and identify the "elbow"
- At which number of PCA components does ARI peak? Explain why
- Visualise cluster centres (inverse-transform to pixel space) for the best configuration
- What do the cluster centres look like? Do they resemble recognisable digits?

YHK

Assignments for Mechanical Engineering

YHK

Predictive Maintenance using Regression Analysis

- **Dataset**: NASA Turbofan Engine Degradation Simulation Data Set
- **Source**: NASA Prognostics Data Repository
- **Description**: Predict the remaining useful life of turbofan engines using regression analysis based on sensor data.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

data = pd.read_csv('turbofan_data.csv')
X = data[['feature1', 'feature2', 'feature3']]
y = data['remaining_useful_life']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
model = LinearRegression()
model.fit(X_train, y_train)
predictions = model.predict(X_test)
```

YHK

Structural Health Monitoring using Anomaly Detection

- **Dataset**: Bridge Health Monitoring Dataset
- **Source**: UCI Machine Learning Repository
- **Description**: Detect anomalies indicating potential structural failures in bridges based on sensor readings.

```
import pandas as pd
from sklearn.ensemble import IsolationForest

data = pd.read_csv('bridge_health_data.csv')
model = IsolationForest(contamination=0.1)
model.fit(data)
anomalies = model.predict(data)

```

YHK

Design Optimization using Genetic Algorithms

- **Dataset**: Engineering Design Optimization Dataset
- **Source**: Kaggle
- **Description**: Optimize engineering designs using genetic algorithms to evaluate performance metrics.

```
from deap import base, creator, tools, algorithms

creator.create("FitnessMin", base.Fitness, weights=(-1.0,))
creator.create("Individual", list, fitness=creator.FitnessMin)

toolbox = base.Toolbox()
# Define operations for crossover, mutation, etc.
population = toolbox.population(n=50)
algorithms.eaSimple(population, toolbox, cxpb=0.5, mutpb=0.2, ngen=40)
```

YHK

Fault Detection in Rotating Machinery using SVM

- **Dataset**: Machinery Fault Diagnosis Dataset
- **Source**: UCI Machine Learning Repository
- **Description**: Classify faults in rotating machinery using support vector machines (SVM).

```
import pandas as pd
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split

data = pd.read_csv('machinery_fault_data.csv')
X = data.drop('fault_type', axis=1)
y = data['fault_type']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25)
model = SVC()
model.fit(X_train, y_train)
accuracy = model.score(X_test, y_test)

```

YHK

Thermal Analysis using Neural Networks

- **Dataset**: Thermal Conductivity Dataset
- **Source**: Kaggle
- **Description**: Predict thermal conductivity of materials using neural networks.

```
from keras.models import Sequential
from keras.layers import Dense

model = Sequential()
model.add(Dense(64, activation='relu', input_dim=10))
model.add(Dense(32, activation='relu'))
model.add(Dense(1))

model.compile(optimizer='adam', loss='mean_squared_error')

```

YHK

Quality Control in Manufacturing using CNNs

- **Dataset**: Manufacturing Defect Dataset
- **Source**: Kaggle
- **Description**: Train a CNN to classify manufacturing defects based on part images.

```
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

model = Sequential()
model.add(Conv2D(32, (3,3), activation='relu', input_shape=(128,128,3)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(10, activation='softmax'))

model.compile(optimizer='adam', loss='categorical_crossentropy',
     metrics=['accuracy'])
```

YHK

Energy Consumption Prediction using Time Series Analysis

- **Dataset**: Energy Consumption Dataset
- **Source**: UCI Machine Learning Repository
- **Description**: Forecast energy consumption patterns using ARIMA models.

```
import pandas as pd
from statsmodels.tsa.arima.model import ARIMA

data = pd.read_csv('energy_consumption.csv')
model = ARIMA(data['consumption'], order=(5,1,0))
model_fit = model.fit()
forecast = model_fit.forecast(steps=10)
```

YHK

Material Property Prediction using Gaussian Processes

- **Dataset**: Materials Project Database
- **Source**: Materials Project
- **Description**: Predict material properties using Gaussian process regression.

```
from sklearn.gaussian_process import GaussianProcessRegressor

gp = GaussianProcessRegressor()
gp.fit(X_train, y_train)
predictions = gp.predict(X_test)
```

YHK

Load Prediction in Structural Engineering using Random Forests

- **Dataset**: Structural Load Data
- **Source**: UCI Machine Learning Repository
- **Description**: Predict structural loads using random forests.

```
import pandas as pd
from sklearn.ensemble import RandomForestRegressor

data = pd.read_csv('structural_load_data.csv')
X = data.drop('load', axis=1)
y = data['load']

model = RandomForestRegressor()
model.fit(X, y)
predictions = model.predict(X)

```

YHK

Simulation of Fluid Dynamics using Deep Learning

- **Dataset**: Fluid Dynamics Simulation Data
- **Source**: OpenFOAM
- **Description**: Simulate fluid dynamics using deep learning models.

```
import tensorflow as tf

model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3,3), activation='relu',
     input_shape=(64,64,1)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1)
])

model.compile(optimizer='adam', loss='mean_squared_error')
```

YHK

Thanks . . .

- Office Hours: Saturdays, 3 to 5 pm (IST); Free-Open to all; email for appointment to yogeshkulkarni at yahoo dot com
- Call + 9 1 9 8 9 0 2 5 1 4 0 6

(https://www.linkedin.com/in/yogeshkulkarni/)

(https://medium.com/@yogeshharibhaukulkarni)

(https://www.github.com/yogeshhk/)

YHK

Temporary page!

LaTeX was unable to guess the total number of pages correctly. As there was some unprocessed data that should have been added to the final page this extra page has been added to receive it.
If you rerun the document (without altering it) this surplus page will go away because LaTeX now knows how many pages to expect for this document.