

# INTRODUCTION TO DEEP LEARNING

Yogesh Kulkarni

February 21, 2020

# Introduction to Machine Learning

## How do we learn?

- ▶ What do we do when we have to prepare for an examination?
- ▶ Study. Learn. Imbibe. Take notes. Practice mock papers.
- ▶ Thus, prepare for the unseen test.

# What is Learning?

*"Learning is any process by which a system improves performance from experience."*

- Herbert Simon, Turing Award 1975, Nobel in Economics 1978.

## What is Machine Learning?

Machine learning is a type of artificial intelligence (AI) which:

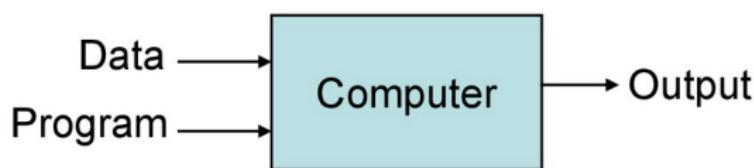
- ▶ Learns function without being explicitly programmed.
- ▶ Can grow and change when exposed to new data.

## So, What is Machine Learning?

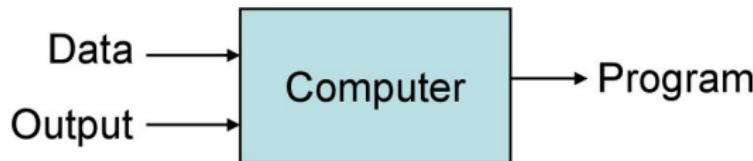
- ▶ Ability of computers to “learn” from “data”
- ▶ Learn: Discover patterns, underlying structure
- ▶ Data: Comes from sensors, transactions, etc.

# Traditional vs. Machine Learning?

## Traditional Programming



## Machine Learning



# Why Machine Learning?

- ▶ Problems with High Dimensionality
- ▶ Hard/Expensive to program manually
- ▶ Techniques to model 'ANY' function given 'ENOUGH' data.
- ▶ Job \$\$\$

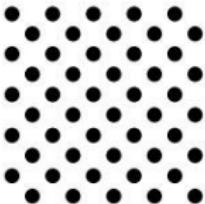
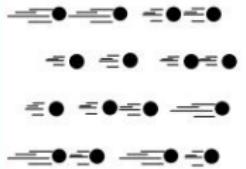
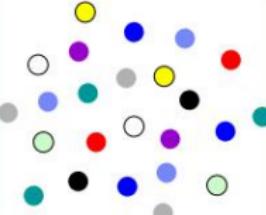
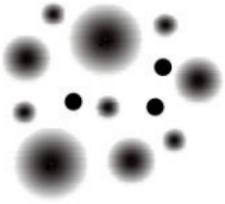
## Why now?

- ▶ Flood of data (Internet, IoT)
- ▶ Increasing computational power
- ▶ Easy/free availability of algorithms
- ▶ Increasing support from industries

## The storm: The Big Data is coming

- ▶ In 2012, HBR put Data Scientists on the radar
- ▶ “The Sexiest Job of the 21st Century”.
- ▶ Industry, trying to be data-driven, than manual.

## (Big) Data Characteristics

Volume	Velocity	Variety	Veracity*
			
<b>Data at Rest</b>  Terabytes to exabytes of existing data to process	<b>Data in Motion</b>  Streaming data, milliseconds to seconds to respond	<b>Data in Many Forms</b>  Structured, unstructured, text, multimedia	<b>Data in Doubt</b>  Uncertainty due to data inconsistency & incompleteness, ambiguities, latency, deception, model approximations

(Image Credit: <http://www.rosebt.com/blog/data-veracity>)

# What's the answer?

## AI-ML-DL

- ▶ Machines showing intelligence of Humans
- ▶ Machine Learning: part of AI
- ▶ Logic is not programmed by hand,
- ▶ Gets emerged in training with data.

# Types of Machine Learning

## Two kinds of learning

- ▶ Supervised
- ▶ Unsupervised

# Supervised

- ▶ Training data with correct answers
- ▶ Both used to train the model
- ▶ Then apply unseen data on model

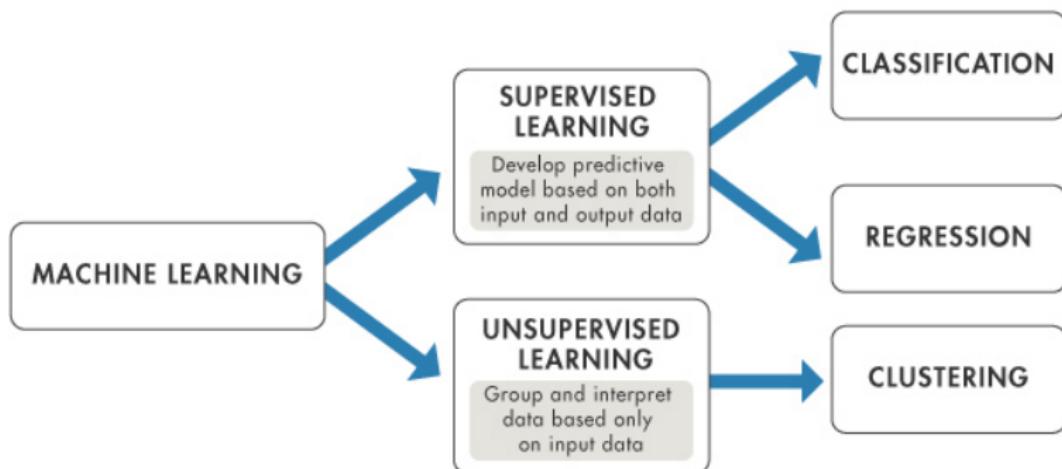
# Unsupervised

- ▶ Training data with no answers
- ▶ Extract patterns, groups

## Some types of algorithms

- ▶ Prediction: predicting a continuous variable from data
- ▶ Classification: assigning records to predefined groups
- ▶ Clustering: splitting records into groups based on similarity
- ▶ Association learning: seeing what often appears together

# Machine Learning Learning Algorithms



(Reference: Machine Learning in MATLAB - MATLAB & Simulink - MathWorks)

# Machine Learning Learning Algorithms

- ▶ Is this A or B? : Classification algorithms
- ▶ Is this weird? : Anomaly detection algorithms
- ▶ How much—or—How many? : Regression algorithms
- ▶ How is this organized? : Clustering algorithms, Dimensionality reduction
- ▶ What should I do next? : Reinforcement learning algorithms

(Ref: Brandon Rohrer's breakdown of the "5 questions data science answers")

# Classification

- ▶ **Description:** Identifying the category an object belongs to.
- ▶ **Applications:** Spam detection, Image recognition.
- ▶ **Algorithms:** SVM, nearest neighbors, random forest, Logistic Regression

# Regression

- ▶ **Description:** Predicting a continuous-valued attribute associated with an object.
- ▶ **Applications:** Drug response, Stock prices.
- ▶ **Algorithms:** Linear Regression

# Clustering

- ▶ **Description:** Automatic grouping of similar objects into sets.
- ▶ **Applications:** Customer segmentation, Grouping experiment outcomes
- ▶ **Algorithms:** k-Means

# Dimensionality Reduction

- ▶ **Description:** Reducing the number of random variables to consider.
- ▶ **Applications:** Visualization, Increased efficiency
- ▶ **Algorithms:** PCA, Singular Value Decomposition

# Popular Algorithms in Machine Learning

- ▶ Linear, Logistic Regression
- ▶ Decision Trees
- ▶ SVM - Support Vector Machines, Naive Bayes
- ▶ K-Means

# Applications of Machine Learning

# Everyday Applications of Machine Learning

- ▶ Face Recognition (Facebook)
- ▶ Spam recognition in Emails
- ▶ Recommender Systems
- ▶ Feelings Analysis, Sentiments
- ▶ Natural language: Translate a sentence from Hindi to English, question answering, etc.
- ▶ Speech: Recognize spoken words, speaking sentences naturally
- ▶ Game playing: Play games like chess
- ▶ Robotics: Walking, jumping, displaying emotions, etc.
- ▶ Driving a car, flying a plane, navigating a maze, etc.

## Cool-down: Summary

SO ...

- ▶ What is Machine learning, after-all?
- ▶ Its usage in your domain?

# Introduction to Deep Learning

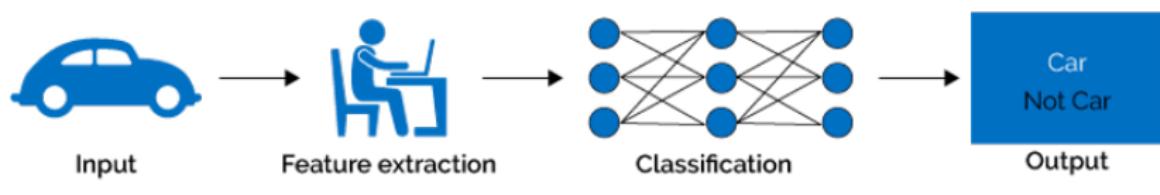
## What is AI-ML-DL?

- ▶ Artificial Intelligence: mimicking human intelligence
- ▶ Machine Learning: Automating Learning with features.
- ▶ ML: human-designed representations and input features. So, its just optimizing weights to best make a final prediction
- ▶ There could be programmed (hand coded) AI, that's not Machine Learning
- ▶ Machine Learning could be for non AI activities, like automation
- ▶ Deep Learning: Neural network with no input features

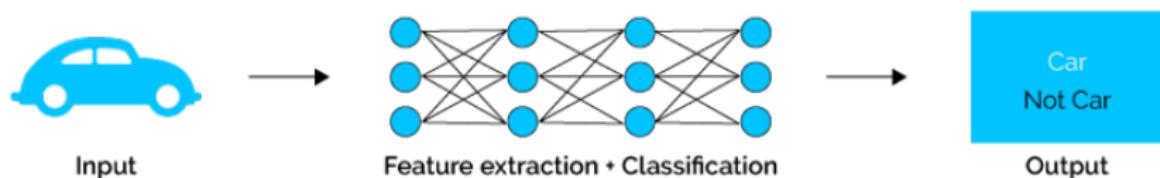
## ML vs DL: What's the difference?

Deep learning algorithms attempt to learn (multiple levels of) representation by using a hierarchy of multiple layers

### Machine Learning

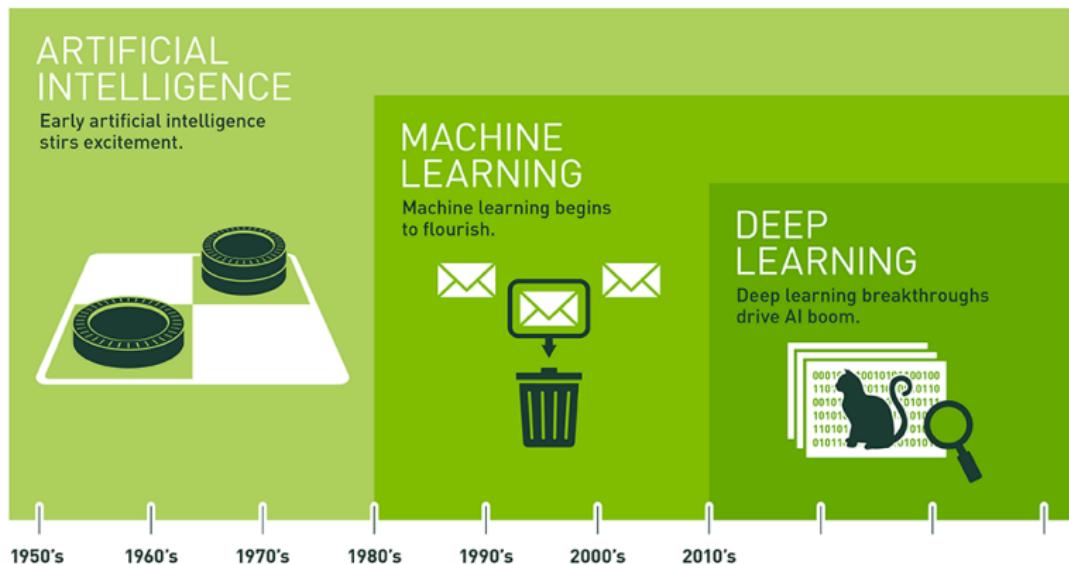


### Deep Learning



(Reference: <https://www.xenonstack.com/blog/static/public/uploads/media/machine-learning-vs-deep-learning.png>)

# AI ML DL: Relationship



Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

(Reference: The Difference Between AI, Machine Learning, and Deep Learning - NVIDIA Blog)

## Use Deep Learning When ...

- ▶ You have lots of data (about 10k+ examples)
- ▶ The problem is “complex” - speech, vision, natural language
- ▶ The data is unstructured
- ▶ You need the absolute “best” model

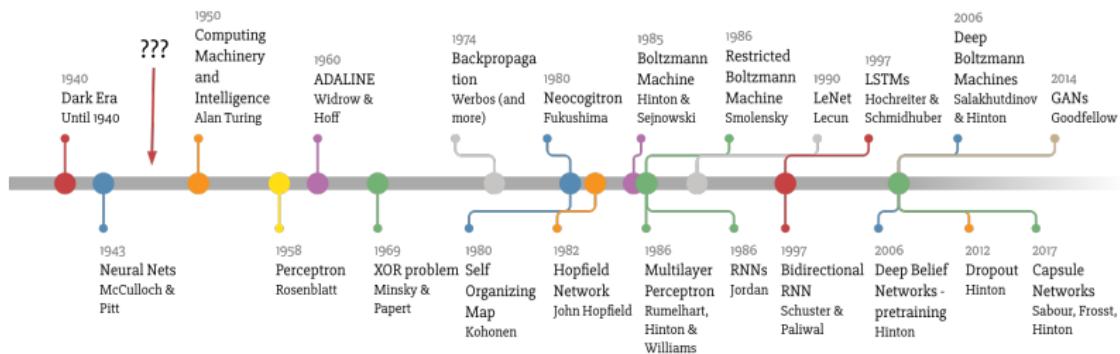
(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

## Don't use Deep Learning When . . .

- ▶ You don't have a large dataset
- ▶ You are performing sufficiently well with traditional ML methods
- ▶ Your data is structured and you possess the proper domain knowledge
- ▶ Your model should be explainable

(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

## History



Made by Favio Vázquez

(Reference: Deep Learning basics - Rodrigo Agundez)

# History

- ▶ 1958 - Perceptron unit - Frank Rosenblatt
- ▶ 1986 - Backpropagation - Geoffrey Hinton
- ▶ 1986 - RNN - Schuster & Paliwal
- ▶ 1989 - LeNet Backpropagation to multi-layer perceptron - Yan LeCun
- ▶ 1997 - LSTM - Sepp Hochreiter and Jürgen Schmidhuber
- ▶ 1998 - LeNet-5 Convolutional neural networks - Yan LeCun
- ▶ 2007 - Fei Fei Li Princeton ImageNet competition
- ▶ 2009 - GPU for deep learning - Andrew Ng
- ▶ 2011 - Demonstration of ReLU for deep neural networks - Yoshua Bengio
- ▶ 2012 - AlexNet wins ImageNet 25% to 16% error
- ▶ 2012 - Dropout technique - Geoffrey Hinton
- ▶ 2014 - Generative adversarial networks - Ian Goodfellow & Yoshua Bengio
- ▶ 2015 - CNN beats human error in ImageNet 5% to 3%
- ▶ 2016 - AlphaGo - Google DeepMind
- ▶ 2016 - Detectic cancer beats human pathologist .96 vs 0.99 AUC
- ▶ 2017 - Capsule networks - Geoffrey Hinton

# History

“



{

**There are many interesting recent development in deep learning...The most important one, in my opinion, is adversarial training (also called GAN for Generative Adversarial Networks). This, and the variations that are now being proposed, is the most interesting idea in the last 10 years in ML.**

Yann LeCun

}

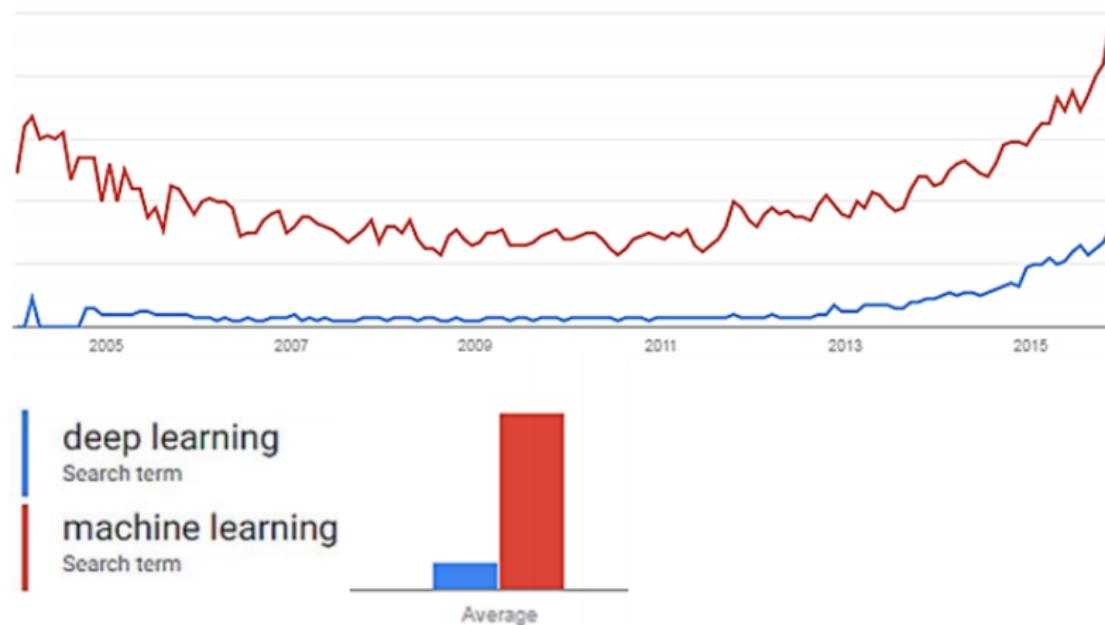
(Reference: Deep Learning basics - Rodrigo Agundez)

## Why is DL useful?

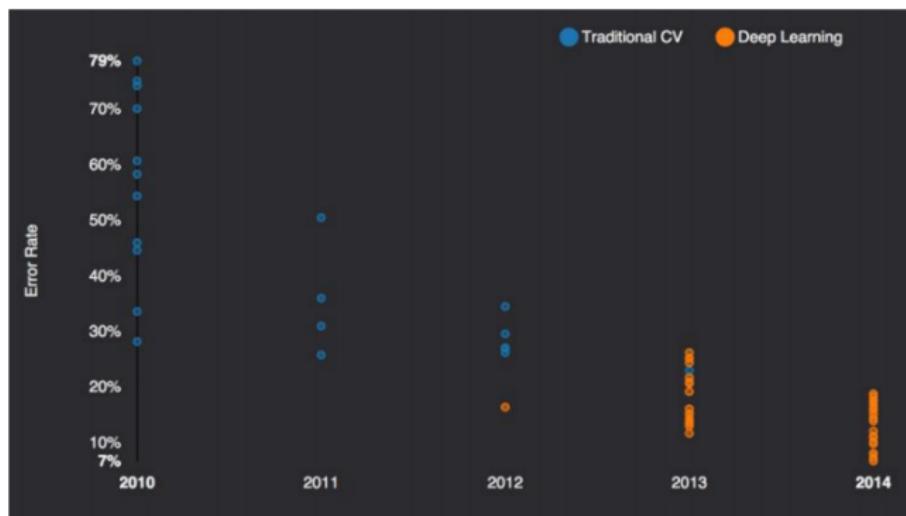
- ▶ ML features could be over-specified, incomplete and take longer to design
- ▶ DL “invents” features.
- ▶ Learned Features are easy to adapt, fast to learn
- ▶ Deep learning provides a very flexible, (almost?) universal, learnable framework for representing world, visual and linguistic information

## Why is DL useful?

- ▶ In 2010 DL started outperforming other ML techniques
- ▶ First in speech and vision, then NLP

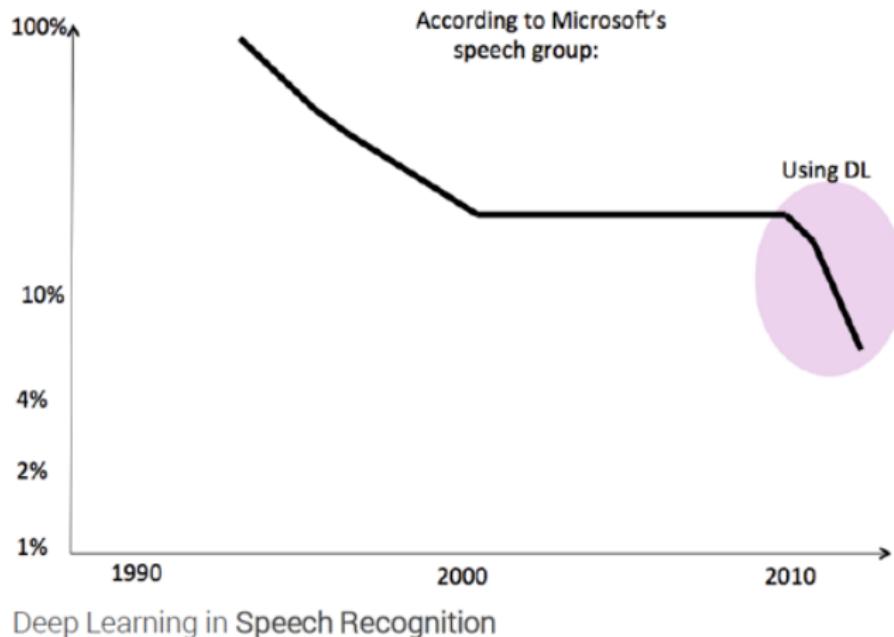


# Big Break-through in Vision



(Reference: Introduction to Deep Learning - Ismini Lourentzou)

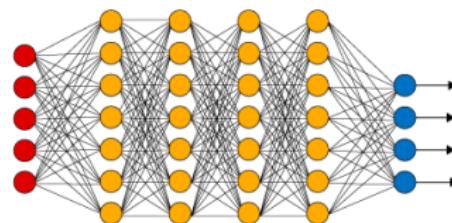
# Big Break-through in Speech



(Reference: Introduction to Deep Learning - Ismini Lourentzou)

# Deep Learning == Neural Nets

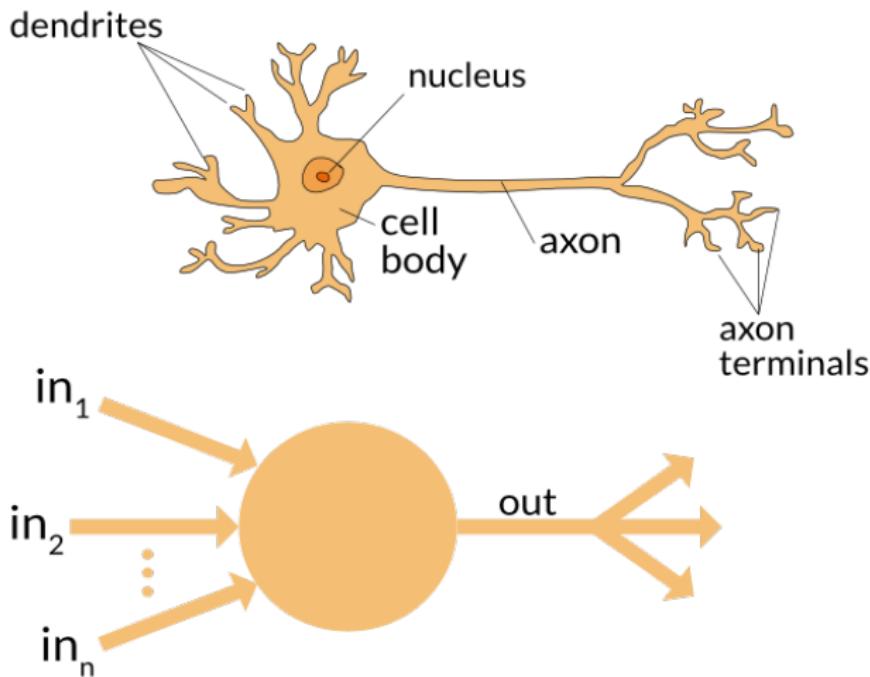
- ▶ Main idea of deep learning: transform the input space into outputs via higher level abstractions.
- ▶ Neural Net architectures are made up of perceptrons (similar to neurons)
- ▶ Each neuron carries certain transformations on inputs coming to it.
- ▶ Collection of such neurons with various types of transformations, can create desired overall transformation.



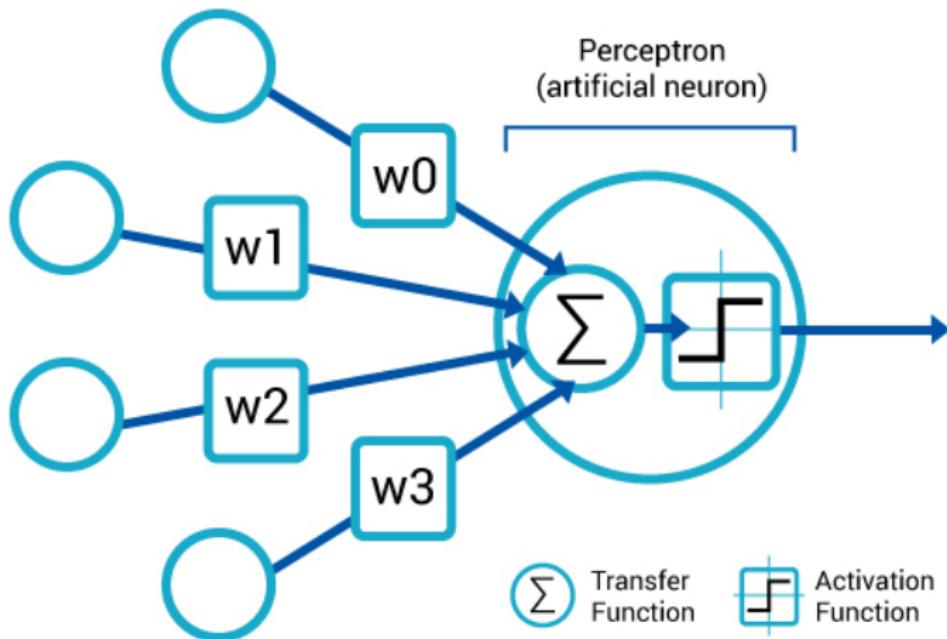
(Reference: Deep Learning basics - Rodrigo Agundez)

# Deep Learning == Neural Nets

First artificial neuron proposed in 1943!



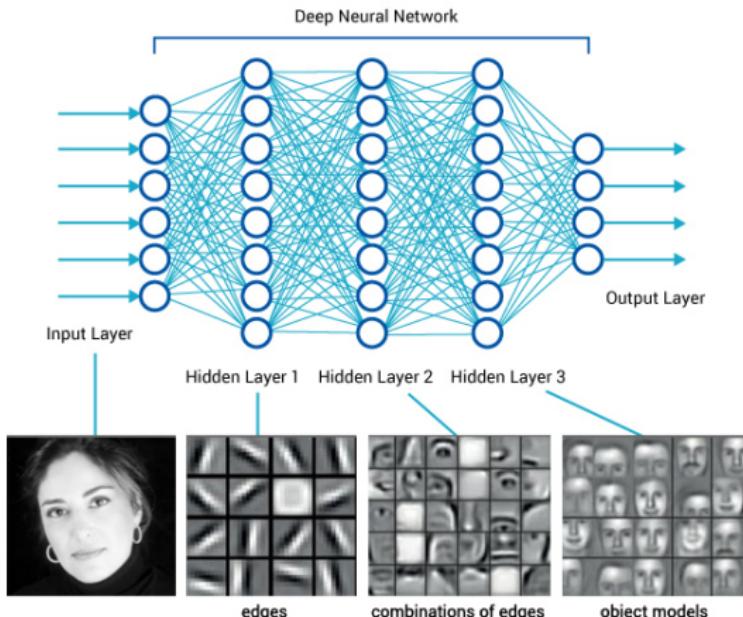
# Artificial neuron



(Reference: Deep Learning basics - Rodrigo Agundez)

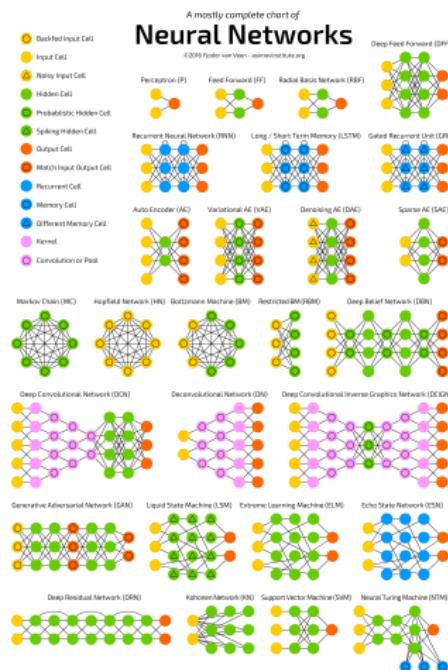
# Layers

## Hierarchical feature representations



(Reference: Deep Learning basics - Rodrigo Agundez)

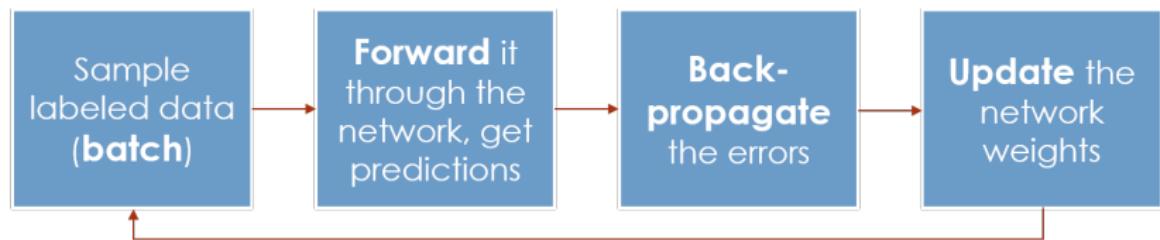
# Neural Networks



(Reference: Deep Learning basics - Rodrigo Agundez)

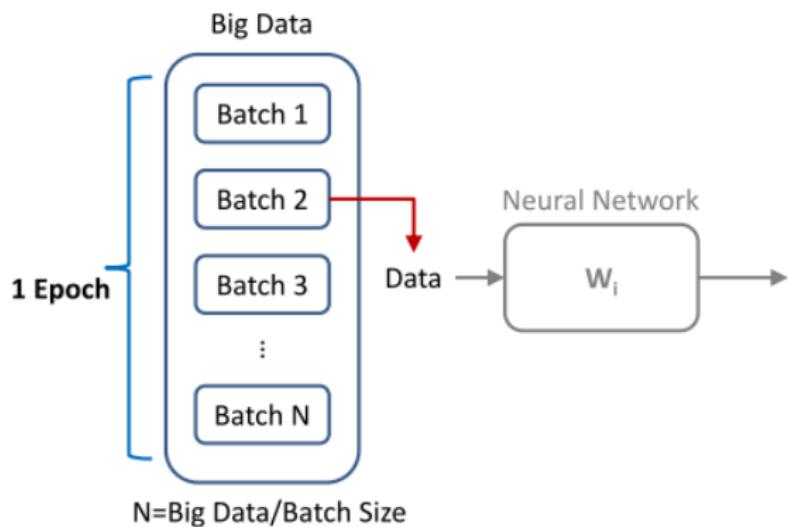
# Neural Networks Training Process: Non Mathematical

## Training Process: Non Mathematical



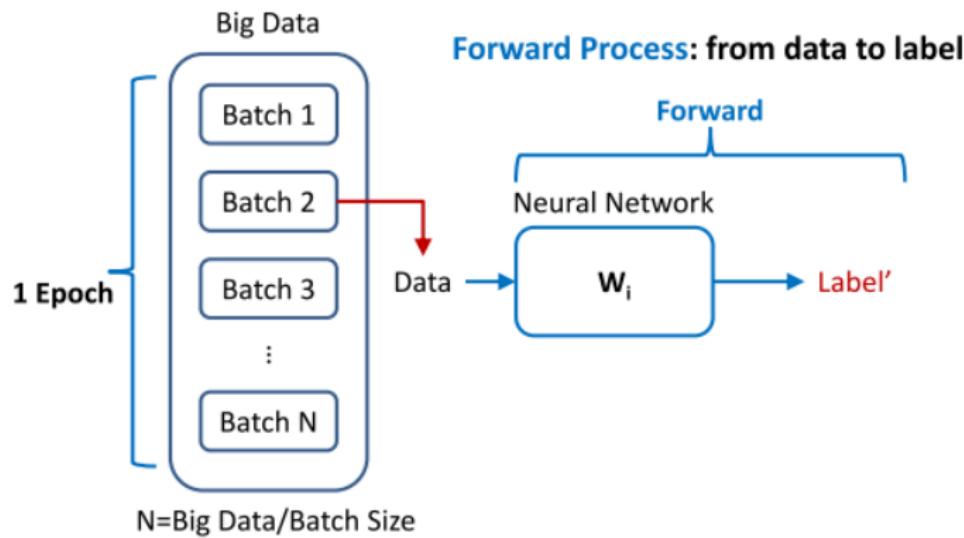
(Reference: Introduction to Deep Learning - Ismini Lourentzou)

# Data Enters



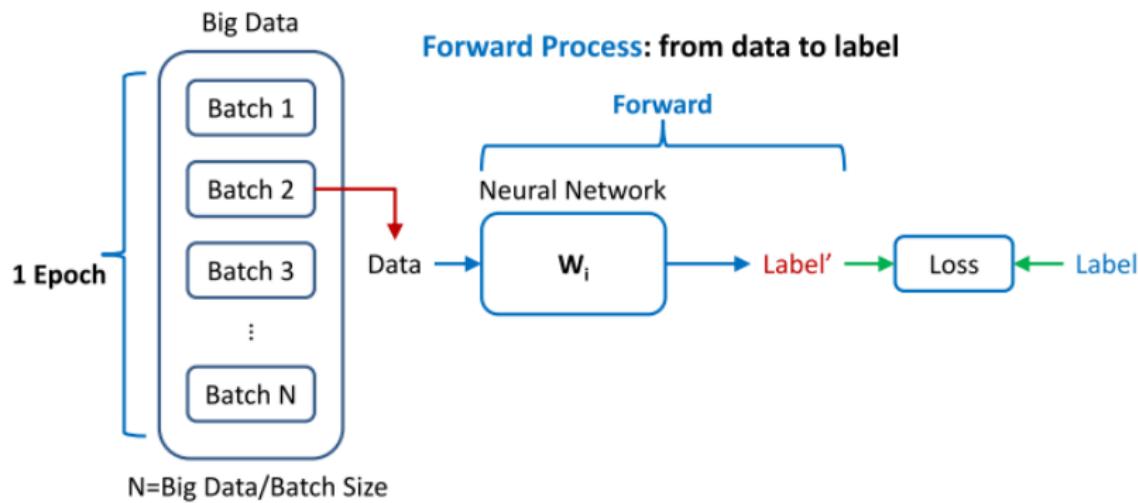
(Reference:PyTorch Tutorial-NTU Machine Learning Course-Lyman Lin )

## Forward Pass



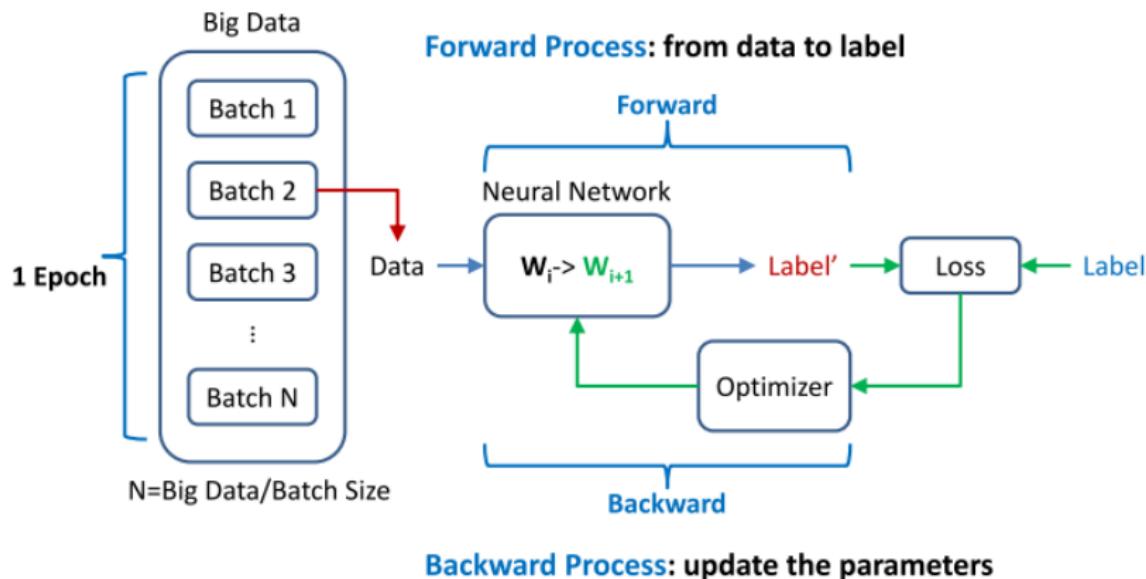
(Reference:PyTorch Tutorial-NTU Machine Learning Course-Lyman Lin )

# Loss Calculations



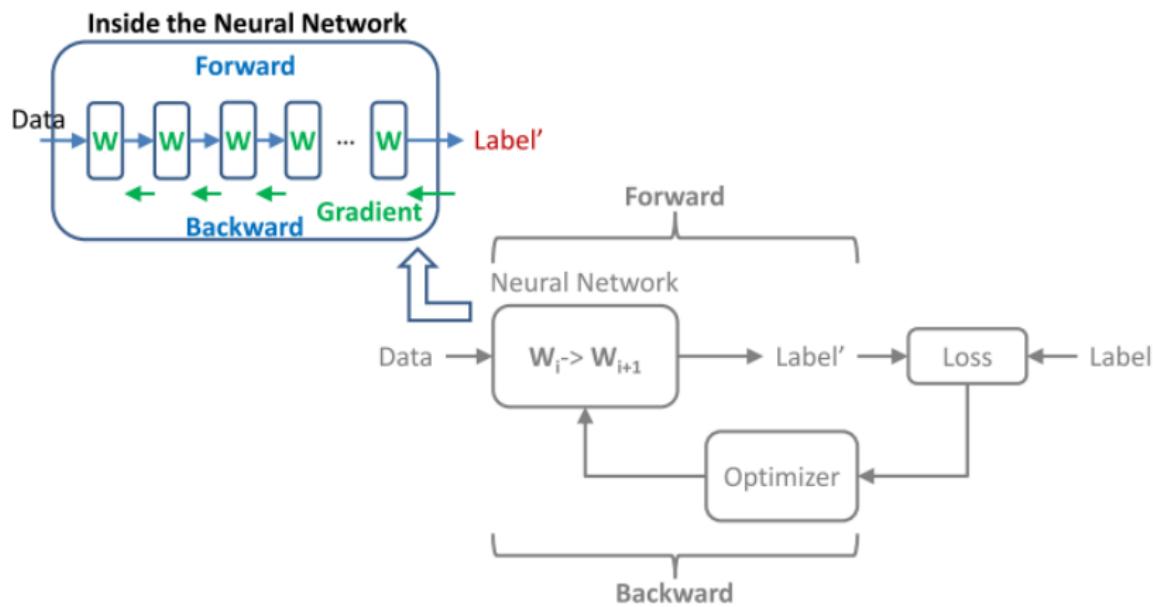
(Reference:PyTorch Tutorial-NTU Machine Learning Course-Lyman Lin )

# Back Propagation



(Reference:PyTorch Tutorial-NTU Machine Learning Course-Lyman Lin )

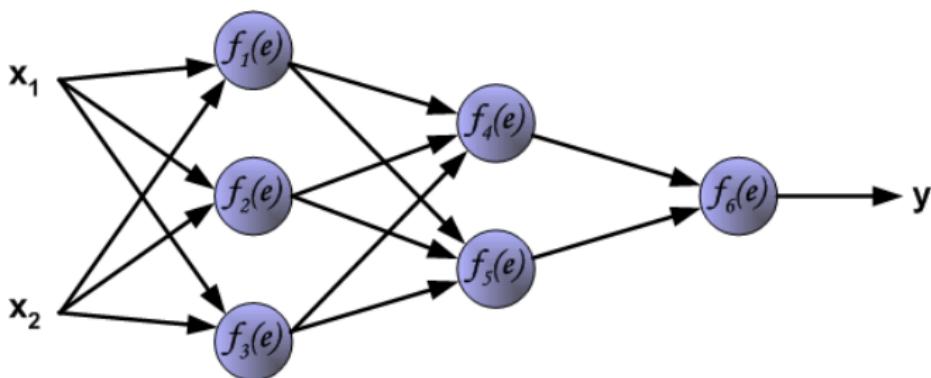
# Overall Process



(Reference:PyTorch Tutorial-NTU Machine Learning Course-Lyman Lin )

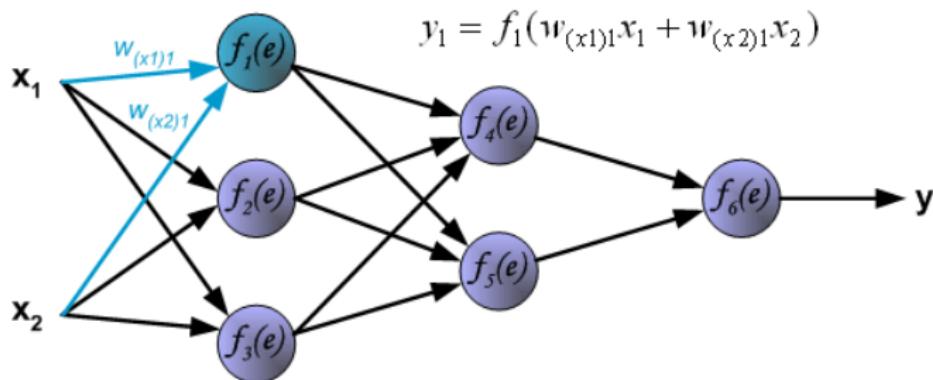
# Neural Networks Training Process: Mathematical

## Start



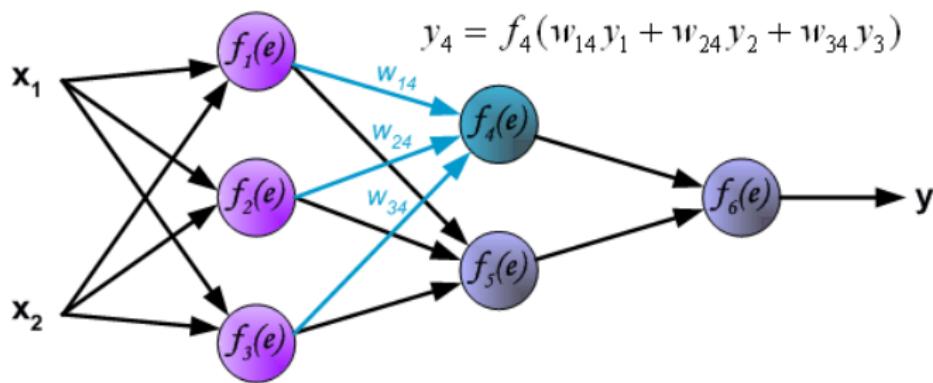
(Reference: Deep Learning basics - Rodrigo Agundez)

## Forward pass



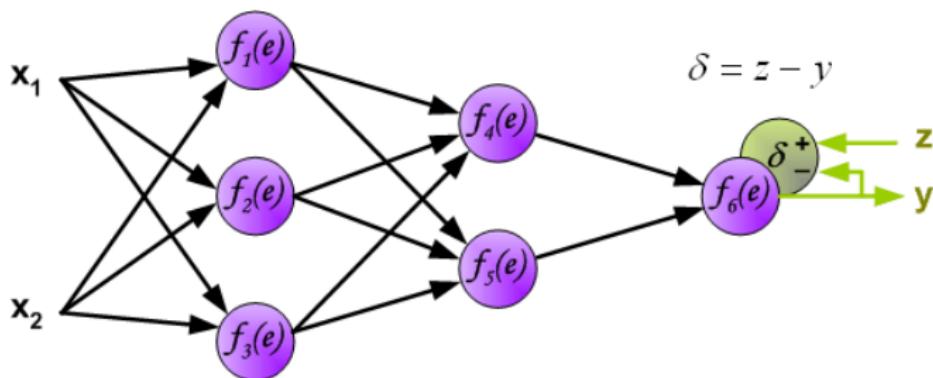
(Reference: Deep Learning basics - Rodrigo Agundez)

## Forward pass



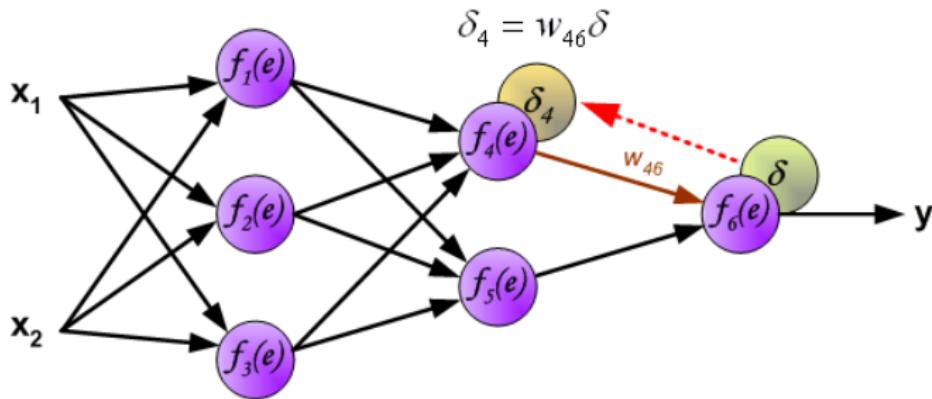
(Reference: Deep Learning basics - Rodrigo Agundez)

## Forward pass



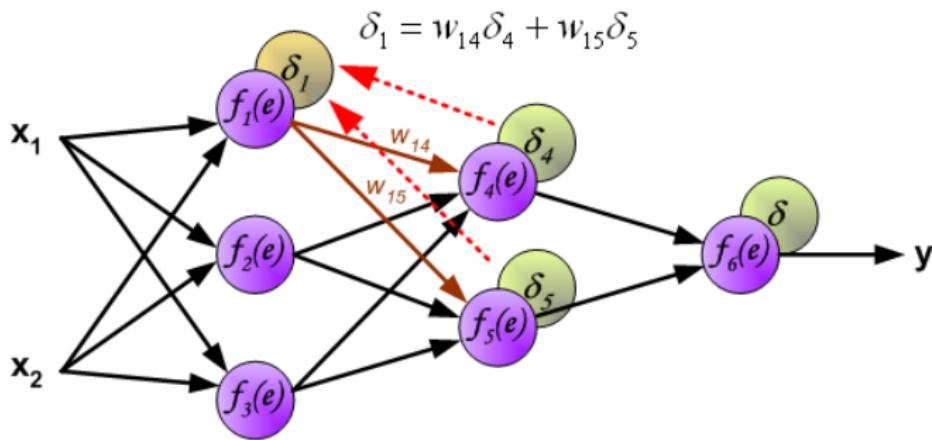
(Reference: Deep Learning basics - Rodrigo Agundez)

# Backpropagation



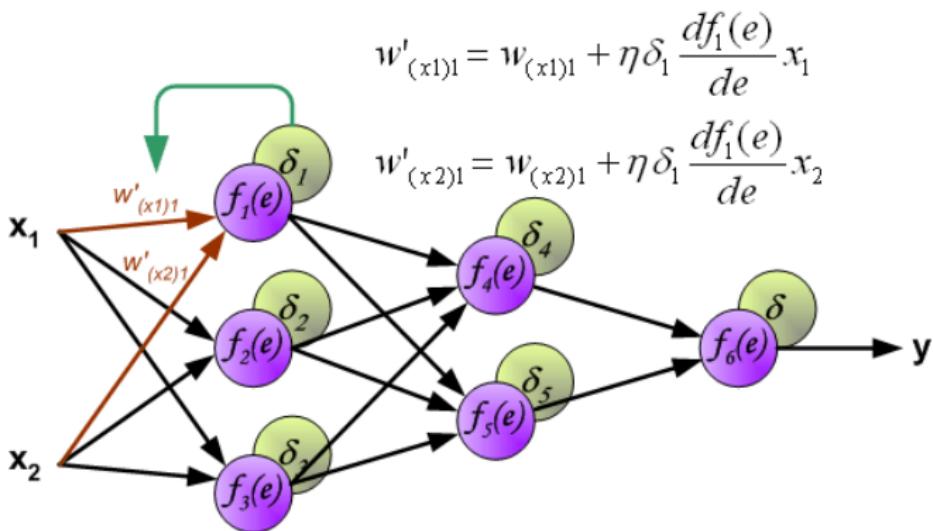
(Reference: Deep Learning basics - Rodrigo Agundez)

# Backpropagation



(Reference: Deep Learning basics - Rodrigo Agundez)

# Backpropagation

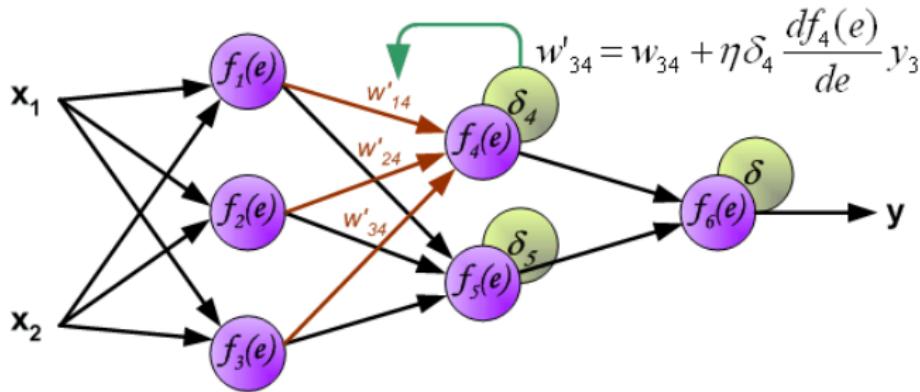


(Reference: Deep Learning basics - Rodrigo Agundez)

# Backpropagation

$$w'_{14} = w_{14} + \eta \delta_4 \frac{df_4(e)}{de} y_1$$

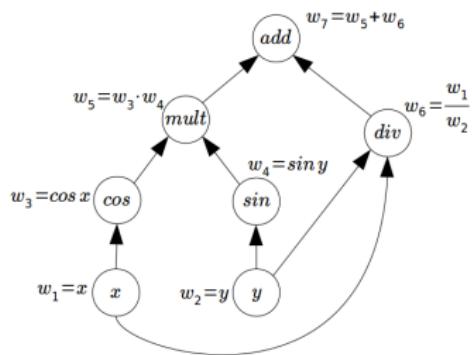
$$w'_{24} = w_{24} + \eta \delta_4 \frac{df_4(e)}{de} y_2$$



(Reference: Deep Learning basics - Rodrigo Agundez)

# Backpropagation

- ▶ Starts at the end of the net and tunes each layer using the gradient of the loss function. Repeatedly applies the chain rule.
- ▶ Numeric approximation:  $f'(x) \approx \frac{f(x+h) - f(x)}{h}$
- ▶ Symbolic differentiation: Symbolic, exact representation of the derivative.
- ▶ Reverse automatic differentiation

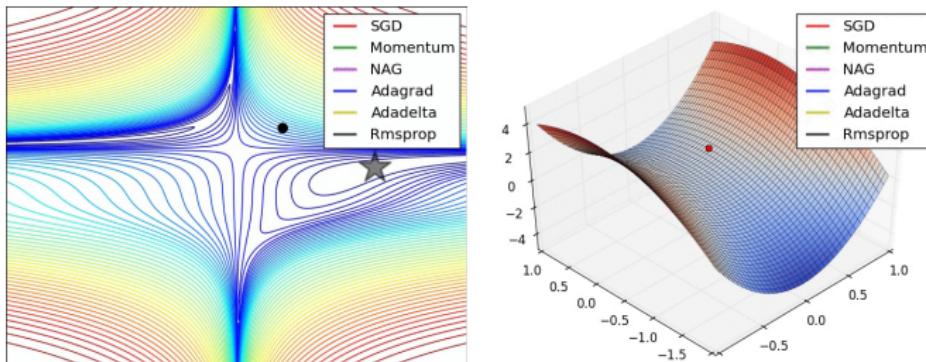


$$\begin{aligned}
 \frac{\partial w_7}{\partial w_1} &= \frac{\partial w_3}{\partial w_1} \frac{\partial w_7}{\partial w_3} + \frac{\partial w_6}{\partial w_1} \frac{\partial w_7}{\partial w_6} = -\sin w_1 \frac{\partial w_7}{\partial w_3} + \frac{1}{w_2} \frac{\partial w_7}{\partial w_6} \\
 \frac{\partial w_7}{\partial w_2} &= \frac{\partial w_4}{\partial w_2} \frac{\partial w_7}{\partial w_4} + \frac{\partial w_6}{\partial w_2} \frac{\partial w_7}{\partial w_6} = \cos w_2 \frac{\partial w_7}{\partial w_4} - \frac{w_1}{w_2^2} \frac{\partial w_7}{\partial w_6} \\
 \frac{\partial w_7}{\partial w_3} &= \frac{\partial w_5}{\partial w_3} \frac{\partial w_7}{\partial w_5} = w_4 \frac{\partial w_7}{\partial w_5} \\
 \frac{\partial w_7}{\partial w_4} &= \frac{\partial w_5}{\partial w_4} \frac{\partial w_7}{\partial w_5} = w_3 \frac{\partial w_7}{\partial w_4} \\
 \frac{\partial w_7}{\partial w_5} &= 1 \\
 \frac{\partial w_7}{\partial w_6} &= 1
 \end{aligned}$$

(Reference: Deep Learning basics - Rodrigo Agundez)

# Optimization by backpropagation

- ▶ Loss/cost function
- ▶ Gradient descent



(Reference: Deep Learning basics - Rodrigo Agundez)

## Other Aspects

# Challenges of Deep Learning

- ▶ Explain-ability - How do you learn what you learn?
- ▶ Debugging - What has gone wrong?
- ▶ Why is this not converging?

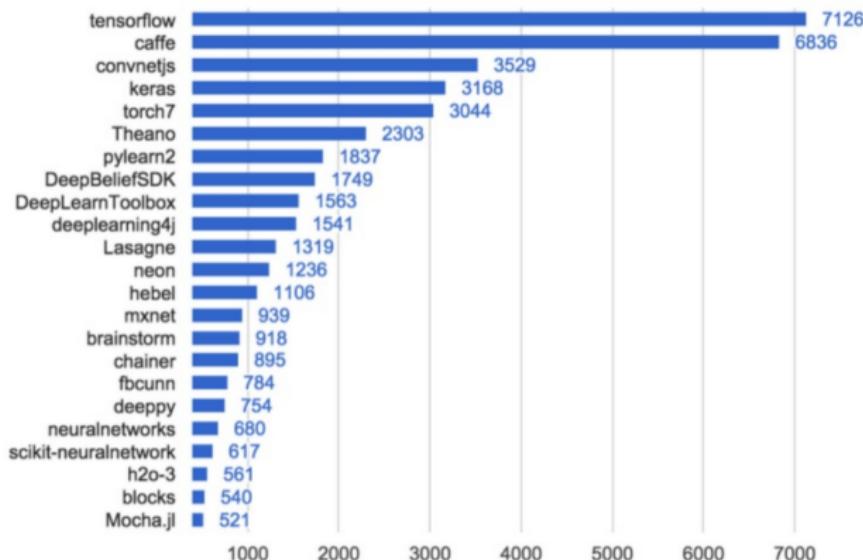
(Reference: AI and Deep Learning - Subrat Panda)

# Usage Requirements

- ▶ Large data set with good quality (input-output mappings)
- ▶ Measurable and describable goals (define the cost)
- ▶ Enough computing power (AWS GPU Instance)
- ▶ Excels in tasks where the basic unit (pixel, word) has very little meaning in itself, but the combination of such units has a useful meaning.

(Deep Learning - The Past, Present and Future of Artificial Intelligence - Lukas Masuch)

# Deep Learning Tools



(Reference: Introduction to Deep Learning - Ismini Lourentzou)

## Deep Learning Outlook

- ▶ Significant advances in deep reinforcement and unsupervised learning
- ▶ Bigger and more complex architectures
- ▶ Harder problems being attempted

# Introduction to TensorFlow 2.0

## TensorFlow is

- ▶ Open source, Free library, with Python bindings, by Google Brain team
- ▶ Other libraries are: Caffe (Berkeley), Torch (Facebook), Cntk (Microsoft),
- ▶ Can deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device with a single API
- ▶ Flexibility: from Raspberry Pi, Android, Windows, iOS, Linux to server farms
- ▶ Till 2019 Keras was popular as a separate library (with back-end as Tensorflow) but with Tensorflow 2.0, Keras has become its default front end API.
- ▶ TensorFlow 2.0 merges keras as "tf.keras". It allows you to design, fit, evaluate deep learning models.

# Open Source Community

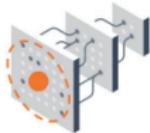
41,000,000+    69,000+    12,000+    2,200+

downloads    commits    pull requests    contributors

As of Oct 2019 ...

(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

# Tensorflow 2.0



## Easy

Simplified APIs.  
Focused on Keras and  
eager execution



## Powerful

Flexibility and performance.  
Power to do cutting edge research  
and scale to > 1 exaflops



## Scalable

Tested at Google-scale.  
Deploy everywhere

(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

# Deploy Anywhere

Servers



Edge devices



JavaScript



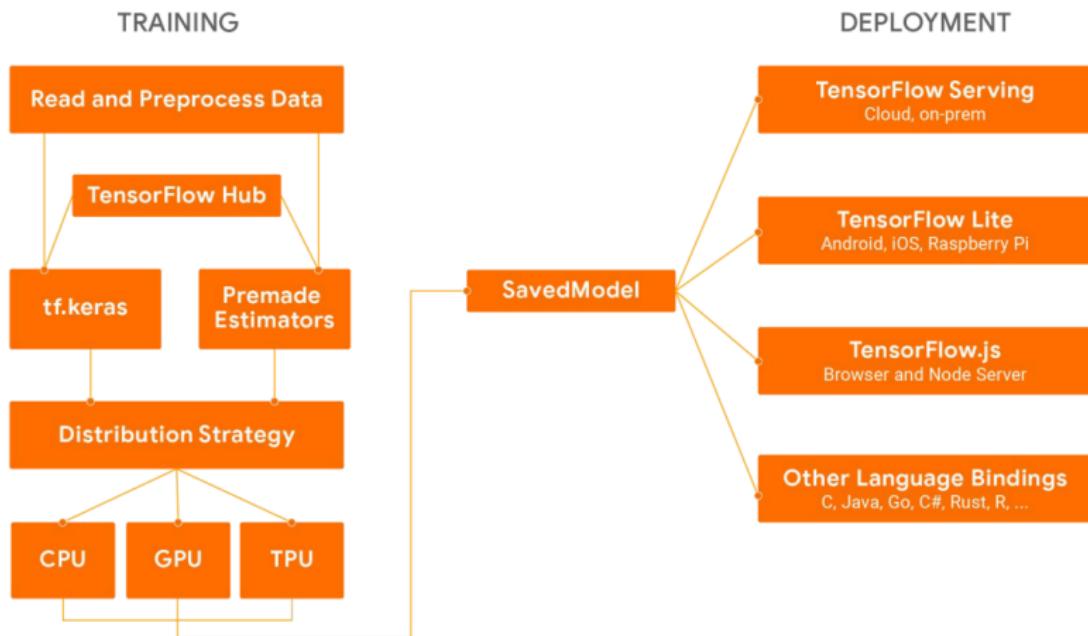
TensorFlow  
Extended

TensorFlow  
Lite

TensorFlow  
.JS

(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

# Training and Deployment



(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

# Ecosystem/Verticals

TF Probability

TF Agents

Tensor2Tensor

TF Ranking

TF Text

TF Federated

TF Privacy

...

# Compared to TF 1.0

## What's Gone

- ▶ `Session.run`
- ▶ `tf.control_dependencies`
- ▶ `tf.global_variables_initializer`
- ▶ `tf.cond`, `tf.while_loop`
- ▶ `tf.contrib`

(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

## What's New

- ▶ Eager execution by default
- ▶ `tf.function`
- ▶ Keras as main high-level api

# tf.keras



(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

# Installation

- ▶ Have Python installed, such as Python 3.6 or higher.
- ▶ Easy way to install TensorFlow
- ▶ Linux:

```
sudo pip install tensorflow
```

- ▶ Windows:

```
1 pip install tensorflow
```

# Installation

An end-to-end open source machine learning platform

TensorFlow For JavaScript For Mobile & IoT For Production

The core open-source library to help you develop and train ML models. Get started quickly by running Colab notebooks directly in your browser.

Get started with TensorFlow

(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

## Installation Check

- ▶ Confirm the installation by:

```
1 # check version
  import tensorflow
3 print(tensorflow.__version__)
```

- ▶ It must be 2.0 onwards
- ▶ If you get warning like below, Don't worry, just IGNORE.

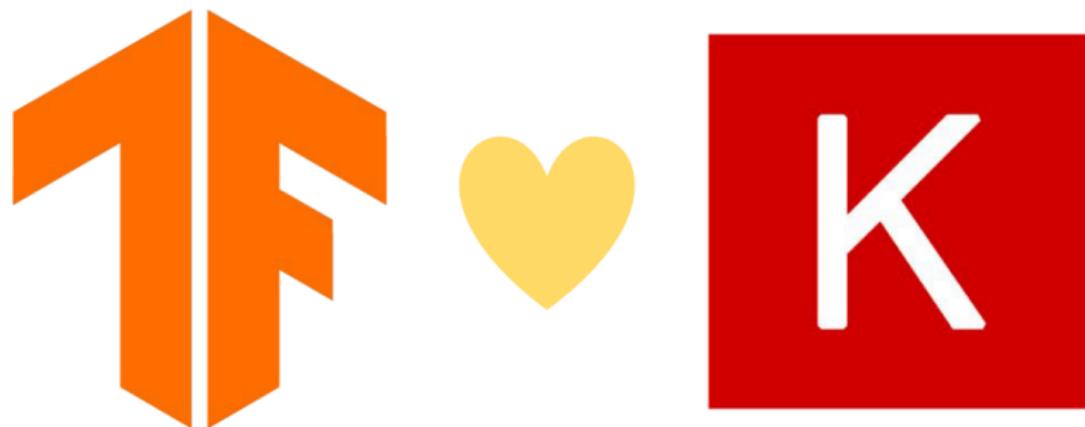
```
1 Your CPU supports instructions that this TensorFlow binary was not
  compiled to use: AVX2 FMA
  XLA service 0x7fde3f2e6180 executing computations on platform Host.
  Devices:
3 StreamExecutor device (0): Host, Default Version
```

## Ecosystem/Verticals

```
1 import tensorflow as tf # Assuming TF 2.0 is installed
2 a = tf.constant([[1, 2],[3, 4]])
3 b = tf.matmul(a, a)
4 print(b)
5 # tf.Tensor( [[ 7 10] [15 22]], shape=(2, 2), dtype=int32)
6 print(type(b.numpy()))
7 # <class 'numpy.ndarray'>
```

(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

# tf.keras



(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

## Keras and tf .keras

- ▶ Fast prototyping, advanced research, and production
- ▶ keras.io = reference implementation `import keras`
- ▶ tf .keras = TensorFlow's implementation (a superset, built-in to TF, no need to install Keras separately) `from tensorflow import keras`

(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

## Steps to use tf.Keras

- ▶ Define the model.
- ▶ Compile the model.
- ▶ Fit the model.
- ▶ Evaluate the model.
- ▶ Make predictions.

## Define the Model

- ▶ First, select the type of the model.
- ▶ Choose architecture or network topology.
- ▶ Meaning, define layers, its parameters.
- ▶ There are multiple API ways to define the model (will look at later)

```
...
2 # define the model
model = ...
```

## Compile the Model

- ▶ Select loss function that you want to optimize, eg Cross Entropy or Mean Squared Error
- ▶ Select Optimization method, eg Adam, Stochastic Gradient Descent
- ▶ Select performance metrics to be used during Training

```
1 ...
2
3 opt = SGD(learning_rate=0.01, momentum=0.9)
model.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])
```

## Fit the Model

- ▶ Select Training configuration (epochs, batch size, etc)
- ▶ \*Epochs: number of full cycles (forward + backward) during training
- ▶ \*Batch Size: Number of samples used to estimate model error, or update the weights
- ▶ Can take minutes to hours to days depending on complexity, hardware, training samples size.
- ▶ Progress bar shows status of each epoch, performance, etc.

```
2 ...
# fit the model
model.fit(X, y, epochs=100, batch_size=32)
```

## Evaluate the Model

- ▶ Select a holdout dataset (cross validation)
- ▶ This is not used for training but just for evaluation as it has correct answers as well.

```
1 ...
# evaluate the model
3 loss = model.evaluate(X, y, verbose=0)
```

## Making Predictions

- ▶ Get Test set for which answers have to be found out.
- ▶ Better to save the model and later load it to make predictions.
- ▶ May choose to fit a model on all of the available data before you start using it.

```
1 ...
# make a prediction
3 yhat = model.predict(X)
```

# Model Definition

## API styles

- ▶ The Sequential Model
  - ▶ Dead simple
  - ▶ Only for single-input, single-output, sequential layer stacks
  - ▶ Good for 70+% of use cases
- ▶ The functional API
  - ▶ Like playing with Lego bricks
  - ▶ Multi-input, multi-output, arbitrary static graph topologies
  - ▶ Good for 95% of use cases
- ▶ Model subclassing
  - ▶ Maximum flexibility
  - ▶ Larger potential error surface

## Sequential Model API (Simple)

- ▶ Called “Sequential” because it involves using Sequential class and adding layers to it one-by-one, in a sequence.
- ▶ E.g. 8 inputs, one hidden layer with 10 nodes, and one output layer with one node to predict numerical value would look:

```
1 # example of a model defined with the sequential api
2 from tensorflow.keras import Sequential
3 from tensorflow.keras.layers import Dense
4 # define the model
5 model = Sequential()
6 model.add(Dense(10, input_shape=(8,)))
7 model.add(Dense(1))
```

Note:

- ▶ Input layer, per say, is NOT added. Its an argument for the first HIDDEN layer.
- ▶ Here ‘input\_shape’ of (8, ) means one sample/row is of 8 values. And such, many samples/rows can come, so left blank.

## Functional Model API (Advanced)

- ▶ Need to explicitly connections between layers.
- ▶ Models may have multiple input/output paths (a word and a number)
- ▶ Input layer needs to be defined explicitly, like:

```
1 x_in = Input(shape=(8,))
```

- ▶ Next, a fully connected layer can be connected to the input by calling the layer and passing the input layer. This will return a reference to the output connection in this new layer.

```
1 x = Dense(10)(x_in)
```

- ▶ Once connected, we define a Model object and specify the input and output layers.

```
1 x_in = Input(shape=(8,))
2 x = Dense(10)(x_in)
3 x_out = Dense(1)(x)
4 # define the model
5 model = Model(inputs=x_in, outputs=x_out)
```

## Sub-classing Model API (Very Advanced)

```
1 class MyModel(tf.keras.Model):
2     def __init__(self, num_classes=10):
3         super(MyModel, self).__init__(name='my_model')
4         self.dense_1 = layers.Dense(32, activation='relu')
5         self.dense_2 = layers.Dense(num_classes, activation='sigmoid')
6
7     def call(self, inputs):
8         # Define your forward pass here,
9         x = self.dense_1(inputs)
10        return self.dense_2(x)
```

## Understanding deferred (symbolic) vs. eager (imperative)

- ▶ Deferred: Build a computation graph that gets compiled first and then once values are filled, executed later
- ▶ Eager: Model is a python exe, Execution is runtime (like Numpy)
- ▶ Deferred: Symbolic tensors don't have a value in your Python code (yet)
- ▶ Eager: tensors have a value in your Python code
- ▶ Eager: can use value-dependent dynamic topologies (tree-RNNs)

## Sample eager execution code

```
lstm_cell = tf.keras.layers.LSTMCell(10)
2
def fn(input, state):
4    return lstm_cell(input, state)

6 input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2
lstm_cell(input, state); fn(input, state) # warm up
8 # benchmark
timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
```

## Let's make this faster

```
1 lstm_cell = tf.keras.layers.LSTMCell(10)
2
3 @tf.function
4 def fn(input, state):
5     return lstm_cell(input, state)
6
7 input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2
8 lstm_cell(input, state); fn(input, state) # warm up
9 # benchmark
10 timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
```

# AutoGraph makes this possible

Say, for a sample function

```
1 @tf.function
2 def f(x):
3     while tf.reduce_sum(x) > 1:
4         x = tf.tanh(x)
5     return x
6 # you never need to run this (unless curious)
7 print(tf.autograph.to_code(f))
```

## Generated code

We need not understand this, but still ...

```
1 def tf__f(x):
2     def loop_test(x_1):
3         with ag__.function_scope('loop_test'):
4             return ag__.gt(tf.reduce_sum(x_1), 1)
5     def loop_body(x_1):
6         with ag__.function_scope('loop_body'):
7             with ag__.utils.control_dependency_on_returns(tf.print(x_1)):
8                 tf_1, x = ag__.utils.alias_tensors(tf, x_1)
9                 x = tf_1.tanh(x)
10                return x,
11 x = ag__.while_stmt(loop_test, loop_body, (x,), (tf,))
12 return x
```

## tf .distribute.Strategy

For the sample code below ...

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, input_shape=[10]),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')])
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

# Multi-GPU

One of the computations distribution strategy could be ...

```
1 strategy = tf.distribute.MirroredStrategy()
with strategy.scope():
2     model = tf.keras.models.Sequential([
3         tf.keras.layers.Dense(64, input_shape=[10]),
4         tf.keras.layers.Dense(64, activation='relu'),
5         tf.keras.layers.Dense(10, activation='softmax')])
6     model.compile(optimizer='adam',
7                     loss='categorical_crossentropy',
8                     metrics=['accuracy'])
```

# TensorFlow Datasets

- audio
  - "nsynth"
- image
  - "cifar10"
  - "diabetic\_retinopathy\_detection"
  - "imagenet2012"
  - "mnist"
- structured
  - "titanic"
- text
  - "imdb\_reviews"
  - "lm1b"
  - "squad"
- translate
  - "wmt\_translate\_ende"
  - "wmt\_translate\_enfr"
- video
  - "bair\_robot\_pushing\_small"
  - "moving\_mnist"
  - "starcraft\_video"

More at [tensorflow.org/datasets](https://tensorflow.org/datasets)

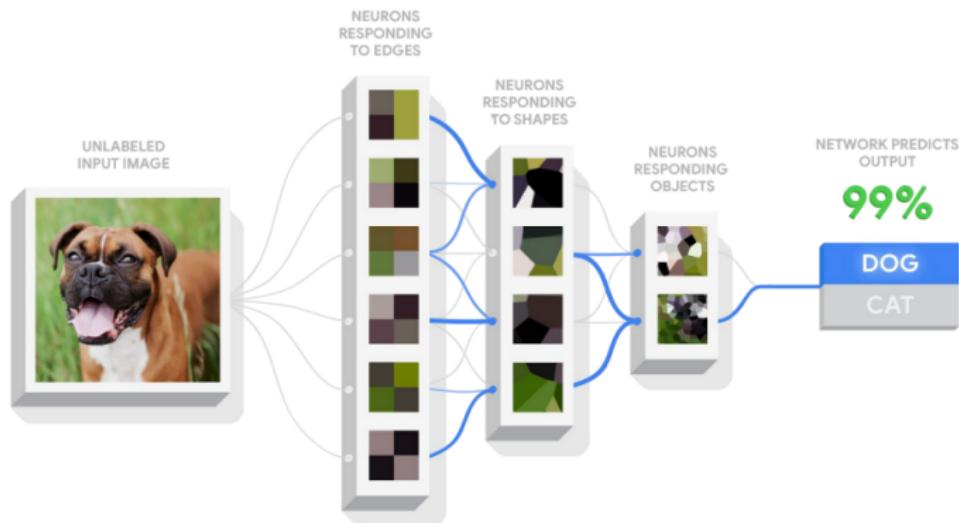
(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

## Terminologies

In the neural network terminology:

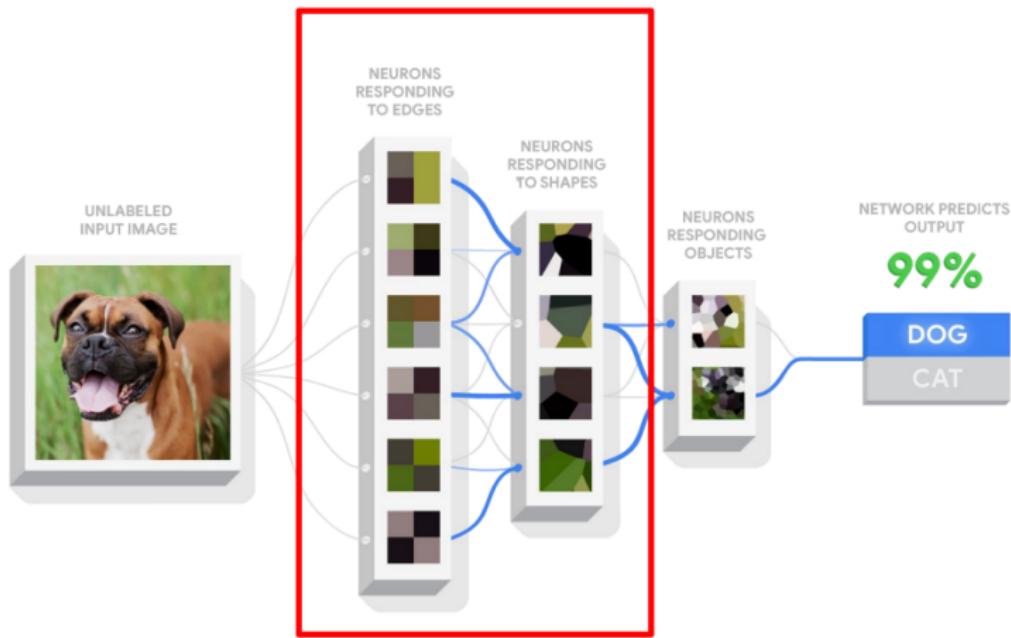
- ▶ one epoch = one forward pass and one backward pass of all the training examples
- ▶ batch size = the number of training examples in one forward/backward pass. The higher the batch size, the more memory space you'll need.
- ▶ number of iterations = number of passes, each pass using [batch size] number of examples. To be clear, one pass = one forward pass + one backward pass (we do not count the forward pass and backward pass as two different passes).
- ▶ Example: if you have 1000 training examples, and your batch size is 500, then it will take 2 iterations to complete 1 epoch.

# Transfer Learning



(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

# Transfer Learning



(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

# Transfer Learning

```
1 import tensorflow as tf
2 base_model = tf.keras.applications.SequentialMobileNetV2(
3     input_shape=(160, 160, 3),
4     include_top=False,
5     weights='imagenet')
6 base_model.trainable = False
7 model = tf.keras.models.Sequential([
8     base_model,
9     tf.keras.layers.GlobalAveragePooling2D(),
10    tf.keras.layers.Dense(1)
11])
12 # Compile and fit
```

(Ref: Introduction to TensorFlow 2.0 - Brad Miro)

# Transfer Learning

≡ TensorFlow Hub 🔍 USER GUIDE

**Text**

Embedding

**Image**

Classification

Feature Vector

Generator

**Video**

Classification

**Publishers**

Google

DeepMind

**Text embedding**

 **universal-sentence-encoder** By Google  
text-embedding DAN en  
Encoder of greater-than-word length text trained on a variety of data.

 **nnlm-en-dim128** By Google  
text-embedding Google News NNLM en  
Token based text embedding trained on English Google News 200B corpus.

 **elmo** By Google  
text-embedding 1 Billion Word Benchmark ELMo en  
Embeddings from a language model trained on the 1 Billion Word Benchmark.

[View more text embeddings](#)

**Image feature vectors**

 **imagenet/inception\_v3/feature\_vector** By Google  
image-feature-vector ImageNet (ILSVRC-2012-CLS) Inception V3  
Feature vectors of images with Inception V3 trained on ImageNet (ILSVRC-2012-CLS).

## Learning More . . .

- ▶ Latest tutorials and guides at <http://tensorflow.org/beta>
- ▶ Book: Hands-on ML with Scikit-Learn, Keras and TensorFlow (2nd edition)

(Ref: Intro to TensorFlow 2.0 - Josh Gordon)

## References

Many publicly available resources have been refereed for making this presentation. Some of the notable ones are:

- ▶ TensorFlow 2 Tutorial: Get Started in Deep Learning With tf.keras - Jason Brownlee
- ▶ Deep Learning using Keras- Alyosamah
- ▶ Introduction to Keras - Francois Chollet
- ▶ Michael Nielsen's Neural Networks and Deep Learning: <http://neuralnetworksanddeeplearning.com/>

Thanks ... yogeshkulkarni@yahoo.com